

This notebook is based on chapter 9 in Statistical Rethinking.

Care and feeding of Markov chains

```

• html"""
• <style>
•   main {
•     margin: 0 auto;
•     max-width: 2000px;
•     padding-left: max(160px, 10%);
•     padding-right: max(100px, 10%);
•   }
• </style>
• """

```

```

• using Pkg ✓

```

```

• begin
•   # Specific to ROSStanPluto
•   using StanSample ✓
•
•   # Graphics related
•   using GLMakie ✓
•
•   # Include basic packages
•   using RegressionAndOtherStories ✓
• end

```

```

Replacing docs for `RegressionAndOtherStories.trankpl
rames.DataFrame, AbstractString}` in module `Regressi
\

```

```

• stan9_2 = "
• data {
•     int n;
•     vector[n] y;
• }
• parameters {
•     real alpha;
•     real<lower=0> sigma;
• }
• model {
•     real mu;
•     alpha ~ normal(0, 1000);
•     sigma ~ exponential(0.0001);
•     mu = alpha;
•     y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std	5%	
1	"alpha"	-22.243	32.0576	478.952	-874.001	0
2	"sigma"	798.287	93.8203	1672.87	10.4871	2

```

• let
•     Random.seed!(123)
•     data = (n=2, y=[-1, 1])
•     global m9_2s = SampleModel("m9_2s", stan9_2)
•     global rc9_2s = stan_sample(m9_2s; data)
•     success(rc9_2s) && describe(m9_2s)
• end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/jl_e
updated.

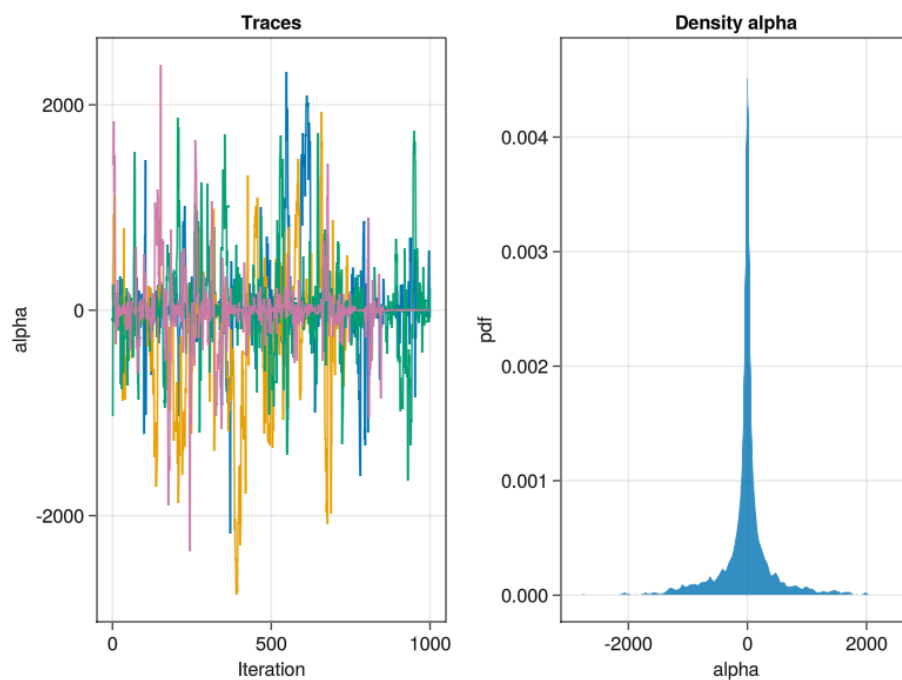
```

	parameters	median	mad_sd	mean	std
1	"alpha"	0.762	127.055	-22.243	478.952
2	"sigma"	231.207	293.574	798.287	1672.87

```

• if success(rc9_2s)
•     post9_2s = read_samples(m9_2s, :dataframe)
•     ms9_2s = model_summary(post9_2s, [:alpha,
•         :sigma])
• end


```



```

• plot_chains(post9_2s, [:alpha])

```



```
• trankplot(post9_2s, "alpha")
```

```
• stan9_3 = "
• data {
•   int n;
•   vector[n] y;
• }
• parameters {
•   real alpha;
•   real<lower=0> sigma;
• }
• model {
•   real mu;
•   alpha ~ normal(0, 1);
•   sigma ~ exponential(1);
•   mu = alpha;
•   y ~ normal(mu, sigma);
• }";
```

	parameters	mean	mcse	std	5
1	"alpha"	-0.000934169	0.0136699	0.650532	-1.0
2	"sigma"	1.37925	0.0139107	0.672726	0.6

```

• let
•     Random.seed!(123)
•     data = (n=2, y=[-1, 1])
•     global m9_3s = SampleModel("m9.3s", stan9_3)
•     global rc9_3s = stan_sample(m9_3s; data)
•     success(rc9_3s) && describe(m9_3s)
• end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gn/T/jl_H
updated.

```

	alpha	sigma
1	0.217657	0.665899
2	0.328003	2.86695
3	0.923458	2.11215
4	0.476047	2.58393
5	1.01984	2.92127
6	-0.384908	0.835844
7	0.0187461	1.16952
8	-1.13623	1.02151
9	0.00714906	1.11492
10	0.246382	0.647248
⋮ more		
4000	-0.0506899	0.604653

```

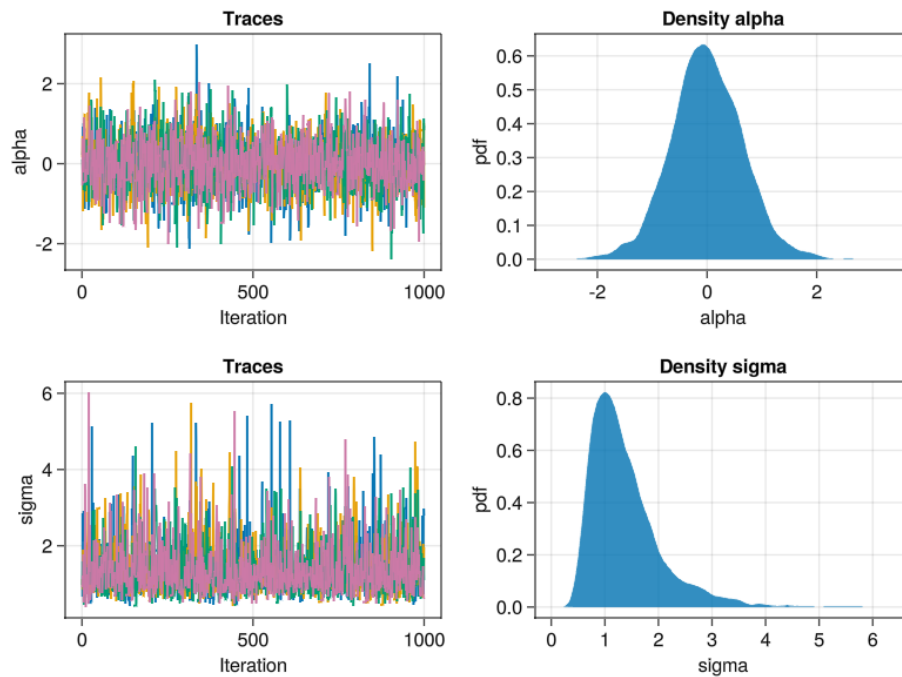
• if success(rc9_3s)
•   post9_3s = read_samples(m9_3s, :dataframe)
• end

```


```
ms9_3s =
```

	parameters	median	mad_sd	mean	std
1	"alpha"	-0.017	0.624	-0.001	0.651
2	"sigma"	1.224	0.518	1.379	0.673

```
• ms9_3s = model_summary(post9_3s, [:alpha, :sigma])
```



```
• plot_chains(post9_3s, [:alpha, :sigma])
```



```
• trankplot(post9_3s, "alpha")
```

```

• stan9_4 = "
• data {
•     int n;
•     vector[n] y;
• }
• parameters {
•     real alpha;
•     real beta;
•     real<lower=0> sigma;
• }
• model {
•     real mu;
•     alpha ~ normal(0, 100);
•     beta ~ normal(0, 1000);
•     sigma ~ exponential(1);
•     mu = alpha + beta;
•     y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std	5%
1	"alpha"	-6.7726	10.2174	104.782	-183.4
2	"beta"	6.77717	10.2173	104.781	-163.7
3	"sigma"	1.09611	0.00527366	0.0783639	0.9755

```

• begin
•     Random.seed!(1)
•     data9_4s = (n = 100, y = rand(Normal(0, 1),
• 100))
•     m9_4s = SampleModel("m9.4s", stan9_4)
•     rc9_4s = stan_sample(m9_4s; data=data9_4s)
•     success(rc9_4s) && describe(m9_4s)
• end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gn/T/jl_d
updated.

```


	alpha	beta	sigma
1	-66.6067	66.668	1.01229
2	-85.607	85.6065	1.04282
3	-85.7241	85.6941	1.01502
4	-82.1121	82.1377	1.10425
5	-52.1738	52.1297	1.08983
6	-22.5539	22.5254	1.0715
7	-9.34204	9.33467	1.05689
8	-16.9432	16.9596	0.975596
9	-25.6934	25.6744	0.967967
10	-35.2967	35.3679	0.977145
⋮ more			
4000	54.3078	-54.3876	1.10962

```

• if success(rc9_4s)
•   post9_4s = read_samples(m9_4s, :dataframe)
• end

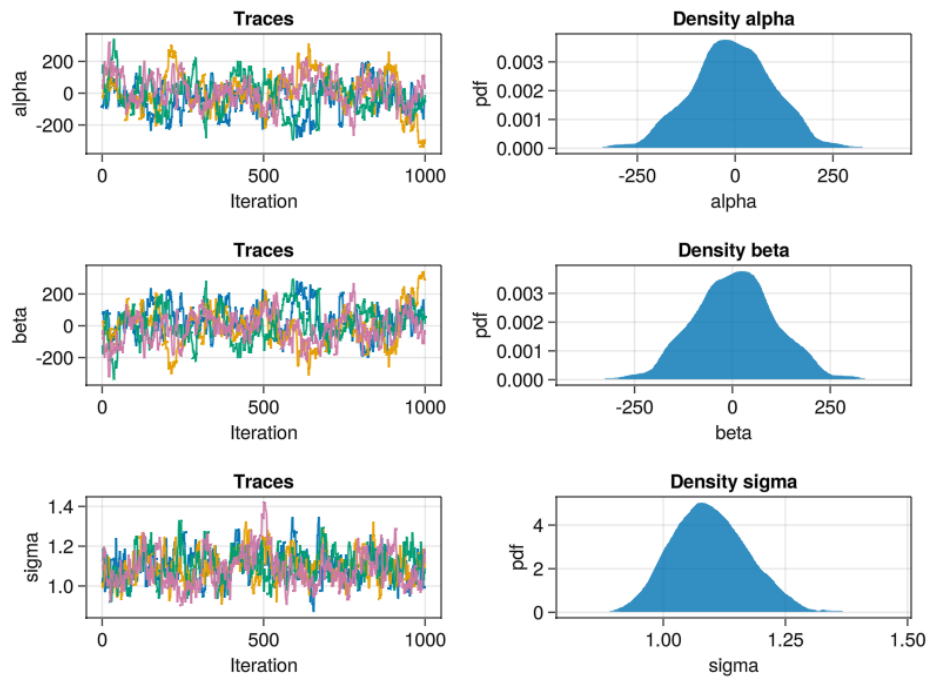
```

	parameters	median	mad_sd	mean	std
1	"alpha"	-7.828	101.512	-6.773	104.782
2	"beta"	7.787	101.535	6.777	104.781
3	"sigma"	1.092	0.08	1.096	0.078


```

• model_summary(post9_4s, [:alpha, :beta, :sigma])


```



```
• plot_chains(post9_4s, [:alpha, :beta, :sigma])
```



```
• frankplot(post9_4s, "alpha")
```



```
• trankplot(post9_4s, "beta")
```

```
• stan9_5 = "
• data {
•   int n;
•   vector[n] y;
• }
• parameters {
•   real alpha;
•   real beta;
•   real<lower=0> sigma;
• }
• model {
•   real mu;
•   alpha ~ normal(0, 10);
•   beta ~ normal(0, 10);
•   sigma ~ exponential(1);
•   mu = alpha + beta;
•   y ~ normal(mu, sigma);
• }";
```

	parameters	mean	mcse	std	5%
1	"alpha"	0.0015794	0.254173	7.23111	-11.1
2	"beta"	0.00349905	0.254632	7.23157	-11.1
3	"sigma"	1.09401	0.0020217	0.0809901	0.96

```

• begin
•     # Re-use data from m9_4s
•     m9_5s = SampleModel("m9.5s", stan9_5)
•     rc9_5s = stan_sample(m9_5s; data=data9_4s)
•     success(rc9_5s) && describe(m9_5s)
• end

```

```

/ Informational Message: The current Metropolis proposal
l be rejected because of the following issue:
Exception: normal_lpdf: Scale parameter is 0, but mean is not
in '/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/jl_1000000000/
stan', line 16, column 4 to column 26)
If this warning occurs sporadically, such as for high-dimensional
variable types like covariance matrices, then the sampler may work,
but if this warning occurs often then your model may be
ly ill-conditioned or misspecified.

Informational Message: The current Metropolis proposal
rejected because of the following issue:
Exception: normal_lpdf: Scale parameter is 0, but mean is not
in '/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/jl_1000000000/
stan', line 16, column 4 to column 26)
If this warning occurs sporadically, such as for high-dimensional
variable types like covariance matrices, then the sampler may work,
but if this warning occurs often then your model may be
ly ill-conditioned or misspecified.

```

	alpha	beta	sigma
1	5.56982	-5.40305	1.0352
2	-3.56507	3.55266	1.17554
3	-8.14825	8.10198	1.11049
4	-8.08085	8.15674	1.1103
5	-6.09395	6.0804	1.09802
6	-0.718372	0.651115	1.15697
7	6.00263	-6.09916	1.24334
8	10.2419	-10.3083	1.20502
9	10.6701	-10.7259	1.18993
10	5.47284	-5.41954	0.948773
⋮ more			
4000	6.57878	-6.69215	1.13901

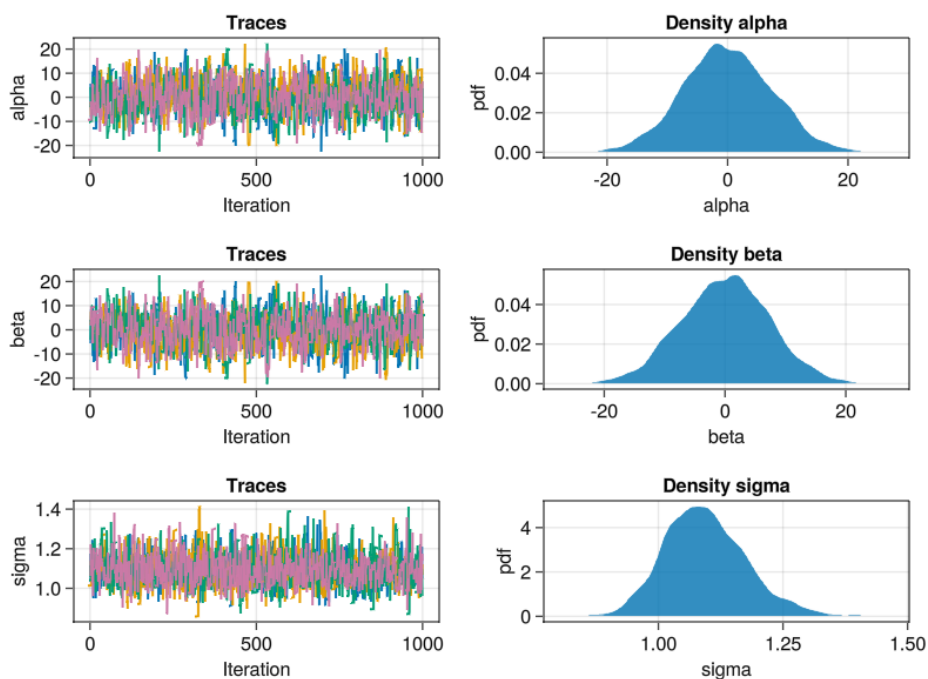
```

• if success(rc9_5s)
•     sdf9_5s = read_summary(m9_5s)
•     post9_5s = read_samples(m9_5s, :dataframe)
• end


```

	parameters	mean	mcse	std
1	:lp__	-60.0227	0.038637	1.23647
2	:accept_stat__	0.934843	0.00144222	0.0981003
3	:stepsize__	0.0110184	0.00037207	0.0005269
4	:treedepth__	6.35575	0.0355616	2.15682
5	:n_leapfrog__	204.535	2.63592	154.241
6	:divergent__	0.0	NaN	0.0
7	:energy__	61.5596	0.0530301	1.79131
8	:alpha	0.0015794	0.254173	7.23111
9	:beta	0.00349905	0.254632	7.23157
10	:sigma	1.09401	0.0020217	0.0809901

• [sdf9_5s](#)



• `plot_chains(post9_5s, [:alpha, :beta, :sigma])`



```
• trankplot(post9_5s, "alpha")
```

	variable	mean	min	median	m
1	:rgdppc_2000	9094.89	466.647	5314.74	5779
2	:log_gdp	8.51712	6.14557	8.57823	10.9
3	:log_gdp_s	1.0	0.721556	1.00718	1.28
4	:rugged	1.33318	0.003	0.9795	6.20
5	:rugged_s	0.21496	0.000483715	0.157933	1.0
6	:cid	1.71176	1	2.0	2

```

• begin
•   df = CSV.read(ros_datadir("SR2",
•     "rugged.csv"), DataFrame)
•   dropmissing!(df, :rgdppc_2000)
•   dropmissing!(df, :rugged)
•   df.log_gdp = log.(df[:, :rgdppc_2000])
•   df.log_gdp_s = df.log_gdp / mean(df.log_gdp)
•   df.rugged_s = df.rugged / maximum(df.rugged)
•   df.cid = [df.cont_africa[i] == 1 ? 1 : 2 for i
•     in 1:size(df, 1)]
•   r̄ = mean(df.rugged_s)
•   describe(df[:, [:rgdppc_2000, :log_gdp,
•     :log_gdp_s, :rugged, :rugged_s, :cid]])
end

```

```

• data8_3s = (N = size(df, 1), K =
•   length(unique(df.cid)), G = df.log_gdp_s, R =
•   df.rugged_s, cid=df.cid);

```



```
stan8_3 = "  
data {  
  int N;  
  int K;  
  vector[N] G;  
  vector[N] R;  
  int cid[N];  
}  
  
parameters {  
  vector[K] a;  
  vector[K] b;  
  real<lower=0> sigma;  
}  
  
transformed parameters {  
  vector[N] mu;  
  for (i in 1:N)  
    mu[i] = a[cid[i]] + b[cid[i]] * (R[i] -  
    $(f));  
}  
  
model {  
  a ~ normal(1, 0.1);  
  b ~ normal(0, 0.3);  
  sigma ~ exponential(1);  
  G ~ normal(mu, sigma);  
}  
";
```

	parameters	mean	mcse	std	5%	
1	"a[1]"	0.89	0.00023	0.016	0.86	0.92
2	"a[2]"	1.1	0.00014	0.01	1.0	1.2
3	"b[1]"	0.13	0.0011	0.076	0.0038	0.28
4	"b[2]"	-0.14	0.00084	0.055	-0.23	-0.05
5	"sigma"	0.11	0.0	0.01	0.1	0.13
6	"mu[1]"	0.88	0.00023	0.016	0.85	0.91
7	"mu[2]"	1.0	0.0003	0.02	0.97	1.03
8	"mu[3]"	1.1	0.00017	0.012	1.0	1.2
9	"mu[4]"	1.1	0.00017	0.012	1.0	1.2
10	"mu[5]"	1.0	0.00021	0.015	1.0	1.1
	⋮ more					
175	"mu[170]"	0.88	0.00023	0.016	0.86	0.90

```

• begin
•     m8_3s = SampleModel("m8.3s", stan8_3)
•     rc8_3s = stan_sample(m8_3s; data=data8_3s)
•     success(rc8_3s) && describe(m8_3s)
• end

```

```
Informational Message: The current Metropolis proposal  
be rejected because of the following issue:  
Exception: normal_lpdf: Scale parameter is 0, but mean is  
in '/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/jl_  
stan', line 25, column 1 to column 23)  
If this warning occurs sporadically, such as for high  
variable types like covariance matrices, then the sampler  
but if this warning occurs often then your model may be  
ly ill-conditioned or misspecified.
```

```
Informational Message: The current Metropolis proposal  
rejected because of the following issue:  
Exception: normal_lpdf: Scale parameter is 0, but mean is  
in '/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/jl_  
stan', line 25, column 1 to column 23)  
If this warning occurs sporadically, such as for high  
variable types like covariance matrices, then the sampler  
but if this warning occurs often then your model may be  
ly ill-conditioned or misspecified.
```

	sigma	a	b
1	0.115896	▶ [0.898105, 1.05505]	▶ [0.226078, -(
2	0.105665	▶ [0.879329, 1.04336]	▶ [0.0432967, .
3	0.108974	▶ [0.876619, 1.04594]	▶ [0.0522773, .
4	0.110399	▶ [0.88959, 1.04211]	▶ [0.156145, -(
5	0.110872	▶ [0.884438, 1.06401]	▶ [0.131963, -(
6	0.106499	▶ [0.88545, 1.0381]	▶ [0.125918, -(
7	0.112394	▶ [0.898236, 1.06481]	▶ [0.178581, -(
8	0.105976	▶ [0.883061, 1.03549]	▶ [0.108057, -(
9	0.111567	▶ [0.869031, 1.06127]	▶ [0.218732, -(
10	0.12295	▶ [0.871215, 1.0709]	▶ [0.205622, -(

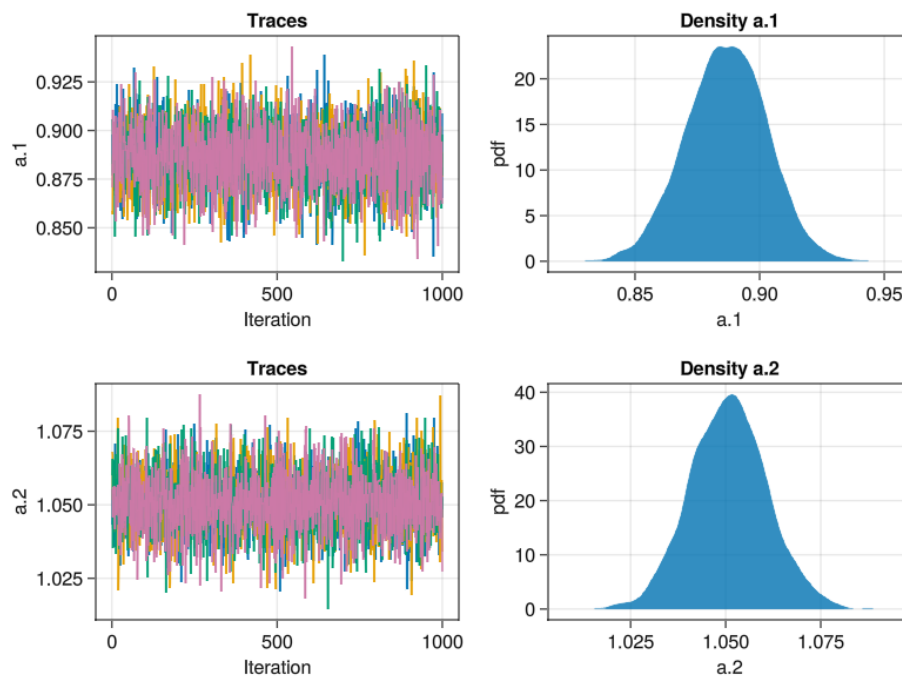
```

• if success(rc8_3s)
•     post8_3s = read_samples(m8_3s, :dataframe)
•     nd8_3s = read_samples(m8_3s, :nesteddataframe)
• end


```

```
4000×170 Matrix{Float64}:
 0.880784  1.00967  1.06728  1.06715  ...  1.0336  0.9
 0.876012  0.98476  1.05915  1.05898  ...  1.01566  0.8
 0.872613  0.977665  1.06433  1.06414  ...  1.01366  0.8
 0.877626  1.01543  1.04929  1.04922  ...  1.0295  0.9
 0.874327  1.04942  1.06794  1.0679  ...  1.05711  0.8
 0.875803  0.963867  1.0581  1.05789  ...  1.00301  0.8
 0.884553  1.03772  1.07211  1.07203  ...  1.05201  0.9
  ...
 0.858055  0.994668  1.06938  1.06921  ...  1.0257  0.8
 0.876568  0.999425  1.06515  1.065  ...  1.02672  0.8
 0.868499  1.00966  1.05866  1.05855  ...  1.03001  0.8
 0.884145  0.994123  1.07014  1.06997  ...  1.0257  0.9
 0.862615  1.00797  1.05886  1.05874  ...  1.02911  0.8
 0.853886  0.978273  1.04031  1.04017  ...  1.00404  0.8
```





```
• array(nd8_3s, :mu)
```




```
• plot_chains(post8_3s, [Symbol("a.1"),
Symbol("a.2")])
```



```
• trankplot(post8_3s, "a.1")
```



```
• plot_chains(post8_3s, [Symbol("b.1"),  
Symbol("b.2")])
```



```
• plot_chains(post8_3s, [:sigma])
```