

See chapter 1 in Regression and Other Stories.

Widen the cells.

```
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

A typical set of Julia packages to include in
notebooks.

```
• using Pkg ✓ , DrWatson ✓
```

```
• begin
•   # Specific to this notebook
•   using GLM ✓
•
•   # Specific to ROSStanPluto
•   using StanSample ✓
•
•   # Graphics related
•   using CairoMakie ✓
•   using AlgebraOfGraphics ✓
•
•   # Include basic packages
•   using RegressionAndOtherStories ✓
• end
```

```
Replacing docs for `RegressionAndOtherStories.
names.DataFrame, AbstractString}` in module `R
\`
```

1.1 The three challenges of statistics.

Note

It is not common for me to copy from the book but this particular section deserves an exception!

The three challenges of statistical inference are:

1. Generalizing from sample to population, a problem that is associated with survey sampling but actually arises in nearly every application of statistical inference;
2. Generalizing from treatment to control group, a problem that is associated with causal inference, which is implicitly or explicitly part of the interpretation of most regressions we have seen; and
3. Generalizing from observed measurements to the underlying constructs of interest, as most of the time our data do not record exactly what we would ideally like to study.

All three of these challenges can be framed as problems of prediction (for new people or new items that are not in the sample, future outcomes under different potentially assigned treatments, and underlying constructs of interest, if they could be measured exactly).

1.2 Why learn regression?

`hibbs =`

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
⋮	more			
16	2012	0.95	52.0	"Obama"

- `hibbs =`
`CSV.read(ros_datadir("ElectionsEconomy",`
`"hibbs.csv"), DataFrame)`

`hibbs_lm =`

`StatsModels.TableRegressionModel{LinearModel{GLM`

`vote ~ 1 + growth`

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	46.2476	1.62193	28.51	<1e-16
growth	3.06053	0.696274	4.40	0.00044

- `hibbs_lm = lm(@formula(vote ~ growth),`
`hibbs)`

► `[-8.99292, 2.66743, 1.0609, 2.20753, -5.89044,`

- `residuals(hibbs_lm)`

2.2744434224582912



- `mad(residuals(hibbs_lm))`

3.635681268522063

- `std(residuals(hibbs_lm))`

► [46.2476, 3.06053]

- `coef(hibbs_lm)`



```

• let
•     fig = Figure()
•     hibbs.label = string.(hibbs.year)
•     xlabel = "Average growth personal
•     income [%]"
•     ylabel = "Incumbent's party vote share"
•     let
•         title = "Forecasting the election
•         from the economy"
•         ax = Axis(fig[1, 1]; title, xlabel,
•         ylabel)
•         for (ind, yr) in
•         enumerate(hibbs.year)
•             annotations!("$ (yr)"; position=
•             (hibbs.growth[ind],
•             hibbs.vote[ind]), fontsize=10)
•         end
•     end
•     let
•         x = LinRange(-1, 4, 100)
•         title = "Data and linear fit"
•         ax = Axis(fig[1, 2]; title, xlabel,
•         ylabel)
•         scatter!(hibbs.growth, hibbs.vote)
•         lines!(x, coef(hibbs_lm)[1] .+
•         coef(hibbs_lm)[2] .* x;
•         color=:darkred)
•         annotations!("vote = 46.2 + 3.0 *
•         growth"; position=(0, 41))
•     end
•     fig
• end

```

1.3 Some examples of regression.

Electric company

	post_test	pre_test	grade	t
1	48.9	13.8	1	1
2	70.5	16.5	1	1
3	89.7	18.5	1	1
4	44.2	8.8	1	1
5	77.5	15.3	1	1
6	84.7	15.0	1	1
7	78.9	19.4	1	1
8	86.8	15.0	1	1
9	60.8	11.8	1	1
10	75.7	16.4	1	1
⋮	more			
192	110.0	102.6	4	0

```

• begin
•     electric =
•     CSV.read(ros_datadir("ElectricCompany"
•     , "electric.csv"), DataFrame)
•     electric = electric[:, [:post_test,
•     :pre_test, :grade, :treatment]]
•     electric.grade =
•     categorical(electric.grade)
•     electric.treatment =
•     categorical(electric.treatment)
•     electric
• end

```

A quick look at the overall values of pre_test and post_test.


	variable	mean	min	median	max
1	:post_test	97.1495	44.2	102.3	122.0
2	:pre_test	72.2245	8.8	80.75	119.8
3	:grade	nothing	1	nothing	4
4	:treatment	nothing	0	nothing	1

- `describe(electric)`

true

- `all(completeness(electric)) == true`


**Post-test density for each grade
conditioned on treatment.**



```

• let
•     f = Figure()
•     axis = (; width = 150, height = 150)
•     el = data(electric) *
•         mapping(:post_test, col=:grade,
•             color=:treatment)
•     plt = el *
•         AlgebraOfGraphics.histogram(;bins=20)
•         * mapping(row=:treatment)
•     draw!(f[1, 1], plt; axis)
•     f
• end

```



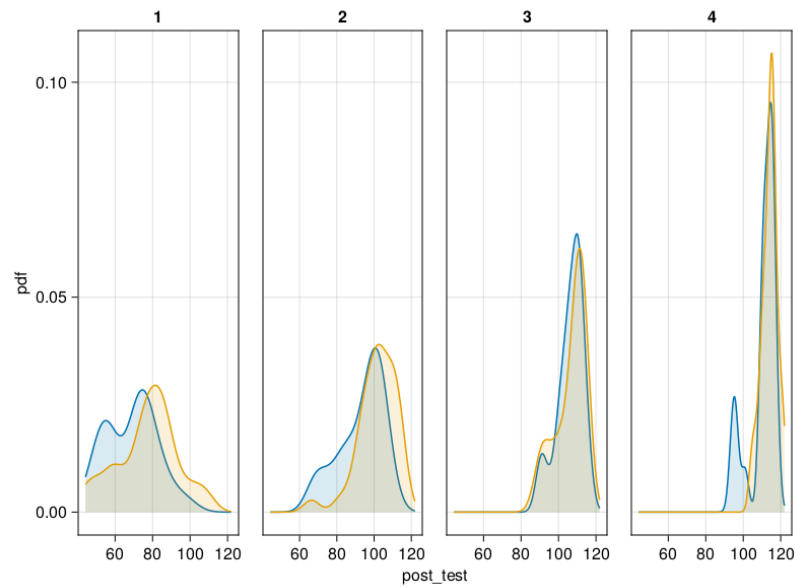
```

• let
•   f = Figure()
•   axis = (; width = 150, height = 150)
•   el = data(electric) *
•   mapping(:post_test, col=:grade,
•   color=:treatment)
•   plt = el * AlgebraOfGraphics.density()
•   * mapping(row=:treatment)
•   draw!(f[1, 1], plt; axis)
•   f
end

```

Note


In above cell, as `density()` is exported by both `GLMakie` and `AlgebraOfGraphics`, it needs to be qualified.



```

• let
•     f = Figure()
•     el = data(electric) *
•     mapping(:post_test, col=:grade)
•     plt = el * AlgebraOfGraphics.density()
•     * mapping(color=:treatment)
•     draw!(f[1, 1], plt)
•     f
• end

```



```
• let
•   plt = data(electric) * visual(Violin)
•   * mapping(:grade, :post_test,
•   dodge=:treatment, color=:treatment)
•   draw(plt)
end
```

Peacekeeping

peace =

	war	cfdate	faildate
1	"Afghanistan-Mujahideen"	8150	8257
2	"Afghanistan-Taliban"	8466	8505
3	"Algeria-FIS/AIS"	10149	12783
4	"Angola"	7820	8319
5	"Angola"	9089	10564
6	"Azerbaijan-N.K."	8643	8678
7	"Azerbaijan-N.K."	8901	12783
8	"Bangladesh-CHT"	8248	12783
9	"Myanmar-Karen"	8153	9282
10	"Myanmar-Karen"	9296	9907
⋮	more		
96	"Yugoslavia-Kosovo"	10751	12783

```
• peace =
  CSV.read(ros_datadir("PeaceKeeping",
    "peacekeeping.csv"), missingstring="NA",
    DataFrame)
```

	variable	mean	min
1	:war	nothing	"Afghanistan-Mujal
2	:cfdate	8925.1	6985
3	:faildate	10795.8	7074
4	:peacekeepers	0.354167	0
5	:badness	-8.15228	-12.26
6	:delay	5.12177	0.04
7	:censored	0.416667	0

- `describe(peace)`

A quick look at this Dates stuff!

8150

- `peace.cfdate[1]`

1992-04-25T00:00:00

- `DateTime(1992, 4, 25)`

107 days

- `Date(1992, 8, 10) - Date(1992, 4, 25)`

1970-01-01

- `Date(1970,1,1)`

1992-04-25

- `Date(1970,1,1) + Dates.Day(8150)`

8150 days

- `Date(1992, 4, 25) - Date(1970, 1, 1)`

107

- `peace.faildate[1] - peace.cfdate[1]`

```

• begin
•   pks_df = peace[peace.peacekeepers .==
•     1, [:cfdate, :faildate]]
•   nopks_df = peace[peace.peacekeepers
•     .== 0, [:cfdate, :faildate]]
end;

```

```
0.4166666666666667
```

```
• mean(peace.censored)
```

```
64
```

```
• length(unique(peace.war))
```

```
0.5588235294117647
```

```
• mean(peace[peace.peacekeepers .== 1,
  :censored])
```

```
0.3387096774193548
```

```
• mean(peace[peace.peacekeepers .== 0,
  :censored])
```

```
1.382
```

```
• mean(peace[peace.peacekeepers .== 1 .&&
  peace.censored .== 0, :delay])
```

```
1.5153658536585364
```

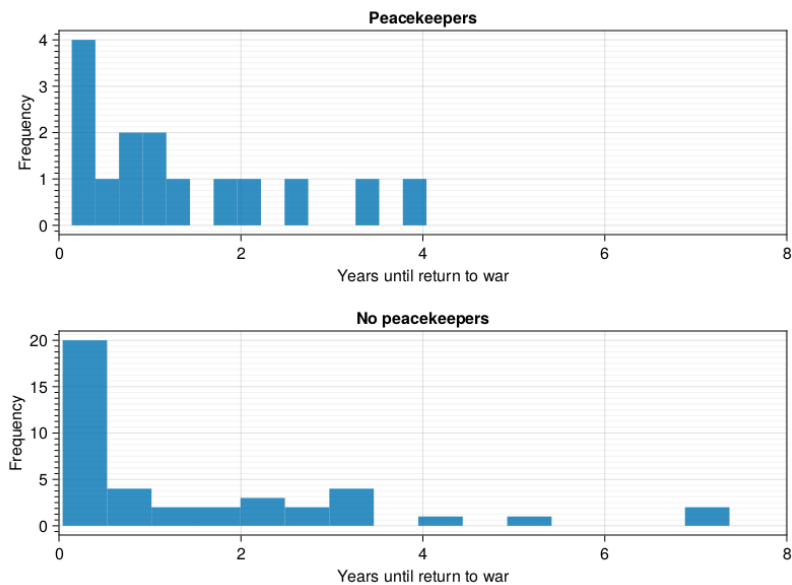
```
• mean(peace[peace.peacekeepers .== 0 .&&
  peace.censored .== 0, :delay])
```

```
1.05
```

```
• median(peace[peace.peacekeepers .== 1 .&&
  peace.censored .== 0, :delay])
```

```
0.59
```

```
• median(peace[peace.peacekeepers .== 0 .&&
  peace.censored .== 0, :delay])
```



```

let
    f = Figure()
    pks = peace[peace.peacekeepers .== 1
    .&& peace.censored .== 0, :]
    nopks = peace[peace.peacekeepers .== 0
    .&& peace.censored .== 0, :]

    for i in 1:2
        title = i == 1 ? "Peacekeepers" :
            "No peacekeepers"

        ax = Axis(f[i, 1]; title,
            xlabel="Years until return to
            war",
            ylabel = "Frequency",
            yminorticksvisible = true,
            yminorgridvisible = true,
            yminorticks = IntervalsBetween(8))

        xlims!(ax, [0, 8])
        hist!(i == 1 ? pks.delay :
            nopks.delay)
    end
end
f
end

```

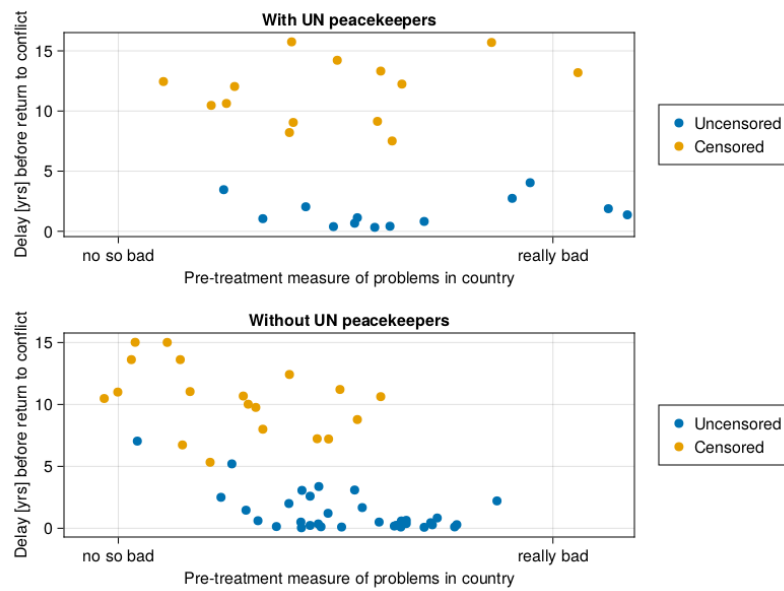

Note

Censored means conflict had not returned until end of observation period (2004).

```

• begin
•   # Filter out missing badness rows.
•   pb = peace[peace.badness .!= missing,
•   :];
•
•   # Delays until return to war for
•   uncensored, peacekeeper cases
•   pks_uc = pb[pb.peacekeepers .== 1 .&&
•   pb.censored .== 0, :delay]
•   # Delays until return to war for
•   censored, peacekeeper cases
•   pks_c = pb[pb.peacekeepers .== 1 .&&
•   pb.censored .== 1, :delay]
•
•   # No peacekeeper cases.
•   nopks_uc = pb[pb.peacekeepers .== 0
•   .&& pb.censored .== 0, :delay]
•   nopks_c = pb[pb.peacekeepers .== 0 .&&
•   pb.censored .== 1, :delay]
•
•   # Crude measure (:badness) used for
•   assessing situation
•   badness_pks_uc = pb[pb.peacekeepers
•   .== 1 .&& pb.censored .== 0,
•   :badness]
•   badness_pks_c = pb[pb.peacekeepers .==
•   1 .&& pb.censored .== 1,
•   :badness]
•   badness_nopks_uc = pb[pb.peacekeepers
•   .== 0 .&& pb.censored .== 0,
•   :badness]
•   badness_nopks_c = pb[pb.peacekeepers
•   .== 0 .&& pb.censored .== 1,
•   :badness]
• end;

```



```

• begin
•   f = Figure()
•   ax = Axis(f[1, 1], title = "With UN
•     peacekeepers",
•       xlabel = "Pre-treatment measure of
•         problems in country",
•       ylabel = "Delay [yrs] before
•         return to conflict")
•   sca1 = scatter!(badness_pks_uc, pks_uc)
•   sca2 = scatter!(badness_pks_c, pks_c)
•   xlims!(ax, [-13, -2.5])
•   Legend(f[1, 2], [sca1, sca2],
•     ["Uncensored", "Censored"])
•   ax.xticks = ([-12, -4], ["no so bad",
•     "really bad"])
•
•
•
•   ax = Axis(f[2, 1], title = "Without UN
•     peacekeepers",
•       xlabel = "Pre-treatment measure of
•         problems in country",
•       ylabel = "Delay [yrs] before
•         return to conflict")
•   sca1 = scatter!(badness_nopks_uc,
•     nopks_uc)
•   sca2 = scatter!(badness_nopks_c,
•     nopks_c)
•   xlims!(ax, [-13, -2.5])
•   Legend(f[2, 2], [sca1, sca2],
•     ["Uncensored", "Censored"])
•   ax.xticks = ([-12, -4], ["no so bad",
•     "really bad"])
•
•
•
•   f
• end

```

1.4 Challenges in building, understanding, and interpreting regression.

Simple causal

Note

In models like below I usually prefer to create 2 separate Stan Language models, one for the continuous case and another for the binary case. But they can be combined in a single model as shown below. I'm using this example to show one way to handle vectors returned from Stan's cmdstan.

```
• stan1_4_1 = "  
• data {  
•   int N;  
•   vector[N] x;  
•   vector[N] x_binary;  
•   vector[N] y;  
• }  
• parameters {  
•   vector[2] a;  
•   vector[2] b;  
•   vector<lower=0>[2] sigma;  
• }  
• model {  
•   // Priors  
•   a ~ normal(10, 10);  
•   b ~ normal(10, 10);  
•   sigma ~ exponential(1);  
•   // Likelihood  
•   y ~ normal(a[1] + b[1] * x, sigma[1]);  
•   y ~ normal(a[2] + b[2] * x_binary,  
•   sigma[2]);  
• }  
• "
```

Note

Aki Vehtari did not include a seed number in his code.

```

• begin
•   Random.seed!(123)
•   n = 50
•   x = rand(Uniform(1, 5), n)
•   x_binary = [x[i] < 3 ? 0 : 1 for i in
•   1:n]
•   y = [rand(Normal(10 + 3x[i], 3), 1)[1]
•   for i in 1:n]
end;

```

	parameters	mean	mcse	std	5'
1	"a[1]"	9.4	0.027	1.4	7.1
2	"a[2]"	16.0	0.013	0.69	15.0
3	"b[1]"	3.3	0.0084	0.43	2.5
4	"b[2]"	7.0	0.019	1.0	5.3
5	"sigma[1]"	3.5	0.0058	0.35	2.9
6	"sigma[2]"	3.7	0.0064	0.37	3.1

```

• let
•   data = (N = n, x = x, x_binary =
•   x_binary, y = y)
•   global m1_4_1s = SampleModel("m1_4_1s",
•   stan1_4_1);
•   global rc1_4_1s = stan_sample(m1_4_1s;
•   data)
•   success(rc1_4_1s) && describe(m1_4_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gr
n updated.

```

Note

This is a good point to take a quick look at Pluto cell metadata: the top left eye symbol and the top right 3-dots in a circle glyph (both only visible when the curser is in the input cell). Both are used quite often in these notebooks. Try them out!

The output of above method of the function `model_summary(::SampleModel)`, called directly on a `SampleModel`, is different from method `model_summary(::DataFrame)`, typically used later on. Above table shows important mcmc diagnostic columns like `n_eff` and `r_hat`.

If Stan parameters are vectors (as in this example), `cmdstan` returns those using ' notation, e.g. `a.1`, `a.2`, ...

	parameters	median	mad_sd	mean	std
1	"a.1"	9.387	1.396	9.368	1.39
2	"a.2"	16.15	0.652	16.133	0.69
3	"b.1"	3.255	0.438	3.252	0.43
4	"b.2"	6.969	1.04	6.972	1.04
5	"sigma.1"	3.43	0.336	3.461	0.34
6	"sigma.2"	3.658	0.366	3.676	0.37

```

• if success(rc1_4_1s)
•   post1_4_1s = read_samples(m1_4_1s,
•   :dataframe)
•   model_summary(post1_4_1s,
•   names(post1_4_1s))
end

```

With vector parameters `read_samples()` can create a nested DataFrame:

`nd1_4_1s =`

	a	b
1	▶ [10.3282, 14.1843]	▶ [2.75128, 8.3895]
2	▶ [10.7021, 16.067]	▶ [2.61631, 8.9163]
3	▶ [6.66245, 17.2466]	▶ [4.09877, 7.9034]
4	▶ [8.92833, 16.0098]	▶ [3.59349, 7.9635]
5	▶ [8.26656, 16.2145]	▶ [3.5647, 7.7423]
6	▶ [7.95563, 16.0481]	▶ [3.72812, 7.0411]
7	▶ [8.45144, 16.6981]	▶ [3.47334, 6.9575]
8	▶ [11.4261, 14.7333]	▶ [2.83714, 7.5298]
9	▶ [10.6172, 15.1755]	▶ [2.957, 7.7749]
10	▶ [8.84695, 15.9684]	▶ [3.54249, 7.9980]
⋮	more	
4000	▶ [10.6575, 17.1363]	▶ [2.87471, 6.0266]

```
nd1_4_1s = read_samples(m1_4_1s,
                        :nesteddataframe)
```

`ms1_4_1s =`

	parameters	median	mad_sd	mean	std
1	"a.1"	9.387	1.396	9.368	1.39
2	"a.2"	16.15	0.652	16.133	0.69
3	"b.1"	3.255	0.438	3.252	0.43
4	"b.2"	6.969	1.04	6.972	1.04
5	"sigma.1"	3.43	0.336	3.461	0.34
6	"sigma.2"	3.658	0.366	3.676	0.37

```
• ms1_4_1s = success(rc1_4_1s) &&
  model_summary(post1_4_1s,
    names(post1_4_1s))
```

1.04

```
• ms1_4_1s["b.2", "mad_sd"]
```

Nested dataframes are handy to obtain a matrix of say the b values:

4000×2 Matrix{Float64}:

```
2.75128  8.38958
2.61631  8.9163
4.09877  7.90345
3.59349  7.96354
3.5647   7.7423
3.72812  7.04112
3.47334  6.95756
⋮
2.79654  4.78503
2.75535  9.23017
3.33143  7.67702
2.83813  6.82789
2.86538  4.61493
2.87471  6.0266
```

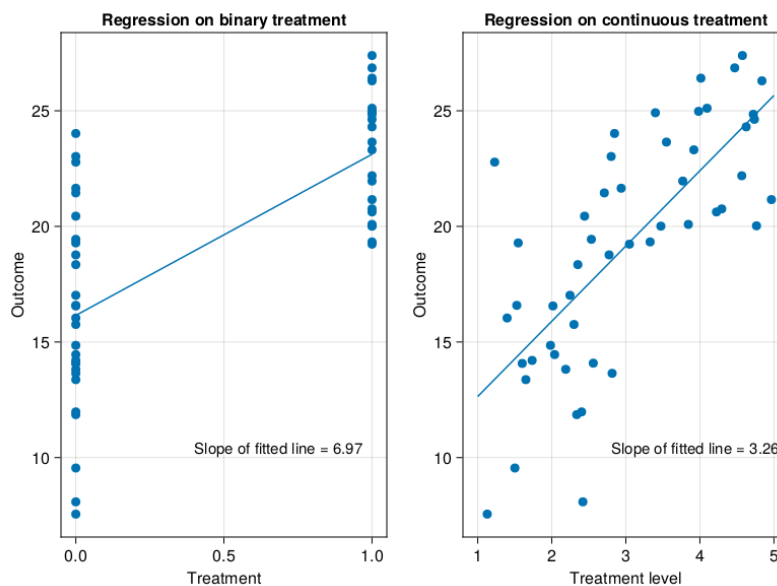
```
• array(nd1_4_1s, :b)
```



```
4000×2 Matrix{Float64}:
```

```
2.75128  8.38958
2.61631  8.9163
4.09877  7.90345
3.59349  7.96354
3.5647   7.7423
3.72812  7.04112
3.47334  6.95756
⋮
2.79654  4.78503
2.75535  9.23017
3.33143  7.67702
2.83813  6.82789
2.86538  4.61493
2.87471  6.0266
```

```
• Array(post1_4_1s[:, ["b.1", "b.2"]])
```



```

• let
•     x1 = 1.0:0.01:5.0
•     f = Figure()
•     medians = ms1_4_1s[:, "median"]
•     ax = Axis(f[1, 2], title = "Regression
•     on continuous treatment",
•         xlabel = "Treatment level", ylabel
•         = "Outcome")
•     sca1 = scatter!(x, y)
•     annotations!("Slope of fitted line =
•     $(round(medians[3], digits=2))",
•         position = (2.8, 10), textsize=15)
•     lin1 = lines!(x1, medians[1] .+
•     medians[3] * x1)
•
•     x2 = 0.0:0.01:1.0
•     ax = Axis(f[1, 1], title="Regression
•     on binary treatment",
•         xlabel = "Treatment", ylabel =
•         "Outcome")
•     sca1 = scatter!(x_binary, y)
•     lin1 = lines!(x2, medians[2] .+
•     medians[4] * x2)
•     annotations!("Slope of fitted line =
•     $(round(medians[4], digits=2))",
•         position = (0.4, 10), textsize=15)
•     f
•
end

```

```
• stan1_4_2 = "  
• data {  
•   int N;  
•   vector[N] x;  
•   vector[N] y;  
• }  
• parameters {  
•   vector[2] a;  
•   real b;  
•   real b_exp;  
•   vector<lower=0>[2] sigma;  
• }  
• model {  
•   // Priors  
•   a ~ normal(10, 5);  
•   b ~ normal(0, 5);  
•   b_exp ~ normal(5, 5);  
•   sigma ~ exponential(1);  
•   // Likelihood  
•   vector[N] mu;  
•   for ( i in 1:N )  
•     mu[i] = a[2] + b_exp * exp(-x[i]);  
•   y ~ normal(mu, sigma[2]);  
•   y ~ normal(a[1] + b * x, sigma[1]);  
• }  
• ";
```

	parameters	mean	mcse	std	5'
1	"a[1]"	15.0	0.02	0.96	14.0
2	"a[2]"	5.7	0.0078	0.41	5.1
3	"b"	-2.5	0.0	0.3	-3.0
4	"b_exp"	24.2	0.1	3.1	19.0
5	"sigma[1]"	2.4	0.0041	0.24	2.0
6	"sigma[2]"	2.2	0.0047	0.26	1.8

```

• let
•     #Random.seed!(1533)
•     n1 = 50
•     x1 = rand(Uniform(1, 5), n1)
•     y1 = [rand(Normal(5 + 30exp(-x1[i]),
•     2), 1)[1] for i in 1:n]
•     data = (N = n1, x = x1, y = y1)
•     global m1_4_2s = SampleModel("m1.4_2s",
•     stan1_4_2);
•     global rc1_4_2s = stan_sample(m1_4_2s;
•     data)
•     success(rc1_4_2s) && describe(m1_4_2s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gr
n updated.

```

	parameters	median	mad_sd	mean	std
1	"a.1"	15.488	0.942	15.48	0.96
2	"a.2"	5.703	0.407	5.722	0.40
3	"b"	-2.492	0.285	-2.489	0.29
4	"b_exp"	24.354	3.014	24.248	3.05
5	"sigma.1"	2.379	0.231	2.399	0.24
6	"sigma.2"	2.209	0.253	2.231	0.25

```
• if success(rc1_4_2s)
•     post1_4_2s = read_samples(m1_4_2s,
•         :dataframe)
•     ms1_4_2s = model_summary(post1_4_2s,
•         ["a.1", "a.2", "b", "b_exp",
•         "sigma.1", "sigma.2"])
end
```

nd1_4_2s =

	b	b_exp	a
1	-2.48576	20.0849	▶ [15.9861, 6.25252]
2	-2.48382	27.7827	▶ [14.8627, 5.18205]
3	-2.51677	27.8339	▶ [16.1262, 5.18631]
4	-2.63137	27.8789	▶ [15.4167, 5.25809]
5	-2.52726	28.7973	▶ [16.002, 4.68289]
6	-2.63422	29.3018	▶ [15.6185, 5.16142]
7	-2.31383	22.5431	▶ [14.7362, 5.84994]
8	-2.10616	29.9681	▶ [14.4546, 5.23276]
9	-2.63249	29.1633	▶ [15.8549, 5.28732]
10	-2.73308	27.838	▶ [16.2086, 5.07141]
⋮	more		
4000	-2.31223	25.1546	▶ [14.6061, 5.48579]

```
• nd1_4_2s = read_samples(m1_4_2s,
    :nesteddataframe)
```

4000×2 Matrix{Float64}:

```
15.9861  6.25252
14.8627  5.18205
16.1262  5.18631
15.4167  5.25809
16.002   4.68289
15.6185  5.16142
14.7362  5.84994
⋮
14.423   6.11107
16.5799  6.11235
14.7391  5.71911
14.2216  6.1607
16.0556  5.70143
14.6061  5.48579
```

```
• array(nd1_4_2s, :a)
```

```

•  $\hat{a}_1, \hat{a}_2, \hat{b}, \hat{b}_{exp}, \hat{\sigma}_1, \hat{\sigma}_2 = [\text{ms1\_4\_2s}[p,$ 
  "median"] for p in ["a.1", "a.2", "b",
    "b_exp", "sigma.1", "sigma.2"]];


```

5.703

```

•  $\hat{a}_2$ 

```



```

• let
•     x1 = LinRange(1, 6, 50)
•     y1 = [rand(Normal(5 + 30exp(-x1[i]),
•         2), 1)[1] for i in 1:length(x1)]
•
•     f = Figure()
•     ax = Axis(f[1, 1], title = "Linear
•         regression",
•         xlabel = "Treatments", ylabel =
•             "Outcomes")
•     scatter!(x1, y1)
•     lines!(x1,  $\hat{a}_1$  .+  $\hat{b}$  .* x1)
•
•     ax = Axis(f[2, 1], title = "Non-linear
•         regression",
•         xlabel = "Treatments", ylabel =
•             "Outcomes")
•     scatter!(x1, y1)
•     lines!(x1,  $\hat{a}_2$  .+  $\hat{b}_{exp}$  .* exp.(-x1))
•     f
• end

```

	xx	z	yy
1	3.1425	0	37.5294
2	0.0335661	1	30.0628
3	1.60408	0	28.1095
4	0.478735	1	31.3638
5	2.65874	0	35.7382
6	1.02705	1	36.0009
7	1.28799	0	24.3213
8	0.052966	1	29.5242
9	0.543994	0	25.4181
10	0.0304007	1	25.1863
⋮	more		
100	0.00156342	1	28.8751

```


• begin
•   Random.seed!(12573)
•   n2 = 100
•   z = repeat([0, 1]; outer=50)
•   df1_8 = DataFrame()
•   df1_8.xx = [(z[i] == 0 ? rand(Normal(0,
•   1.2), 1).^2 : rand(Normal(0, 0.8),
•   1).^2)[1] for i in 1:n2]
•   df1_8.z = z
•   df1_8.yy = [rand(Normal(20 .+
•   5df1_8.xx[i] .+ 10df1_8.z[i], 3), 1)
•   [1] for i in 1:n2]
•   df1_8
• end

```


	Coef.	Std. Error	t	Pr(> t)
(Intercept)	20.1093	0.529823	37.95	<1e-06
xx	4.97503	0.213492	23.30	<1e-06
z	9.625	0.604978	15.91	<1e-06

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	20.0337	0.544062	36.82	<1e-16
xx	5.01957	0.226965	22.12	<1e-16

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	29.8841	0.49051	60.92	<1e-16
xx	4.66553	0.609796	7.65	<1e-16



```

let
    â₁, b̂₁ = coef(lm1_8_0)
    â₂, b̂₂ = coef(lm1_8_1)
    x = LinRange(0, maximum(df1_8.xx), 40)

    f = Figure()
    ax = Axis(f[1, 1]; title="Figure 1.8")
    scatter!(df1_8.xx[df1_8.z .== 0],
             df1_8.yy[df1_8.z .== 0])
    scatter!(df1_8.xx[df1_8.z .== 1],
             df1_8.yy[df1_8.z .== 1])
    lines!(x, â₁ .+ b̂₁ * x, label =
            "Control")
    lines!(x, â₂ .+ b̂₂ * x, label =
            "Treated")
    axislegend(; position=(right,
                           :bottom))
    current_figure()
end

```

1.5 Classical and Bayesian inference.

1.6 Computing least-squares and Bayesian regression.

1.8 Exercises.

Helicopters

`helicopters =`

	Helicopter_ID	width_cm	length_cm	time_sec
1	1	4.6	8.2	1.64
2	1	4.6	8.2	1.74
3	1	4.6	8.2	1.68
4	1	4.6	8.2	1.62
5	1	4.6	8.2	1.68
6	1	4.6	8.2	1.7
7	1	4.6	8.2	1.62
8	1	4.6	8.2	1.66
9	1	4.6	8.2	1.69
10	1	4.6	8.2	1.62
	⋮ more			
20	2	4.6	8.2	1.61

```
• helicopters =
  CSV.read(ros_datadir("Helicopters",
    "helicopters.csv"), DataFrame)
```

Simulate 40 helicopters.

	width_cm	length_cm	time_sec
1	10.0236	9.59097	1.78566
2	2.33684	7.81096	1.41044
3	7.68879	2.34587	0.794504
4	6.67829	13.6578	1.81387
5	8.79925	9.27474	1.57939
6	2.40055	7.55651	1.2363
7	5.9089	16.177	2.04483
8	5.12956	14.864	1.8629
9	3.6735	15.5807	1.7891
10	2.18695	7.53194	1.15288
⋮	more		
40	8.39115	5.10451	1.32113

```

• begin
•     helis = DataFrame(width_cm =
•         rand(Normal(5, 2), 40), length_cm =
•         rand(Normal(10, 4), 40))
•     helis.time_sec = 0.5 .+ 0.04 .*
•         helis.width_cm .+ 0.08 .*
•         helis.length_cm .+ 0.1 .*
•         rand(Normal(0, 1), 40)
•     helis
• end

```

```
stan1_5 = "  
data {  
  int N;  
  vector[N] w;  
  vector[N] l;  
  vector[N] y;  
}  
parameters {  
  real a;  
  real b;  
  real c;  
  real<lower=0> sigma;  
}  
model {  
  // Priors  
  a ~ normal(10, 5);  
  b ~ normal(0, 5);  
  sigma ~ exponential(1);  
  
  // Likelihood time on width  
  vector[N] mu;  
  for ( i in 1:N )  
    mu[i] = a + b * w[i] + c * l[i];  
  y ~ normal(mu, sigma);  
}  
";
```

	parameters	mean	mcse	std
1	"a"	0.438333	0.00132571	0.05575
2	"b"	0.0415044	0.000159679	0.00701
3	"c"	0.0835149	7.1496e-5	0.00351
4	"sigma"	0.0946768	0.00026859	0.01124

```

• let
•   data = (N = nrow(helis), y =
•   helis.time_sec, w = helis.width_cm, l
•   = helis.length_cm)
•   global m1_5s = SampleModel("m1.5s",
•   stan1_5);
•   global rc1_5s = stan_sample(m1_5s;
•   data)
•   success(rc1_5s) && describe(m1_5s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gr
updated.

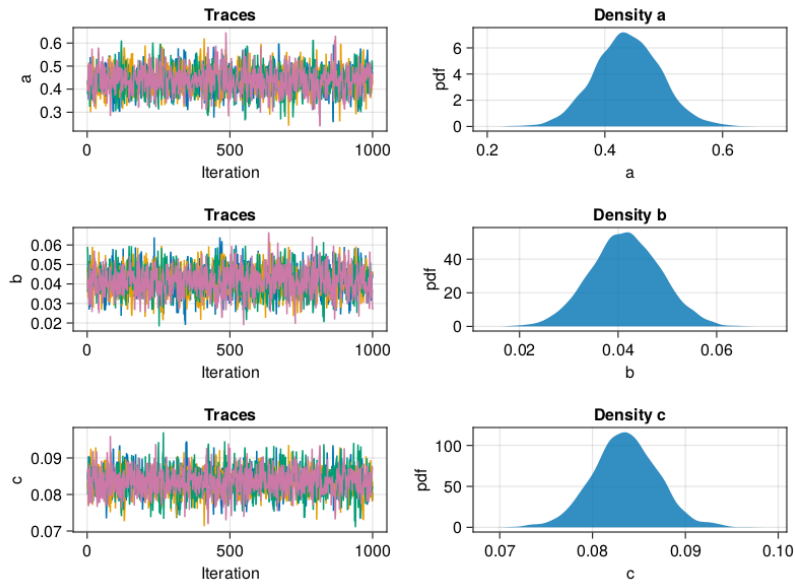
```

	parameters	median	mad_sd	mean	std
1	"a"	0.4378	0.0554	0.4383	0.05
2	"b"	0.0415	0.007	0.0415	0.00
3	"c"	0.0835	0.0034	0.0835	0.00
4	"sigma"	0.0936	0.0108	0.0947	0.01


```

• if success(rc1_5s)
•   post1_5s = read_samples(m1_5s,
•   :dataframe)
•   model_summary(post1_5s, [:a, :b, :c,
•   :sigma]; digits=4)
end


```



```
• plot_chains(post1_5s, [:a, :b, :c])
```



```
• plot_chains(post1_5s, [:sigma])
```



```

let
    w_range = LinRange(1.0, 8.0, 100)
    w_times = mean.(link(post1_5s, (r, w) -
    > r.a + r.c + r.b * w, w_range))
    l_range = LinRange(6.0, 15.0, 100)
    l_times = mean.(link(post1_5s, (r, l) -
    > r.a + r.b + r.c * l, l_range))

    f = Figure()
    ax = Axis(f[1, 1], title = "Time in
    the air on width and length",
    xlabel = "Width/Length", ylabel =
    "Time in the air")

    lines!(w_range, w_times; label="Width")
    lines!(l_range, l_times;
    label="Length")

    f[1, 2] = Legend(f, ax, "Regression
    lines", framevisible = false)

    current_figure()
end

```

```
lnk1_5s =
```

```
► [[0.894902, 0.886343, 0.876715, 0.872754, 0.84
```

```

    lnk1_5s = link(post1_5s, (r, l) -> r.a +
    r.b + r.c * l, [5, 10, 12])

```



```
▶ [0.897544, 1.31378, 1.48082]
```

```
• median.(lnk1_5s)
```

```
▶ [0.0388104, 0.0333603, 0.0332957]
```

```
• mad.(lnk1_5s)
```

```
▶ [0.897412, 1.31499, 1.48202]
```

```
• mean.(link(post1_5s, (r, l) -> r.a + r.b  
+ r.c * l, [5, 10, 12]))
```

	a	b	c	sigma
1	0.436716	0.0369489	0.0842475	0.08961
2	0.418862	0.0358816	0.0863198	0.10801
3	0.414789	0.051869	0.0820115	0.10764
4	0.390606	0.0437647	0.0876767	0.09686
5	0.35434	0.0420388	0.0905117	0.10267
6	0.404372	0.0515824	0.0810212	0.07660
7	0.512649	0.0268483	0.0845897	0.11502
8	0.46746	0.0448521	0.0787241	0.08613
9	0.448316	0.0290421	0.0883864	0.10002
10	0.397243	0.0549251	0.0808393	0.08290
⋮	more			
4000	0.465282	0.0268118	0.0873861	0.10298

```
• read_samples(m1_5s, :nesteddataframe)
```

```
No nested columns found.
```

