

See chapter 9 in Regression and Other Stories.

```

• html"""
• <style>
•   main {
•     margin: 0 auto;
•     max-width: 2000px;
•     padding-left: max(160px, 10%);
•     padding-right: max(160px, 10%);
•   }
• </style>
• """
•

```

```

• using Pkg ✓ , DrWatson ✓

```

```

• begin
•   using GLM ✓
•
•   # Specific to ROSStanPluto
•   using StanSample ✓
•
•   # Graphics related
•   using GLMakie ✓
•
•   # Common data files and functions
•   using RegressionAndOtherStories ✓
• end

```

```

Replacing docs for `RegressionAndOtherStories.
frames.DataFrame, AbstractString}` in module `R
`

```

9.1 Propagating uncertainty in inference using posterior simulations.

`hibbs =`

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
⋮	more			
16	2012	0.95	52.0	"Obama"

```
• hibbs =
  CSV.read(ros_datadir("ElectionsEconomy",
    "hibbs.csv"), DataFrame)
```



```
hibbs_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
vote ~ 1 + growth
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	46.2476	1.62193	28.51	<1e-16
growth	3.06053	0.696274	4.40	0.00011

```
• hibbs_lm = lm(@formula(vote ~ growth),  
hibbs)
```



```

• let
•     fig = Figure()
•     hibbs.label = string(hibbs.year)
•     xlabel = "Average growth personal
•     income [%]"
•     ylabel = "Incumbent's party vote share"
•     let
•         title = "Forecasting the election
•         from the economy"
•         ax = Axis(fig[1, 1]; title, xlabel,
•         ylabel)
•         for (ind, yr) in
•         enumerate(hibbs.year)
•             annotations!("$(yr)"; position=
•             (hibbs.growth[ind],
•             hibbs.vote[ind]), fontsize=10)
•         end
•     end
•     let
•         x = LinRange(-1, 4, 100)
•         title = "Data and linear fit"
•         ax = Axis(fig[1, 2]; title, xlabel,
•         ylabel)
•         scatter!(hibbs.growth, hibbs.vote)
•         lines!(x, coef(hibbs_lm)[1] .+
•         coef(hibbs_lm)[2] .* x;
•         color=:darkred)
•         annotations!("vote = 46.2 + 3.0 *
•         growth"; position=(0, 41))
•     end
•     fig
• end

```

```
• stan7_1 = "  
• data {  
•   int<lower=1> N;           // total number  
•   of observations  
•   vector[N] growth;       // Independent  
•   variable: growth  
•   vector[N] vote;         // Dependent  
•   variable: votes  
• }  
• parameters {  
•   real b;                  // Coefficient  
•   independent variable  
•   real a;                  // Intercept  
•   real<lower=0> sigma;     // dispersion  
•   parameter  
• }  
• model {  
•   vector[N] mu;  
•  
•   // priors including constants  
•   a ~ normal(50, 20);  
•   b ~ normal(2, 10);  
•   sigma ~ exponential(1);  
•  
•   mu = a + b * growth;  
  
•   // likelihood including constants  
•   vote ~ normal(mu, sigma);  
• }";
```

	parameters	mean	mcse	std	
1	"b"	3.04897	0.0162602	0.654979	1
2	"a"	46.2878	0.0365967	1.5194	4
3	"sigma"	3.57757	0.0159071	0.613961	2

```

• let
•     data = (N=nrow(hibbs), vote=hibbs.vote,
•           growth=hibbs.growth)
•     global m7_1s = SampleModel("hibbs",
•           stan7_1)
•     global rc7_1s = stan_sample(m7_1s;
•           data)
•           success(rc7_1s) && describe(m7_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gr
updated.

```

1489.69

```

• let
•     ss7_1s = describe(m7_1s)
•     ss7_1s[:sigma, :ess]
• end

```

1489.69

```

• let
•     ss7_1s = describe(m7_1s; showall=true)
•     ss7_1s[:sigma, :ess]
• end

```


	parameters	mean	mcse	
1	"lp__"	-31.9128	0.0383475	1.3
2	"accept_stat__"	0.917887	0.00977819	0.1
3	"stepsize__"	0.404689	0.0258293	0.0
4	"treedepth__"	2.666	0.0150726	0.0
5	"n_leapfrog__"	7.88	0.205667	4.1
6	"divergent__"	0.0	NaN	0.0
7	"energy__"	33.3892	0.0489152	1.7
8	"b"	3.04897	0.0162602	0.0
9	"a"	46.2878	0.0365967	1.5
10	"sigma"	3.57757	0.0159071	0.0

• `describe(m7_1s; showall=true)`

`post7_1s =`

	<code>b</code>	<code>a</code>	<code>sigma</code>
1	1.90114	48.6989	2.77909
2	3.10169	47.0837	3.10472
3	4.40427	44.2424	4.01044
4	4.60825	45.0937	4.07456
5	1.63412	47.5276	3.7154
6	1.74813	47.3943	3.79968
7	3.18921	45.6586	3.21523
8	2.71001	47.9793	4.27731
9	2.62185	47.9963	4.25759
10	2.98121	46.9598	4.32174
⋮ more			
4000	2.04827	49.2821	3.21009

- `post7_1s = success(rc7_1s) &&
read_samples(m7_1s, :dataframe)`



- `trankplot(post7_1s, "b")`

`ms7_1s =`

	parameters	median	mad_sd	mean	std
1	"a"	46.303	1.504	46.288	1.51
2	"b"	3.036	0.622	3.049	0.65
3	"sigma"	3.508	0.577	3.578	0.61

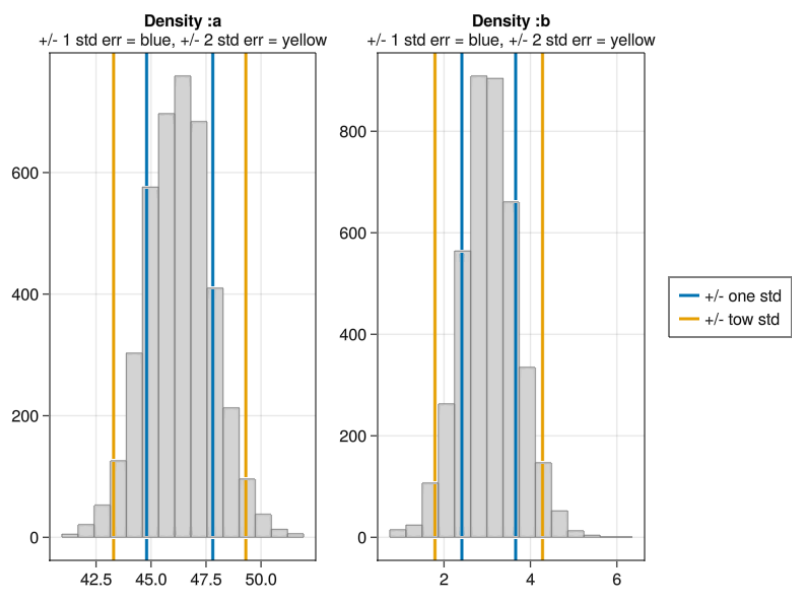
- `ms7_1s = model_summary(post7_1s, [:a, :b, :sigma])`

```
sims = 4000×3 Matrix{Float64}:
 1.90114  48.6989  2.77909
 3.10169  47.0837  3.10472
 4.40427  44.2424  4.01044
 4.60825  45.0937  4.07456
 1.63412  47.5276  3.7154
 1.74813  47.3943  3.79968
 3.18921  45.6586  3.21523
 ⋮
 3.0609   46.2324  3.51815
 3.5505   45.7074  3.81199
 2.95468  45.8916  3.8862
 2.86051  46.8871  5.21757
 1.58464  48.6092  3.16146
 2.04827  49.2821  3.21009
```

- `sims = Array(post7_1s)`

```
1×3 Matrix{Float64}:
 3.03601  46.3033  3.50782
```

- `median(sims; dims=1)`





```

• let
•     f = Figure()
•     ax = Axis(f[1, 1]; title="Density :a",
•     subtitle="+/- 1 std err = blue, +/- 2
•     std err = yellow")
•     hist!(post7_1s.a; bins=15, color =
•     :lightgrey, strokewidth = 1,
•     strokecolor = :grey)
•     one = vlines!([ms7_1s[:a, :median] -
•     ms7_1s[:a, :mad_sd], ms7_1s[:a,
•     :median] + ms7_1s[:a, :mad_sd]]);
•     linewidth=3)
•     two = vlines!([ms7_1s[:a, :median] -
•     2ms7_1s[:a, :mad_sd], ms7_1s[:a,
•     :median] + 2ms7_1s[:a, :mad_sd]]);
•     linewidth=3)

•     ax = Axis(f[1, 2]; title="Density :b",
•     subtitle="+/- 1 std err = blue, +/- 2
•     std err = yellow")
•     hist!(post7_1s.b; bins=15, color =
•     :lightgrey, strokewidth = 1,
•     strokecolor = :grey)
•     vlines!([ms7_1s[:b, :median] -
•     ms7_1s[:b, :mad_sd], ms7_1s[:b,
•     :median] + ms7_1s[:b, :mad_sd]]);
•     linewidth=3)
•     vlines!([ms7_1s[:b, :median] -
•     2ms7_1s[:b, :mad_sd], ms7_1s[:b,
•     :median] + 2ms7_1s[:b, :mad_sd]]);
•     linewidth=3)
•     Legend(f[1, 3], [one, two], ["+/- one
•     std", "+/- tow std"])
•     f
• end

```



```

• let
•     growth_range =
•     LinRange(minimum(hibbs.growth),
•     maximum(hibbs.growth), 200)
•     votes = mean.(link(post7_1s, (r,x) ->
•     r.a + x * r.b, growth_range))
•
•     xlabel = "Average growth personal
•     income [%]"
•     ylabel = "Incumbent's party vote share"
•
•     fig = Figure()
•
•     ax = Axis(fig[1, 1]; title="Plot the
•     mcmc draws for :a and :b",
•     xlabel=":a", ylabel=":b")
•     scatter!(post7_1s.a, post7_1s.b;
•     markersize=4)
•
•     xlabel = "Average growth personal
•     income [%]"
•     ylabel="Incumbent's party vote share"
•     ax = Axis(fig[1, 2]; title="Regression
•     line based on 4000 posterior samples",
•     subtitle = "(grey lines are based
•     on first 200 draws of :a and :b)",
•     xlabel, ylabel)
•     for i in 1:100
•         lines!(growth_range, post7_1s.a[i]
•         .+ post7_1s.b[i] .* growth_range,
•         color = :lightgrey)
•     end
•     scatter!(hibbs.growth, hibbs.vote)
•     lines!(growth_range, votes, color =
•     :red)
•     fig
• end

```

9.2 Prediction and uncertainty.

	x	y
1	-2.0	50
2	-1.0	44
3	0.0	50
4	1.0	47
5	2.0	56

```
• let
•   x = LinRange(-2, 2, 5)
•   y = [50, 44, 50, 47, 56]
•   global sexratio = DataFrame(x = x, y =
•   y)
• end
```

```
• stan9_1 = "  
• data {  
•   int<lower=1> N; // total number of  
•   observations  
•   vector[N] x;    // Independent  
•   variable: growth  
•   vector[N] y;    // Dependent variable:  
•   votes  
• }  
• parameters {  
•   real b;          // Coefficient  
•   independent variable  
•   real a;          // Intercept  
•   real<lower=0> sigma; // dispersion  
•   parameter  
• }  
• model {  
•   vector[N] mu;  
•  
•   // priors including constants  
•   a ~ normal(50, 5);  
•   b ~ normal(0, 5);  
•   sigma ~ uniform(0, 10);  
•  
•   mu = a + b * x;  
  
•   // likelihood including constants  
•   y ~ normal(mu, sigma);  
• }";
```

	parameters	mean	mcse	std	
1	"b"	1.25238	0.0676678	1.74543	-1
2	"a"	49.4801	0.0416165	2.41213	45
3	"sigma"	5.72269	0.0849459	1.88623	3.

```

let
  data = (N = nrow(sexratio), x =
sexratio.x, y = sexratio.y)
  global m9_1s = SampleModel("m9_1s",
stan9_1)
  global rc9_1s = stan_sample(m9_1s;
data)
  success(rc9_1s) && describe(m9_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gr
updated.



```

	parameters	median	mad_sd	mean	std
1	"a"	49.474	2.147	49.48	2.41
2	"b"	1.304	1.572	1.252	1.74
3	"sigma"	5.514	2.058	5.723	1.88

```

if success(rc9_1s)
  post9_1s = read_samples(m9_1s,
:dataframe)
  sm9_1s = model_summary(post9_1s, [:a,
:b, :sigma])
end

```

```

• let
•     x_range = LinRange(minimum(sexratio.x),
•     maximum(sexratio.x), 200)
•     y = mean.(link(post9_1s, (r,x) -> r.a
•     + x * r.b, x_range))
•
•     xlabel = "x"
•     ylabel = "y"
•
•     fig = Figure()
•
•     ax = Axis(fig[1, 1]; title="Posterior
•     simulation under default prior",
•     xlabel="Intercept, a", ylabel="Slope,
•     b")
•     scatter!(post9_1s.a, post9_1s.b;
•     markersize=4)
•
•     ax = Axis(fig[1, 2]; title="Bayes
•     regression (4000 posterior samples)",
•     subtitle = "(grey lines are based
•     on first 100 draws of a and b)",
•     xlabel, ylabel)
•     for i in 1:100
•         lines!(x_range, post9_1s.a[i] .+
•         post9_1s.b[i] .* x_range, color =
•         :lightgrey)
•     end
•     scatter!(sexratio.x, sexratio.y)
•     lines!(x_range, y, color = :red)
•     fig
• end

```

```
• stan9_2 = "  
• data {  
•   int<lower=1> N; // total number of  
•   observations  
•   vector[N] x;    // Independent  
•   variable: growth  
•   vector[N] y;    // Dependent variable:  
•   votes  
• }  
• parameters {  
•   real b;          // Coefficient  
•   independent variable  
•   real a;          // Intercept  
•   real<lower=0> sigma; // dispersion  
•   parameter  
• }  
• model {  
•   vector[N] mu;  
•  
•   // priors including constants  
•   a ~ normal(48.8, 0.2);  
•   b ~ normal(0, 0.2);  
•   sigma ~ uniform(0, 10);  
•  
•   mu = a + b * x;  
  
•   // likelihood including constants  
•   y ~ normal(mu, sigma);  
• }";
```

	parameters	mean	mcse	std
1	"b"	0.0308248	0.0034984	0.19705
2	"a"	48.8087	0.00329835	0.19532
3	"sigma"	5.06993	0.0303117	1.66184

```

let
  data = (N = nrow(sexratio), x =
sexratio.x, y = sexratio.y)
  global m9_2s = SampleModel("m9_2s",
stan9_2)
  global rc9_2s = stan_sample(m9_2s;
data)
  success(rc9_2s) && describe(m9_2s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gr
updated.



```

	parameters	median	mad_sd	mean	std
1	"a"	48.809	0.195	48.809	0.19
2	"b"	0.03	0.194	0.031	0.19
3	"sigma"	4.758	1.573	5.07	1.66

```

if success(rc9_2s)
  post9_2s = read_samples(m9_2s,
:dataframe)
  sm9_2s = model_summary(post9_2s, [:a,
:b, :sigma])
end

```



```

• let
•     x_range = LinRange(minimum(sexratio.x),
•     maximum(sexratio.x), 200)
•     y = mean.(link(post9_2s, (r,x) -> r.a
•     + x * r.b, x_range))
•
•     xlabel = "x"
•     ylabel = "y"
•
•     fig = Figure()
•
•     ax = Axis(fig[1, 1]; title="Posterior
•     simulation under informative prior",
•     xlabel="Intercept, a", ylabel="Slope,
•     b")
•     ylims!(ax, -8, 8)
•     scatter!(post9_2s.a, post9_2s.b;
•     markersize=4)
•
•     ax = Axis(fig[1, 2]; title="Bayes
•     regression (4000 posterior samples)",
•     subtitle = "(grey lines are based
•     on first 100 draws of :a and :b)",
•     xlabel, ylabel)
•     for i in 1:100
•         lines!(x_range, post9_2s.a[i] .+
•         post9_2s.b[i] .* x_range, color =
•         :lightgrey)
•     end
•     scatter!(sexratio.x, sexratio.y)
•     lines!(x_range, y, color = :red)
•     fig
• end

```

9.3 Prior information and Bayesian synthesis.

Prior based on a previously-fitted model using economic and political condition.

```

• begin
•   theta_hat_prior = 0.524
•   se_prior = 0.041
• end;

```

Survey of 400 people, of whom 190 say they will vote for the Democratic candidate.

```

• begin
•   n = 400
•   y = 190
• end;

```

Data estimate.

```
theta_hat_data = 0.475
```

```
• theta_hat_data = y/n
```

```
se_data = 0.02496873044429772
```

```
• se_data = sqrt((y/n)*(1-y/n)/n)
```

Bayes estimate.

```
theta_hat_bayes = 0.48825635323153693
```

```

• theta_hat_bayes =
•   (theta_hat_prior/se_prior^2 +
•    theta_hat_data/se_data^2)
•   /(1/se_prior^2 + 1/se_data^2)


```

```
se_bayes = 0.02132543263776925
```

```

• se_bayes = sqrt(1/(1/se_prior^2 +
•   1/se_data^2))

```




```

• let
•     x = 0.3:0.001:0.7
•     f = Figure()
•     ax = Axis(f[1, 1], title="Prior,
•     likelihood & posterior")
•     prior = lines!(f[1, 1], x, pdf.
•     (Normal(theta_hat_prior, se_prior),
•     x), color=:gray)
•     data = lines!(x, pdf.
•     (Normal(theta_hat_data, se_data),
•     x), color=:darkred)
•     bayes = lines!(x, pdf.
•     (Normal(theta_hat_bayes, se_bayes),
•     x), color=:darkblue)
•     Legend(f[1, 2], [prior, data, bayes],
•     ["Prior", "Data", "Bayesian"])

•     current_figure()

end

```

```

• let
•     f = Figure()
•     ax = Axis(f[1, 1], title="Prior,
•         likelihood & posterior (using
•         `density()`)" )
•     density!(rand(Normal(theta_hat_prior,
•         se_prior), Int(1e6)), lab="prior")
•     density!(rand(Normal(theta_hat_data,
•         se_data), Int(1e6)), lab="likelihood")
•     density!(rand(Normal(theta_hat_bayes,
•         se_bayes), Int(1e6)), lab="bayes")
•     current_figure()
• end

```

9.4 Example of Bayesian inference: beauty and sex ratio.

