

# See chapter 3 in Regression and Other Stories.

---

## Widen the notebook

```
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

```
• using Pkg ✓, DrWatson ✓
```

```
• begin  
•     # Specific to this notebook  
•     using GLM ✓  
•  
•     # Graphics related  
•     using GLMakie ✓  
•  
•     # Specific to ROSStanPluto  
•     using StanSample ✓  
•  
•     # Common data files and functions  
•     using RegressionAndOtherStories ✓  
• end
```

## 3.1 - Weighted averages

pop =

	stratum	country	population	average
1	1	"United States"	310000000	36.8
2	2	"Mexico"	112000000	26.7
3	3	"Canada"	34000000	40.7

```
• pop = DataFrame(stratum=1:3, country=
  • ["United States", "Mexico", "Canada"],
    population=Int[310e6, 112e6, 34e6],
    average_age=[36.8, 26.7, 40.7])
```

34.61008771929824

```
• mean(pop.average_age,
  weights(pop.population))
```


► [0.679825, 0.245614, 0.0745614]

```
• weights(pop.population)/sum(pop.population
  )
```

	variable	mean	min	median
1	:stratum	2.0	1	2.0
2	:country	nothing	"Canada"	nothing
3	:population	1.52e8	34000000	1.12e8
4	:average_age	34.7333	26.7	36.8

```
• describe(pop)
```


## 3.3 - Graphing a line



```

• let
•     data = LinRange(0.01, 0.99, 200)
•     f = Figure(resolution = (800, 800))
•
•     for (i, scale) in enumerate([identity,
•     log10, log2, log, sqrt, Makie.logit])
•
•         row, col = fldmod1(i, 2)
•         Axis(f[row, col], yscale = scale,
•         title = string(scale),
•             yminorticksvisible = true,
•             yminorgridvisible = true,
•             yminorticks =
•             IntervalsBetween(8))
•
•         lines!(data, color = :blue)
•     end
•
•     f
• end

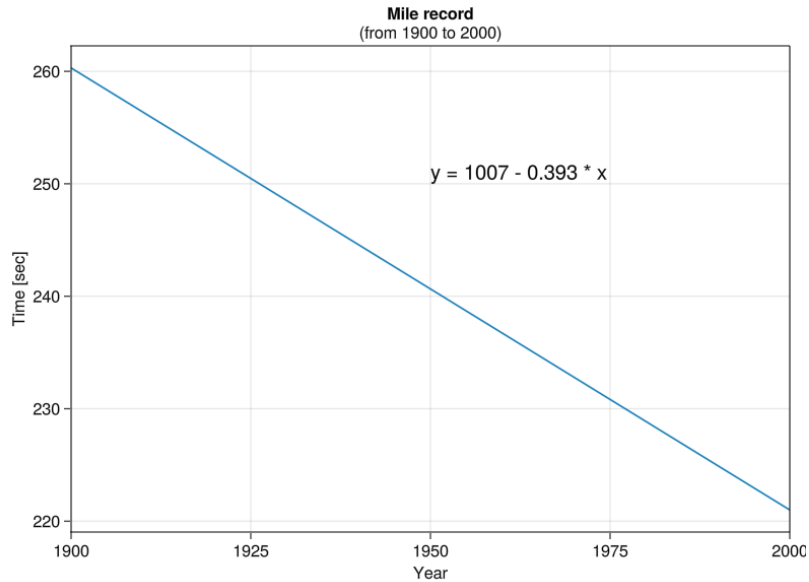
```



```

• let
•     data = 10 .^ LinRange(0.01, 5.0, 200)
•     f = Figure(resolution = (800, 300))
•
•     for (i, scale) in enumerate([log10,
•     log2, log])
•         row, col = fldmod1(i, 2)
•         Axis(f[1, i], yscale = scale,
•         xscale = scale, title = string(scale),
•         yminorticksvisible = true,
•         yminorgridvisible = true,
•         yminorticks =
•         IntervalsBetween(8))
•
•         lines!(data, color = :blue)
•     end
•     f
• end

```



```

let
  f = Figure()
  ax = Axis(f[1, 1]; title = "Mile
  record", subtitle = "(from 1900 to
  2000)", xlabel = "Year", ylabel =
  "Time [sec]")
  xlims!(ax, 1900, 2000)
  ax.xticks = 1900:25:2000
  x = 1900:2000
  y = 1007 .- 0.393 .* x
  lines!(x, y)
  annotations!("y = 1007 - 0.393 * x",
  position=(1950, 250))
  current_figure()
end

```

## 3.4 - Log and exponential scales

Simulated data for metabolic.

	body_mass	rate
1	2.98408	3.76147
2	4.08312	4.44374
3	4.16122	4.48304
4	4.47076	4.77568
5	4.60702	4.89275
6	4.72867	4.89577
7	5.41953	5.48691
8	5.85765	5.76041
9	5.93956	5.77697
10	5.9717	5.88113
⋮	more	
200	9.1991	8.22749

```

• begin
•     x = sort(rand(Uniform(0.01, 10000),
•     200))
•     y = 4.1 * x.^0.74 .+ [rand.(Normal.(0,
•     sqrt(x[i])), 1)[1] for i in
•     1:length(x)]
•     metabolic = DataFrame(:body_mass =>
•     log.(x), :rate => log.(y))
end

```

```

• stan3_1 = "
• data {
•     int N;
•     vector[N] m;
•     vector[N] r;
• }
• parameters {
•     real a;
•     real b;
•     real sigma;
• }
• model {
•     vector[N] mu;
•     a ~ normal(0, 0.3);
•     b ~ normal(0, 0.3);
•     sigma ~ exponential(1);
•     mu = a + b * m;
•     r ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std
1	"a"	1.45414	0.000701633	0.02043
2	"b"	0.734723	8.56162e-5	0.00249
3	"sigma"	0.0384015	4.3462e-5	0.00189

```

• let
•     data = (N =
•         length(metabolic.body_mass), m =
•         metabolic.body_mass, r =
•         metabolic.rate)
•     global m3_1s = SampleModel("m3.1s",
•         stan3\_1)
•     global rc3_1s = stan_sample(m3_1s;
•         data)
•     success(rc3_1s) && describe(m3_1s)
• end

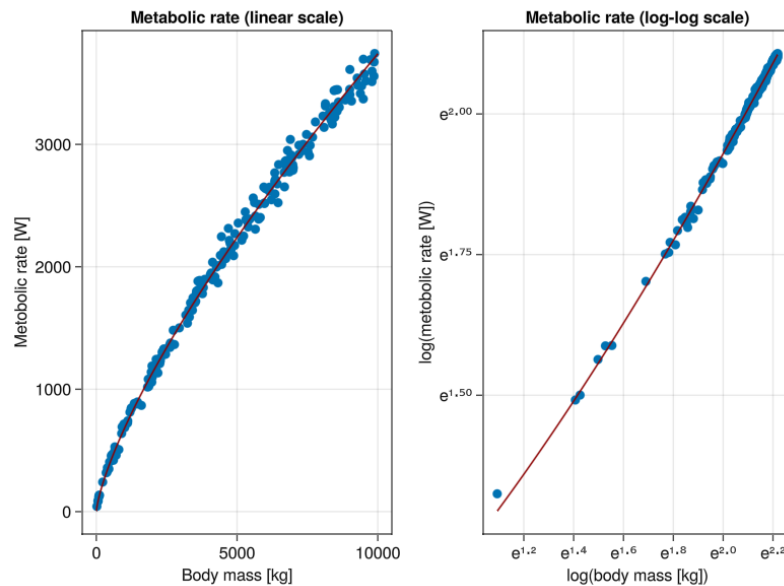
```

	parameters	median	mad_sd	mean	std
1	"a"	1.455	0.02	1.454	0.02
2	"b"	0.735	0.002	0.735	0.00
3	"sigma"	0.038	0.002	0.038	0.00

```

• if success(rc3_1s)
•     post3_1s = read_samples(m3_1s,
•         :dataframe)
•     ms3_1s = model_summary(post3_1s, [:a,
•         :b, :sigma])
end

```





```

• let
•     x = LinRange(1, 10000, 1000)
•     y = 4.1 * x.^0.74
•     f = Figure()
•     ax = Axis(f[1, 1]; title="Metabolic
•         rate (linear scale)", xlabel="Body
•         mass [kg]", ylabel="Metobolic rate
•         [W]")
•     scatter!(exp.(metabolic.body_mass),
•              exp.(metabolic.rate))
•     lines!(x, y; color=:darkred)
•
•     ax = Axis(f[1, 2]; title="Metabolic
•         rate (log-log scale)", xscale=log,
•         yscale=log,
•             xlabel="log(body mass [kg])",
•             ylabel="log(metobolic rate [W])")
•     x =
•         LinRange(minimum(metabolic.body_mass),
•                 maximum(metabolic.body_mass), 100)
•     scatter!(metabolic.body_mass,
•              metabolic.rate)
•     lines!(x, ms3_1s[:a, :mean] .+
•             ms3_1s[:b, :mean] * x; color=:darkred)
•     current_figure()
• end

```

```

100-element LinRange{Float64, Int64}:
 2.98408, 3.04686, 3.10964, 3.17241, 3.23519, ..

```

```

• LinRange(minimum(metabolic.body_mass),
            maximum(metabolic.body_mass), 100)

```

```
57.69669926958961
```

```
• exp(exp(1.4))
```

## 3.5 - Probability distributions

	height	sex
1	61.0116	"female"
2	63.1106	"female"
3	60.6683	"female"
4	62.4183	"female"
5	65.9579	"female"
6	64.6202	"female"
7	66.4275	"female"
8	61.6656	"female"
9	62.6349	"female"
10	66.2756	"female"
⋮ more		
200000	65.4376	"male"


```

• begin
•     N = 100000
•     heights = DataFrame()
•     height = vcat(rand(Normal(63.7, 2.7),
•         N),
•         rand(Normal(69.1, 2.9), N))
•     sex = repeat(["female", "male"],
•         inner=N)
•     heights.height = height
•     heights.sex = sex
•     heights
• end

```

```
▶ (mean = 63.6951, var = 7.31388, std = 2.70442
```

```
• begin
•   menHeights = heights[heights.sex .==
•   "male", :height]
•   womenHeights = heights[heights.sex .==
•   "female", :height]
•   (mean=mean(womenHeights),
•   var=var(womenHeights),
•       std=std(womenHeights),
•       median=median(womenHeights),
•       mad_sd=mad(womenHeights))
• end
```



```
• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Density
•   heights")
•   density!(womenHeights; color=color
•   = (:lightgreen, 0.4),
•   label="Women")
•   density!(menHeights; color=color =
•   (:lightblue, 0.4), label="Men")
•   axislegend()
•   f
• end
```

0.49714220980937984

```

• begin
•   wdf = Normal(63.65, 2.68)
•   cdf(wdf, 63.65 + 0.67 * 2.68) -
•   cdf(wdf, 63.65 - 0.67 * 2.68)
• end

```

0.6826894921370859

```

• cdf(wdf, 63.65 + 2.68) - cdf(wdf, 63.65 -
  2.68)


```

0.9544997361036417

```

• cdf(wdf, 63.65+2*2.68) - cdf(wdf, 63.65-
  2*2.68)

```



```

• let
•   wdf = Normal(63.65, 2.68)
•   x = range(55.0, 72.5 ; length=100)
•   lines(x, pdf.(wdf, x); color=:darkblue)
•
•   x1 = range(63.65 - 3 * 2.68, 63.65 - 2
•   * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•   (wdf, x1);
•   color = (:blue, 0.25), label =
•   "Label")
•
•   x1 = range(63.65 + 2 * 2.68, 63.65 + 3
•   * 2.68; length=20)

```

```

x1 = range(63.65 - 2 * 2.68, 63.65 - 1 * 2.68; length=20)
band!(x1, fill(0, length(x1)), pdf.
(wdf, x1);
color = (:blue, 0.25), label =
"Label")

x1 = range(63.65 - 2 * 2.68, 63.65 - 1 * 2.68; length=20)
band!(x1, fill(0, length(x1)), pdf.
(wdf, x1);
color = (:blue, 0.45), label =
"Label")

x1 = range(63.65 + 1 * 2.68, 63.65 + 2 * 2.68; length=20)
band!(x1, fill(0, length(x1)), pdf.
(wdf, x1);
color = (:blue, 0.45), label =
"Label")

x1 = range(63.65 - 1 * 2.68, 63.65; length=20)
band!(x1, fill(0, length(x1)), pdf.
(wdf, x1);
color = (:blue, 0.55), label =
"Label")

x1 = range(63.65, 63.65 + 2.68; length=20)
band!(x1, fill(0, length(x1)), pdf.
(wdf, x1);
color = (:blue, 0.55), label =
"Label")

text!("68%", position = (63.65, 0.05),
align = (:center, :center),
textsize = 30)
text!("13.5%", position = (67.5, 0.02),
align = (:center, :center),
textsize = 20)
text!("13.5%", position = (59.6, 0.02),
align = (:center, :center),
textsize = 20)
text!("2.5%", position = (69.75,
0.0045), align = (:center, :center),
textsize = 15)

```

```

text!("2.5%", position = (57.7,
0.0045), align = (:center, :center),
    textsize = 15)
current_figure()
end

```

► ( $\hat{m} = 6.0422$ ,  $m = 6.0$ ,  $\hat{\sigma} = 2.06428$ ,  $\sigma = 2.0493$ )

```

• let
•     n = 20; p = 0.3
•     y = rand(Binomial(n, p), 10000)
•     ( $\hat{m} = \text{mean}(y)$ ,  $m = 20 * 0.3$ ,  $\hat{\sigma} = \text{std}(y)$ ,
•      $\sigma = \text{sqrt}(n * p * (1 - p))$ )
• end

```


► ( $\hat{m} = 0.3053$ ,  $m = 0.3$ )

```

• let
•     n = 20; p = 0.3
•     y = rand(Bernoulli(p), 10000)
•     ( $\hat{m} = \text{mean}(y)$ ,  $m = 0.3$ )
• end

```

## LogNormal



```

• let
•   menw = rand(LogNormal(5.13, 0.17),
•   10000)
•   menwl = log.(menw)
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Log weights
•   of men\n(Normal distribution)",
•   xlabel="Logarithm of weight [lb]",
•   ylabel="Density")
•   density!(menwl)
•   ax = Axis(f[1, 2]; title="Weights of
•   men\n(LogNormal distribution",
•   xlabel="Weight of men [lb]",
•   ylabel="Density")
•   density!(menw)
•   current_figure()
end

```

## Binomial

	parameters	median	mad_sd	mean	std
1	"bv"	6.0	1.483	6.099	2.03

```

• let
•     df = DataFrame(bv = rand(Binomial(20,
•         0.3), 1000))
•     model_summary(df, [:bv])
• end

```

► (mean = 6.0, std = 2.04939)

```

• begin
•     n = 20
•     p = 0.3
•     (mean = n * p, std = √(n * p * (1 -
•         p)))
• end

```

## Poisson

► [3, 3, 6, 2, 1, 3, 3, 4, 6, 7]

```

• rand(Poisson(4.52), 10)

```

## 3.6 - Probability modeling

20689.577579211084

```

• 1 / (pdf(Normal(0.49, 0.04), 0.5) /
•     200000)

```

20.689577579211083

```

• 1 / (1000pdf(Normal(0.49, 0.04), 0.5) /
•     200000)

```



