

See chapter 4 in Regression and Other Stories.

Widen the notebook.

```

•
• html"""
• <style>
•   main {
•     margin: 0 auto;
•     max-width: 2000px;
•     padding-left: max(160px, 10%);
•     padding-right: max(160px, 10%);
•   }
• </style>
• """

```

```

• using Pkg ✓ , DrWatson ✓

```

A typical set of Julia packages to include in notebooks.

```

• begin
•   # Specific to this notebook
•   using GLM ✓
•
•   # Specific to ROSStanPluto
•   using StanSample ✓
•
•   # Graphics related
•   using GLMakie ✓
•
•   # Common data files and functions
•   using RegressionAndOtherStories ✓
• end


```

4.1 Sampling distributions and generative models.

```

• begin
•   Random.seed!(1)
•   a = 1.0
•   b = 2.0
•   x = LinRange(-2, 2, 100)
•   y = a .+ b .* x .+ rand(Normal(0.0,
•   0.2), 100)
• end;

```



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Linear
•   regression")
•   scatter!(x, y)
•   lines!(x, a .+ b .* x; color=:darkred)
•   annotations!("yi = 1.0 + 2.0 * xi +
•   εi", position=(0, -2), fontsize=20)
•   current_figure()
• end

```

```

• stan4_1 = "
• data {
•     int N;
•     vector[N] x;
•     vector[N] y;
• }
• parameters {
•     real a;
•     real b;
•     real<lower=0> sigma;
• }
• model {
•     vector[N] mu;
•     a ~ normal(0.0, 1.5);
•     b ~ normal(1.0, 1.5);
•     sigma ~ exponential(1);
•     mu = a + b * x;
•     y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std
1	"a"	1.00107	0.000306437	0.021698
2	"b"	1.97538	0.000296873	0.018729
3	"sigma"	0.218255	0.000234626	0.015618

```

• let
•     data = (N = length(x), x = x, y = y)
•     global m4_1s = SampleModel("m4.1s",
•         stan4_1)
•     global rc4_1s = stan_sample(m4_1s;
•         data)
•     success(rc4_1s) && describe(m4_1s)
end

```

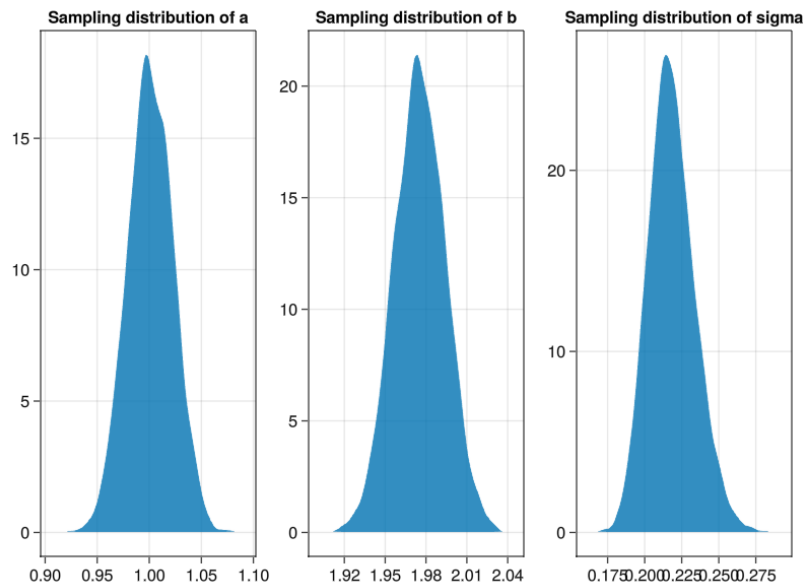
```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gr
updated.

```

	parameters	median	mad_sd	mean	std
1	"a"	1.001	0.022	1.001	0.02
2	"b"	1.975	0.019	1.975	0.01
3	"sigma"	0.217	0.015	0.218	0.01


```
• if success(rc4_1s)
•     post4_1s = read_samples(m4_1s,
•         :dataframe)
•     ms4_1s = model_summary(post4_1s, [:a,
•         :b, :sigma])
end
```



```

• let
•     f = Figure()
•     ax = Axis(f[1, 1]; title="Sampling
•         distribution of a")
•     density!(post4_1s.a)
•     ax = Axis(f[1, 2]; title="Sampling
•         distribution of b")
•     density!(post4_1s.b)
•     ax = Axis(f[1, 3]; title="Sampling
•         distribution of sigma")
•     density!(post4_1s.sigma)
•     current_figure()
• end

```




```

let
  f = Figure()
  ax = Axis(f[1, 1]; title="Linear
  regression")
  scatter!(x, y)
  lines!(x, ms4_1s[:a, :mean] .+
  ms4_1s[:b, :mean] .* x;
  color=:darkred)
  mean_a = round(ms4_1s[:a, :mean];
  digits=2)
  mean_b = round(ms4_1s[:b, :mean];
  digits=2)
  mean_σ = round(ms4_1s[:sigma, :mean];
  digits=2)
  annotations!("y = $(mean_a) + $(mean_b)
  * x + $(mean_σ)", position=(0, -2),
  fontsize=20)
  current_figure()
end

```


4.2 Estimates, standard errors, and confidence intervals.



```

• let
•     f = Figure()
•     ax = Axis(f[1, 1]; title="Sampling
•     distribution of b (revisited)")
•      $\hat{b}$  = ms4_1s[:b, :mean]
•      $\hat{\sigma}$  = ms4_1s[:b, :std]
•     x = LinRange( $\hat{b}$  - 4 $\hat{\sigma}$ ,  $\hat{b}$  + 4 $\hat{\sigma}$ , 100)
•     y = pdf.(Normal( $\hat{b}$ ,  $\hat{\sigma}$ ), x)
•     ylims!(ax, [0, maximum(y) + 1.0])
•     ax.xticks =  $\hat{b}$  - 3 $\hat{\sigma}$  :  $\hat{\sigma}$  :  $\hat{b}$  + 3 $\hat{\sigma}$ 
•     ax.xtickformat = xs -> ["$(i) s.e." for
•     i in -3:3]
•     lines!(x, y)
•     vlines!(ax,  $\hat{b}$ ;
•     ymax=maximum(y)/(maximum(y) + 1.0),
•     color=:grey)
•     vlines!(ax,  $\hat{b}$ - $\hat{\sigma}$ ; ymax=pdf.(Normal( $\hat{b}$ ,
•      $\hat{\sigma}$ ),  $\hat{b}$ - $\hat{\sigma}$ )/(maximum(y) + 1.0),
•     color=:grey)
•     vlines!(ax,  $\hat{b}$ + $\hat{\sigma}$ ; ymax=pdf.(Normal( $\hat{b}$ ,
•      $\hat{\sigma}$ ),  $\hat{b}$ + $\hat{\sigma}$ )/(maximum(y) + 1.0),
•     color=:grey)
•     annotations!("b  $\pm$  1 s.e.", position=
•     ( $\hat{b}$ -0.008, 7.5), textsize=20)
•     x1 = range( $\hat{b}$  -  $\hat{\sigma}$ ,  $\hat{b}$  +  $\hat{\sigma}$ ; length=60)
•     band!(x1, fill(0, length(x1)), pdf.
•     (Normal( $\hat{b}$ ,  $\hat{\sigma}$ ), x1); color = (:blue,
•     0.25))
•     f
end

```

```

• let
•     n = 100
•     b = 2.0
•
•     f = Figure()
•     ax = Axis(f[1,1]; title="Simulation of
•     confidence intervals",
•     xlabel="Simulation",
•     ylabel="Assumed 50% and 95%
•     confidence intervals")
•
•     x = 1:n
•     y = [rand(Uniform(b - 2.1 , b + 2.1),
•     1)[1] for i in 1:n]
•
•     # Assumed s.e. = 1.0
•     lowerrors = fill(0.66, n)
•     higherrors = fill(2, n)
•
•     errorbars!(x, y, lowerrors, color =
•     :red) # same low and high error
•     errorbars!(x, y, higherrors, color =
•     :grey) # same low and high error
•
•     scatter!(x, y, markersize = 3, color =
•     :black)
•     hlines!(ax, [2])
•
•     f
•
end

```

► (estimate = 0.7, se = 0.0144914, int_95 = [0.

```

• let
•     n = 1000
•     yes = 700
•     no = n - yes
•     est = yes/n
•     se = sqrt(est * (1 - est)/n)
•
•     (estimate = est, se = se, int_95 = est
•     .+ quantile.(Normal(0, 1), [0.025,
•     0.975]) * se)
•
end

```


```
► (estimate = 35.8, se = 0.734847, int_50 = [35
```

```
• let
•   y = [35, 34, 38, 35, 37]
•   n = length(y)
•   est = mean(y)
•   se = std(y)/sqrt(n)
•   int_50 = est .+ quantile.(TDist(n-1),
•   [0.25, 0.75]) * se
•   int_95 = est .+ quantile.(TDist(n-1),
•   [0.025, 0.975]) * se
•
•   (estimate = est, se = se, int_50 =
•   int_50, int_95 = int_95)
end
```

```
df_poll =
```

	poll1	poll2	poll3	poll4	poll5
1	2002	10.0	70.0	25.0	5.0
2	2002	5.0	72.0	25.0	3.0
3	2001	10.0	68.0	26.0	6.0
4	2001	5.0	65.0	27.0	8.0
5	2001	2.0	67.0	25.0	8.0
6	2000	8.0	67.0	28.0	5.0
7	2000	6.0	66.0	26.0	8.0
8	2000	2.0	66.0	28.0	6.0
9	1999	5.0	71.0	22.0	7.0
10	1995	9.0	77.0	13.0	10.0
⋮	more				
32	1937	12.0	60.0	33.0	7.0

```
• df_poll = CSV.read(ros_datadir("Death",
"polls.csv"), DataFrame)
```



```

• begin
•     f = Figure()
•     ax = Axis(f[1, 1]; title="Death
•         penalty opinions", xlabel="Year",
•         ylabel="Percentage support for the
•         death penalty")
•     scatter!(df_poll.year, df_poll.support
•         .* 100)
•     err_lims =
•         [100(sqrt(df_poll.support[i]*(1-
•             df_poll.support[i])/1000)) for i in
•             1:nrow(df_poll)]
•     errorbars!(df_poll.year,
•         df_poll.support .* 100, err_lims,
•         color = :red)
•     f
• end

```

4.3 Bias and unmodeled uncertainty.

4.4 Statistical significance, hypothesis testing, and statistical errors.

4.5 Problems with the concept of statistical significance.

4.6 Example of hypothesis testing: 55,000 residents need your help!

4.7 Moving beyond hypothesis testing.