

## See chapter 3 in Regression and Other Stories.

.....

### Widen the notebook

```
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

```
• using Pkg ✓ , DrWatson ✓
```

```
• begin  
•     # Specific to this notebook  
•     using GLM ✓  
•   
•     # Graphics related  
•     using GLMakie ✓  
•   
•     # Specific to ROSStanPluto  
•     using StanSample ✓  
•   
•     # Common data files and functions  
•     using RegressionAndOtherStories ✓  
• end
```

### 3.1 - Weighted averages

pop =

|   | stratum | country         | population | avera |
|---|---------|-----------------|------------|-------|
| 1 | 1       | "United States" | 310000000  | 36.8  |
| 2 | 2       | "Mexico"        | 112000000  | 26.7  |
| 3 | 3       | "Canada"        | 34000000   | 40.7  |

```
• pop = DataFrame(stratum=1:3, country=[  
  "United States", "Mexico", "Canada"],  
  population=Int[310e6, 112e6, 34e6],  
  average_age=[36.8, 26.7, 40.7])
```

34.61008771929824

```
• mean(pop.average_age,  
  weights(pop.population))
```


► [0.679825, 0.245614, 0.0745614]

```
• weights(pop.population)/sum(pop.population)
```

|   | variable     | mean    | min      | median  |   |
|---|--------------|---------|----------|---------|---|
| 1 | :stratum     | 2.0     | 1        | 2.0     | 3 |
| 2 | :country     | nothing | "Canada" | nothing | " |
| 3 | :population  | 1.52e8  | 34000000 | 1.12e8  | 3 |
| 4 | :average_age | 34.7333 | 26.7     | 36.8    | 4 |

```
• describe(pop)
```


## 3.3 - Graphing a line



```

• let
•   data = LinRange(0.01, 0.99, 200)
•   f = Figure(resolution = (800, 800))
•
•   for (i, scale) in enumerate([identity,
•                                log10, log2, log, sqrt, Makie.logit])
•
•       row, col = fldmod1(i, 2)
•       Axis(f[row, col], yscale = scale,
•            title = string(scale),
•            yminorticksvisible = true,
•            yminorgridvisible = true,
•            yminorticks =
•            IntervalsBetween(8))
•
•       lines!(data, color = :blue)
•   end
•
•   f
• end

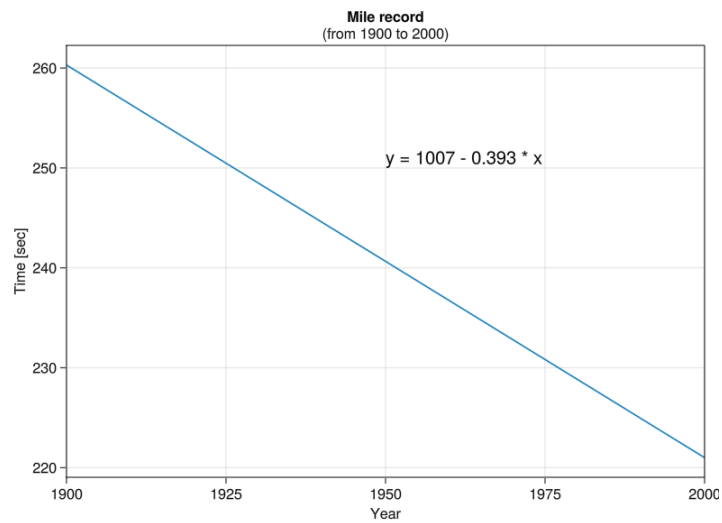
```



```

• let
•   data = 10 .^ LinRange(0.01, 5.0, 200)
•   f = Figure(resolution = (800, 300))
•
•   for (i, scale) in enumerate([log10,
•   log2, log])
•       row, col = fldmod1(i, 2)
•       Axis(f[1, i], yscale = scale, xscale
• = scale, title = string(scale),
•         yminorticksvisible = true,
•         yminorgridvisible = true,
•         yminorticks =
•         IntervalsBetween(8))
•
•       lines!(data, color = :blue)
•   end
• f
end

```



```

let
    f = Figure()
    ax = Axis(f[1, 1]; title = "Mile
record", subtitle = "(from 1900 to
2000)", xlabel = "Year", ylabel = "Time
[sec]")
    xlims!(ax, 1900, 2000)
    ax.xticks = 1900:25:2000
    x = 1900:2000
    y = 1007 .- 0.393 .* x
    lines!(x, y)
    annotations!("y = 1007 - 0.393 * x",
position=(1950, 250))
    current_figure()
end

```

## 3.4 - Log and exponential scales

Simulated data for metabolic.

|            | body_mass | rate    |
|------------|-----------|---------|
| <b>1</b>   | 4.27509   | 4.56164 |
| <b>2</b>   | 5.10826   | 5.20084 |
| <b>3</b>   | 5.17192   | 5.35231 |
| <b>4</b>   | 5.22901   | 5.25554 |
| <b>5</b>   | 5.66336   | 5.65967 |
| <b>6</b>   | 5.67954   | 5.51303 |
| <b>7</b>   | 6.10486   | 5.93264 |
| <b>8</b>   | 6.11608   | 5.88396 |
| <b>9</b>   | 6.19732   | 5.94754 |
| <b>10</b>  | 6.24436   | 5.98998 |
| ⋮ more     |           |         |
| <b>200</b> | 9.19991   | 8.19632 |

```

• begin
•   x = sort(rand(Uniform(0.01, 10000),
•   200))
•   y = 4.1 * x.^0.74 .+ [rand.(Normal.(0,
•   sqrt(x[i])), 1)[1] for i in 1:length(x)]
•   metabolic = DataFrame(:body_mass => log.
•   (x), :rate => log.(y))
end

```

```

• stan3_1 = "
• data {
•   int N;
•   vector[N] m;
•   vector[N] r;
• }
• parameters {
•   real a;
•   real b;
•   real sigma;
• }
• model {
•   vector[N] mu;
•   a ~ normal(0, 0.3);
•   b ~ normal(0, 0.3);
•   sigma ~ exponential(1);
•   mu = a + b * m;
•   r ~ normal(mu, sigma);
• };

```

|   | parameters | mean      | mcse        | std     |
|---|------------|-----------|-------------|---------|
| 1 | "a"        | 1.38631   | 0.000827076 | 0.02180 |
| 2 | "b"        | 0.742607  | 9.95526e-5  | 0.00263 |
| 3 | "sigma"    | 0.0325087 | 3.80206e-5  | 0.00163 |

```

• let
•   data = (N =
•     length(metabolic.body\_mass), m =
•     metabolic.body\_mass, r = metabolic.rate)
•   global m3_1s = SampleModel("m3.1s",
•     stan3\_1)
•   global rc3_1s = stan_sample(m3_1s; data)
•   success(rc3_1s) && describe(m3_1s)
• end

```


|   | parameters | median | mad_sd | mean  | std   |
|---|------------|--------|--------|-------|-------|
| 1 | "a"        | 1.386  | 0.023  | 1.386 | 0.023 |
| 2 | "b"        | 0.743  | 0.003  | 0.743 | 0.003 |
| 3 | "sigma"    | 0.032  | 0.002  | 0.033 | 0.002 |

```

• if success(rc3_1s)
•   post3_1s = read_samples(m3_1s,
•     :dataframe)
•   ms3_1s = model_summary(post3_1s, [:a,
•     :b, :sigma])
end

```





```

• let
•   x = LinRange(1, 10000, 1000)
•   y = 4.1 * x.^0.74
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Metabolic
•   rate (linear scale)", xlabel="Body mass
•   [kg]", ylabel="Metobolic rate [W]")
•   scatter!(exp.(metabolic.body_mass), exp.
•   (metabolic.rate))
•   lines!(x, y; color=:darkred)

•   ax = Axis(f[1, 2]; title="Metabolic
•   rate (log-log scale)", xscale=log,
•   yscale=log,
•   xlabel="log(body mass [kg])",
•   ylabel="log(metobolic rate [W])")
•   x =
•   LinRange(minimum(metabolic.body_mass),
•   maximum(metabolic.body_mass), 100)
•   scatter!(metabolic.body_mass,
•   metabolic.rate)
•   lines!(x, ms3_1s[:a, :mean] .+
•   ms3_1s[:b, :mean] * x; color=:darkred)
•   current_figure()

end

```

```

100-element LinRange{Float64, Int64}:
 4.27509, 4.32483, 4.37458, 4.42432, 4.47407, ...,

```

```

• LinRange(minimum(metabolic.body_mass),
  maximum(metabolic.body_mass), 100)

```

```

57.69669926958961

```

```

• exp(exp(1.4))

```


## 3.5 - Probability distributions

|        | height  | sex      |
|--------|---------|----------|
| 1      | 65.3008 | "female" |
| 2      | 62.9675 | "female" |
| 3      | 63.3444 | "female" |
| 4      | 60.8675 | "female" |
| 5      | 56.5783 | "female" |
| 6      | 65.3281 | "female" |
| 7      | 63.075  | "female" |
| 8      | 65.6519 | "female" |
| 9      | 61.0603 | "female" |
| 10     | 65.0997 | "female" |
| ⋮ more |         |          |
| 200000 | 71.0511 | "male"   |

```
• begin
•   N = 100000
•   heights = DataFrame()
•   height = vcat(rand(Normal(63.7, 2.7),
•   N),
•   rand(Normal(69.1, 2.9), N))
•   sex = repeat(["female", "male"],
•   inner=N)
•   heights.height = height
•   heights.sex = sex
•   heights
end
```

```
► (mean = 63.6881, var = 7.29857, std = 2.70159, i
```

```
• begin
•   menHeights = heights[heights.sex ==
•   "male", :height]
•   womenHeights = heights[heights.sex ==
•   "female", :height]
•   (mean=mean(womenHeights),
•   var=var(womenHeights),
•   std=std(womenHeights),
•   median=median(womenHeights),
•   mad_sd=mad(womenHeights))
end
```



```
• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Density
•   heights")
•   density!(womenHeights; color=color
•   = (:lightgreen, 0.4), label="Women")
•   density!(menHeights; color=color =
•   (:lightblue, 0.4), label="Men")
•   axislegend()
•   f
end
```

```
0.49714220980937984
```


```
• begin
•   wdf = Normal(63.65, 2.68)
•   cdf(wdf, 63.65 + 0.67 * 2.68) -
•   cdf(wdf, 63.65 - 0.67 * 2.68)
end
```

0.6826894921370859

- $\text{cdf}(\text{wdf}, 63.65 + 2.68) - \text{cdf}(\text{wdf}, 63.65 - 2.68)$

0.9544997361036417

- $\text{cdf}(\text{wdf}, 63.65 + 2 \cdot 2.68) - \text{cdf}(\text{wdf}, 63.65 - 2 \cdot 2.68)$



```

• let
•   wdf = Normal(63.65, 2.68)
•   x = range(55.0, 72.5 ; length=100)
•   lines(x, pdf.(wdf, x); color=:darkblue)
•
•   x1 = range(63.65 - 3 * 2.68, 63.65 - 2
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.25), label =
•             "Label")
•
•   x1 = range(63.65 + 2 * 2.68, 63.65 + 3
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.25), label =
•             "Label")
•
•   x1 = range(63.65 - 2 * 2.68, 63.65 - 1
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.45), label =
•             "Label")
•
•   x1 = range(63.65 + 1 * 2.68, 63.65 + 2
•             * 2.68; length=20)
•   band!(x1, fill(0, length(x1)), pdf.
•         (wdf, x1);
•         color = (:blue, 0.45), label =
•             "Label")
•
•   x1 = range(63.65 - 1 * 2.68, 63.65;
•             length=20)
•
•

```

```

• band!(x1, fill(0, length(x1)), pdf.
• (wdf, x1);
• color = (:blue, 0.55), label =
• "Label")
•

x1 = range(63.65, 63.65 + 2.68;
length=20)
band!(x1, fill(0, length(x1)), pdf.
(wdf, x1);
color = (:blue, 0.55), label =
"Label")

text!("68%", position = (63.65, 0.05),
align = (:center, :center),
textsize = 30)
text!("13.5%", position = (67.5, 0.02),
align = (:center, :center),
textsize = 20)
text!("13.5%", position = (59.6, 0.02),
align = (:center, :center),
textsize = 20)
text!("2.5%", position = (69.75,
0.0045), align = (:center, :center),
textsize = 15)
text!("2.5%", position = (57.7,
0.0045), align = (:center, :center),
textsize = 15)
current_figure()

end

```

► ( $\hat{m} = 6.01$ ,  $m = 6.0$ ,  $\hat{\sigma} = 2.06807$ ,  $\sigma = 2.04939$ )

```

• let
• n = 20; p = 0.3
• y = rand(Binomial(n, p), 10000)
• ( $\hat{m} = \text{mean}(y)$ ,  $m = 20 * 0.3$ ,  $\hat{\sigma} = \text{std}(y)$ ,
•  $\sigma = \text{sqrt}(n * p * (1 - p))$ )
end

```


► ( $\hat{m} = 0.3075$ ,  $m = 0.3$ )

```

• let
• n = 20; p = 0.3
• y = rand(Bernoulli(p), 10000)
• ( $\hat{m} = \text{mean}(y)$ ,  $m = 0.3$ )
• end

```

## LogNormal



```

• let
•   menw = rand(LogNormal(5.13, 0.17),
•   10000)
•   menwl = log.(menw)
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Log weights
•   of men\n(Normal distribution)",
•   xlabel="Logarithm of weight [lb]",
•   ylabel="Density")
•   density!(menwl)
•   ax = Axis(f[1, 2]; title="Weights of
•   men\n(LogNormal distribution",
•   xlabel="Weight of men [lb]",
•   ylabel="Density")
•   density!(menw)
•   current_figure()
end

```

## Binomial

|   | parameters | median | mad_sd | mean  | st   |
|---|------------|--------|--------|-------|------|
| 1 | "bv"       | 6.0    | 1.483  | 6.174 | 2.00 |

```

• let
•   df = DataFrame(bv = rand(Binomial(20,
•   0.3), 1000))
•   model_summary(df, [:bv])
end

```

```
► (mean = 6.0, std = 2.04939)
```

```
• begin  
•   n = 20  
•   p = 0.3  
•   (mean = n * p, std =  $\sqrt{n * p * (1 - p)}$ )  
• end
```

## Poisson

```
► [4, 5, 7, 7, 0, 2, 3, 5, 2, 3]
```

```
• rand(Poisson(4.52), 10)
```

## 3.6 - Probability modeling

```
20689.577579211084
```

```
• 1 / (pdf(Normal(0.49, 0.04), 0.5) / 200000)
```

```
20.689577579211083
```

```
• 1 / (1000pdf(Normal(0.49, 0.04), 0.5) /  
200000)
```