

See Chapter 6 in Regression and Other Stories.

Widen the notebook.

```
• html"""  
• <style>  
•     main {  
•         margin: 0 auto;  
•         max-width: 2000px;  
•         padding-left: max(160px, 10%);  
•         padding-right: max(160px, 10%);  
•     }  
• </style>  
• """
```

```
• using Pkg ✓ , DrWatson ✓
```

A typical set of Julia packages to include in notebooks.

```
• begin  
•     # Specific to this notebook  
•     using GLM ✓  
•  
•     # Specific to ROSStanPluto  
•     using StanSample ✓  
•  
•     # Graphics related  
•     using GLMakie ✓  
•  
•     # Common data files and functions  
•     using RegressionAndOtherStories ✓  
• end
```

```
Replacing docs for `RegressionAndOtherStories.tr  
DataFrame, AbstractString}` in module `Regressio
```

6.1 Regression models.

6.2 Fitting a simple regression to fake data.

| | x | y |
|----|------|----------|
| 1 | 1.0 | 1.18017 |
| 2 | 2.0 | 0.66976 |
| 3 | 3.0 | 0.951537 |
| 4 | 4.0 | 0.941839 |
| 5 | 5.0 | 1.90931 |
| 6 | 6.0 | 1.6079 |
| 7 | 7.0 | 2.41162 |
| 8 | 8.0 | 3.01127 |
| 9 | 9.0 | 2.70929 |
| 10 | 10.0 | 2.80496 |
| ⋮ | more | |
| 20 | 20.0 | 6.03491 |

```
• let
•   n = 20
•   x = LinRange(1, n, 20)
•   a = 0.2
•   b = 0.3
•   sigma = 0.5
•   y = a .+ b .* x .+ rand(Normal(0,
•   sigma), n)
•   global fake = DataFrame(x=x, y=y)
end
```

```

• stan6_1 = "
• data {
•   int N;
•   vector[N] x;
•   vector[N] y;
• }
• parameters {
•   real a;
•   real b;
•   real<lower=0> sigma;
• }
• model {
•   vector[N] mu;
•   a ~ uniform(-2, 2);
•   b ~ uniform(-2, 2);
•   sigma ~ uniform(0, 10);
•   mu = a + b * x;
•   y ~ normal(mu, sigma);
• }";

```

| | parameters | mean | mcse | std |
|---|------------|----------|-------------|----------|
| 1 | "a" | 0.284945 | 0.00587742 | 0.230988 |
| 2 | "b" | 0.29781 | 0.000481178 | 0.019324 |
| 3 | "sigma" | 0.474901 | 0.00220934 | 0.090591 |

```

• let
•   data = (N=nrow(fake), x=fake.x,
•   y=fake.y)
•   global m6_1s = SampleModel("m6_1s",
•   stan6_1)
•   global rc6_1s = stan_sample(m6_1s; data)
•   success(rc6_1s) && describe(m6_1s)
end

```

```

/
c
Informational Message: The current Metropolis
rejected because of the following issue:
Exception: normal_lpdf: Scale parameter is 0,
r/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
7, column 1 to column 23)
If this warning occurs sporadically, such as f
types like covariance matrices, then the sampl
but if this warning occurs often then your mod
conditioned or misspecified.

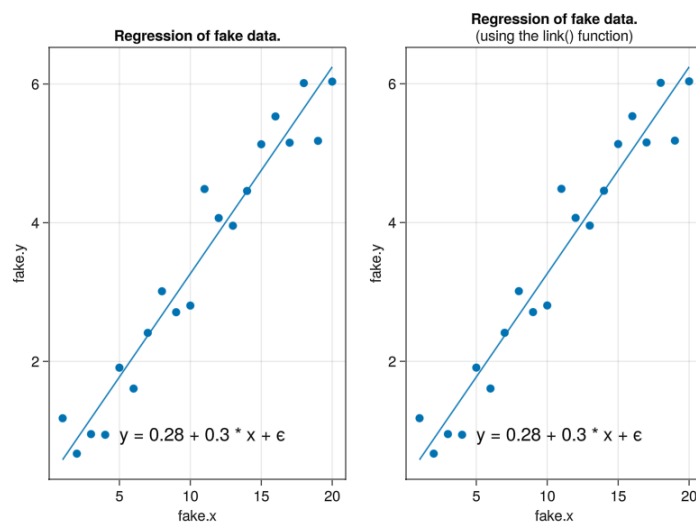
```

| | parameters | median | mad_sd | mean | st |
|---|------------|--------|--------|-------|------|
| 1 | "a" | 0.282 | 0.229 | 0.285 | 0.23 |
| 2 | "b" | 0.298 | 0.018 | 0.298 | 0.01 |
| 3 | "sigma" | 0.462 | 0.082 | 0.475 | 0.09 |

```

• if success(rc6_1s)
•   post6_1s = read_samples(m6_1s,
•     :dataframe)
•   ms6_1s = model_summary(post6_1s, [:a,
•     :b, :sigma])
end

```



```

let
  f = Figure()

  ax = Axis(f[1, 1]; title="Regression of
fake data.", xlabel="fake.x",
ylabel="fake.y")
  scatter!(fake.x, fake.y)
  x = 1:0.01:20
  y = ms6_1s[:, :mean] .+ ms6_1s[:, :b,
:mean] .* x
  lines!(x, y)
  â = round(ms6_1s[:, :mean]; digits=2)
  b̂ = round(ms6_1s[:, :b, :mean]; digits=2)
  annotations!("y = $(â) + $(b̂) * x + ε";
position=(5, 0.8))

  ax = Axis(f[1, 2]; title="Regression of
fake data.", subtitle="(using the
link() function)",
xlabel="fake.x", ylabel="fake.y")
  scatter!(fake.x, fake.y)
  xrange = LinRange(1, 20, 200)
  y = mean.(link(post6_1s, (r,x) -> r.a +
x * r.b, xrange))
  lines!(xrange, y)
  annotations!("y = $(â) + $(b̂) * x + ε";
position=(5, 0.8))

  current_figure()
end

```

| | parameters | simulated | median | mad_sd |
|---|------------|-----------|--------|--------|
| 1 | :a | 0.2 | 0.282 | 0.229 |
| 2 | :b | 0.3 | 0.298 | 0.018 |
| 3 | :sigma | 0.5 | 0.462 | 0.082 |

```
• DataFrame(parameters = Symbol.  
  (names(post6_1s)), simulated = [0.2, 0.3,  
  0.5], median = ms6_1s[:, :median], mad_sd =  
  ms6_1s[:, :mad_sd])
```

6.3 Interpret coefficients as comparisons, not effects.

| | earnk | height | male |
|--------|-------|--------|------|
| 1 | 50.0 | 74 | 1 |
| 2 | 60.0 | 66 | 0 |
| 3 | 30.0 | 64 | 0 |
| 4 | 25.0 | 65 | 0 |
| 5 | 50.0 | 63 | 0 |
| 6 | 62.0 | 68 | 0 |
| 7 | 51.0 | 63 | 0 |
| 8 | 9.0 | 64 | 0 |
| 9 | 29.0 | 62 | 0 |
| 10 | 32.0 | 73 | 1 |
| ⋮ more | | | |
| 1816 | 6.0 | 68 | 1 |

```

• begin
•   earnings =
•   CSV.read(ros_datadir("Earnings",
•   "earnings.csv"), DataFrame)
•   earnings[:, [:earnk, :height, :male]]
end

```

| | variable | mean | min | median | max | n |
|---|----------|----------|-----|--------|-------|---|
| 1 | :earnk | 21.1473 | 0.0 | 16.0 | 400.0 | 0 |
| 2 | :height | 66.5688 | 57 | 66.0 | 82 | 0 |
| 3 | :male | 0.371696 | 0 | 0.0 | 1 | 0 |

```

• describe(earnings[:, [:earnk, :height,
• :male]])

```

```

• stan6_2 = "
• data {
•   int N;
•   vector[N] male;
•   vector[N] height;
•   vector[N] earnk;
• }
• parameters {
•   real a;
•   real b;
•   real c;
•   real<lower=0> sigma;
• }
• model {
•   vector[N] mu;
•   sigma ~ exponential(1);
•   mu = a + b * height + c * male;
•   earnk ~ normal(mu, sigma);
• };

```

| | parameters | mean | mcse | std |
|---|------------|----------|------------|----------|
| 1 | "a" | -25.8136 | 0.342803 | 11.8855 |
| 2 | "b" | 0.645866 | 0.00533156 | 0.184075 |
| 3 | "c" | 10.6543 | 0.0382441 | 1.48058 |
| 4 | "sigma" | 21.2877 | 0.00837708 | 0.35878 |

```

• let
•   data = (N=nrow(earnings),
•           height=earnings.height,
•           male=earnings.male,
•           earnk=earnings.earnk)
•   global m6_2s = SampleModel("m6_2s",
•                               stan6_2)
•   global rc6_2s = stan_sample(m6_2s; data)
•   success(rc6_2s) && describe(m6_2s)
• end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gn/T/
d.

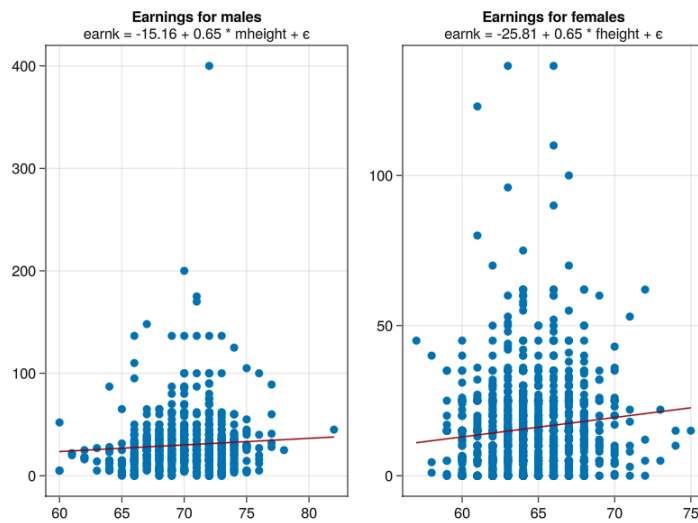
```


| | parameters | median | mad_sd | mean | st |
|---|------------|---------|--------|---------|------|
| 1 | "a" | -25.412 | 11.818 | -25.814 | 11.8 |
| 2 | "b" | 0.64 | 0.184 | 0.646 | 0.18 |
| 3 | "c" | 10.678 | 1.469 | 10.654 | 1.48 |
| 4 | "sigma" | 21.274 | 0.366 | 21.288 | 0.36 |

```

• if success(rc6_2s)
•   post6_2s = read_samples(m6_2s,
•     :dataframe)
•   ms6_2s = model_summary(post6_2s, [:a,
•     :b, :c, :sigma])
end

```



```

let
  â, b̂, ĉ = round.(ms6_2s[:, :mean];
    digits=2)

  fig = Figure()

  ax = Axis(fig[1, 1]; title="Earnings
    for males", subtitle="earnk = $(round(ĉ
    + â; digits=2)) + $(b̂) * mheight + e")
  m = sort(earnings[earnings.male .== 1,
    [:height, :earnk]])
  scatter!(m.height, m.earnk)
  mheight_range =
    LinRange(minimum(m.height),
    maximum(m.height), 200)
  earnk = mean.(link(post6_2s, (r,x) ->
    r.c + r.a + x * r.b, mheight_range))

  #earnk = ms6_2s[:, :c, "mean"] +
  #ms6_2s[:, :a, "mean"] .* ms6_2s[:, :b,
  #"mean"] .* mheight
  lines!(mheight_range, earnk;
    color=:darkred)

  ax = Axis(fig[1, 2]; title="Earnings
    for females", subtitle="earnk = $(â) +
    $(b̂) * fheight + e")
  f = sort(earnings[earnings.male .== 0,
    [:height, :earnk]])
  scatter!(f.height, f.earnk)
  fheight_range =
    LinRange(minimum(f.height),
    maximum(f.height), 200)
  earnk = mean.(link(post6_2s, (r,x) ->
    r.a + x * r.b, fheight_range))

```

```

    lines!(fheight_range, earnk;
           color=:darkred)

    fig
end

```

```
R2 = 0.10735394024175804
```

```

• R2 = 1 - ms6_2s[:sigma, :mean]^2 /
  std(earnings.earnk)^2

```

6.4 Historical origins of regression.

```

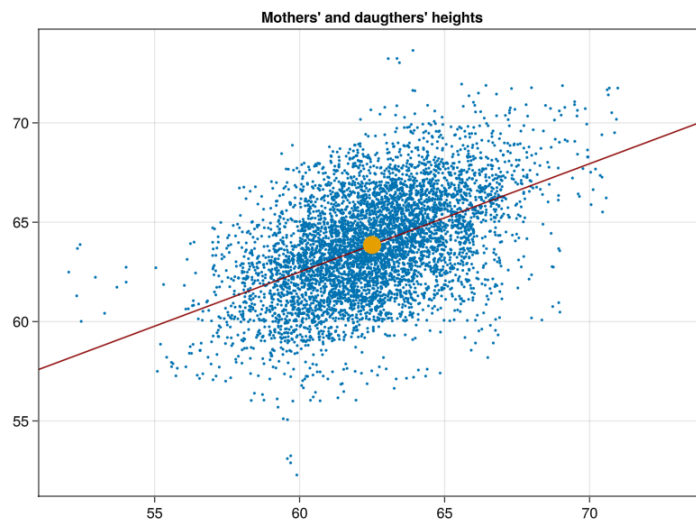
• stan6_3 = "
• data {
•   int N;
•   vector[N] m_height;
•   vector[N] d_height;
• }
• parameters {
•   real a;
•   real b;
•   real<lower=0> sigma;
• }
• model {
•   vector[N] mu;
•   sigma ~ exponential(1);
•   mu = a + b * m_height;
•   d_height ~ normal(mu, sigma);
• }";

```

```
heights =
```

| | daughter_height | mother_height |
|-------------|-----------------|---------------|
| 1 | 52.5 | 59.5 |
| 2 | 52.5 | 59.5 |
| 3 | 53.5 | 59.5 |
| 4 | 53.5 | 59.5 |
| 5 | 55.5 | 59.5 |
| 6 | 55.5 | 59.5 |
| 7 | 55.5 | 59.5 |
| 8 | 55.5 | 59.5 |
| 9 | 56.5 | 58.5 |
| 10 | 56.5 | 58.5 |
| ⋮ more | | |
| 5524 | 73.5 | 63.5 |

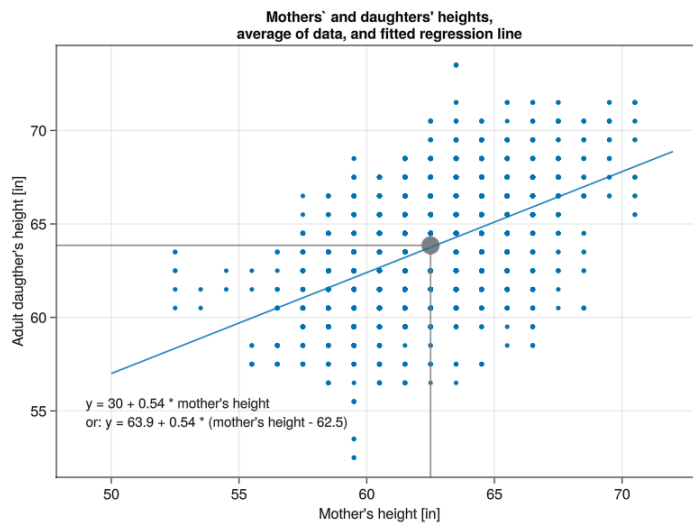
```
• heights =  
  CSV.read(ros_datadir("PearsonLee",  
    "heights.csv"), DataFrame)
```

```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Mothers' and
•   daughters' heights")
•   xlims!(ax, 51, 74)
•   scatter!(jitter.
•   (heights.mother_height), jitter.
•   (heights.daughter_height); markersize=3)
•   x_range = LinRange(51, 74, 100)
•   lines!(x_range, mean.(link(post6_3s,
•   (r, x) -> r.a + r.b * x, x_range));
•   color=:darkred)
•   scatter!([mean(heights.mother_height)],
•   [mean(heights.daughter_height)];
•   markersize=20)
•   f
• end

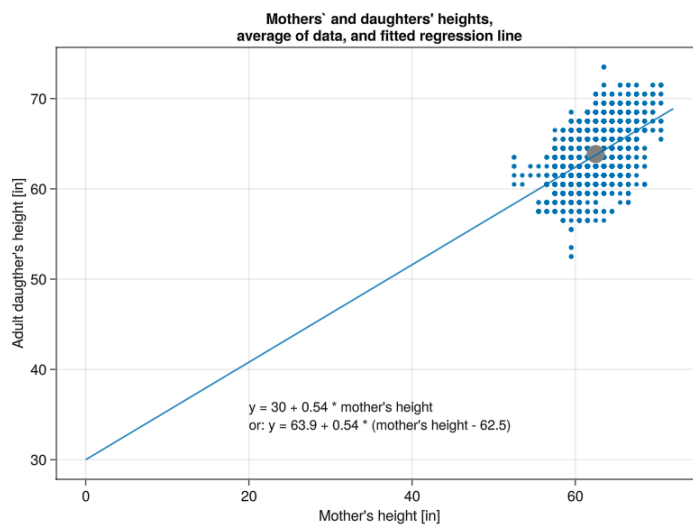
```



```

let
  f = Figure()
  ax = Axis(f[1, 1]; title="Mothers' and
    daughters' heights,\naverage of data,
    and fitted regression line",
    xlabel="Mother's height [in]",
    ylabel="Adult daughter's height
      [in]")
  scatter!(heights.mother_height,
    heights.daughter_height; markersize=5)
  xrange = LinRange(50, 72, 100)
  y = 30 .+ 0.54 .* xrange
  m̄ = mean(heights.mother_height)
  d̄ = mean(heights.daughter_height)
  scatter!([m̄], [d̄]; markersize=20,
    color=:gray)
  lines!(xrange, y)
  vlines!(ax, m̄; ymax=0.55, color=:grey)
  hlines!(ax, d̄; xmax=0.58, color=:grey)
  annotations!("y = 30 + 0.54 * mother's
    height", position=(49, 55), fontsize=15)
  annotations!("or: y = 63.9 + 0.54 *
    (mother's height - 62.5)", position=
    (49, 54), fontsize=15)
  f
end

```



```

let
  f = Figure()
  ax = Axis(f[1, 1]; title="Mothers' and
    daughters' heights,\naverage of data,
    and fitted regression line",
    xlabel="Mother's height [in]",
    ylabel="Adult daughter's height
    [in]")
  scatter!(heights.mother_height,
    heights.daughter_height; markersize=5)
  xrange = LinRange(0, 72, 100)
  y = 30 .+ 0.54 .* xrange
  m̄ = mean(heights.mother_height)
  d̄ = mean(heights.daughter_height)
  scatter!([m̄], [d̄]; markersize=20,
    color=:gray)
  lines!(xrange, y)
  annotations!("y = 30 + 0.54 * mother's
    height", position=(20, 35), fontsize=15)
  annotations!("or: y = 63.9 + 0.54 *
    (mother's height - 62.5)", position=
    (20, 33), fontsize=15)
  f
end

```



```

• stan6_4 = "
• data {
•   int N;
•   vector[N] m;
•   vector[N] d;
• }
• parameters {
•   real a;
•   real b;
•   real<lower=0> sigma;
• }
• model {
•   vector[N] mu;
•   a ~ normal(25, 3);
•   b ~ normal(0, 0.5);
•   sigma ~ exponential(1);
•   mu = a + b * m;
•   d ~ normal(mu, sigma);
• }";

```

| | parameters | mean | mcse | std |
|---|------------|---------|-------------|-----------|
| 1 | "a" | 29.5157 | 0.0235803 | 0.769221 |
| 2 | "b" | 0.54945 | 0.000376458 | 0.0122975 |
| 3 | "sigma" | 2.26274 | 0.000620423 | 0.0223053 |

```

• let
•   data = (N = nrow(heights), m =
•           heights.mother_height, d =
•           heights.daughter_height)
•   global m6_4s = SampleModel("m6_4s",
•                               stan6_4)
•   global rc6_4s = stan_sample(m6_4s; data)
•   success(rc6_4s) && describe(m6_4s)
• end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
d.

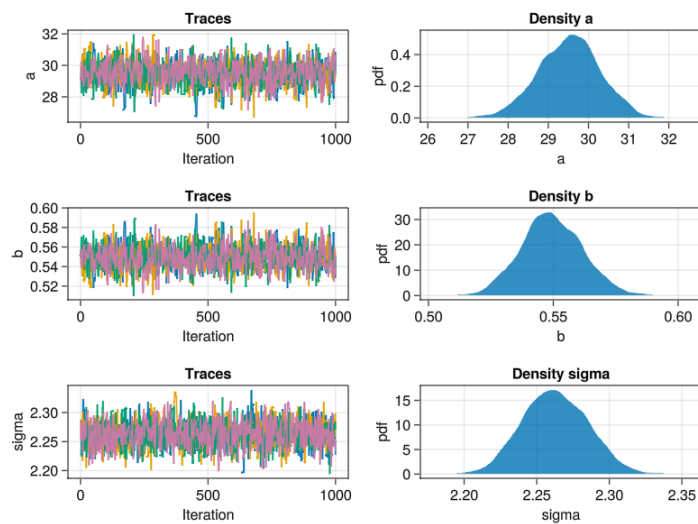
```

| | parameters | median | mad_sd | mean | std |
|---|------------|--------|--------|--------|-------|
| 1 | "a" | 29.541 | 0.769 | 29.516 | 0.769 |
| 2 | "b" | 0.549 | 0.012 | 0.549 | 0.012 |
| 3 | "sigma" | 2.262 | 0.023 | 2.263 | 0.023 |

```

• if success(rc6_4s)
•   post6_4s = read_samples(m6_4s,
•     :dataframe)
•   ms6_4s = model_summary(post6_4s, [:a,
•     :b, :sigma])
end

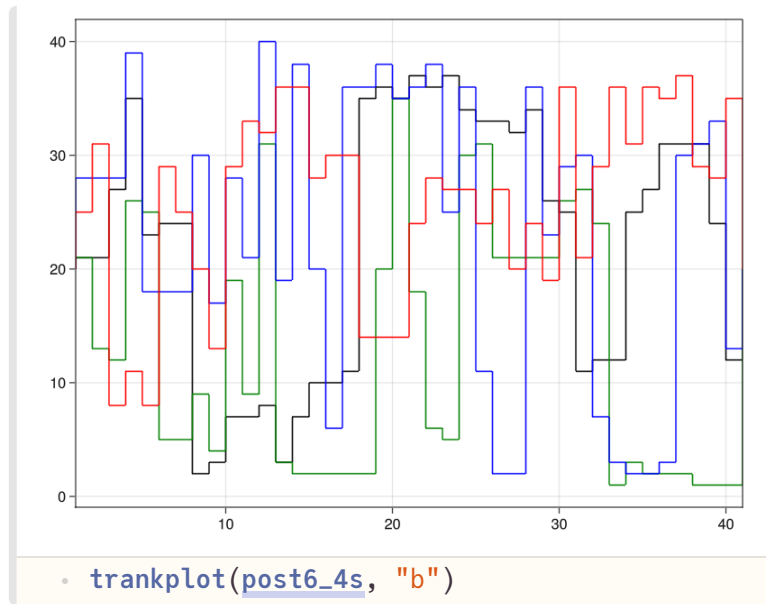
```



```

• plot_chains(post6_4s, [:a, :b, :sigma])

```



Above trankplot and the low ess numbers a couple of cells earlier do not look healthy.

6.5 The paradox of regression to the mean.

| | midterm | final |
|-------------|---------|---------|
| 1 | 39.0331 | 31.4929 |
| 2 | 58.12 | 51.6225 |
| 3 | 58.9264 | 47.0342 |
| 4 | 47.4465 | 35.3108 |
| 5 | 62.3286 | 38.493 |
| 6 | 50.4649 | 66.393 |
| 7 | 52.7304 | 44.1828 |
| 8 | 43.209 | 37.7982 |
| 9 | 23.9576 | 36.4188 |
| 10 | 55.5456 | 77.6372 |
| : more | | |
| 1000 | 40.0598 | 29.367 |

```

• let
•   n = 1000
•   true_ability = rand(Normal(50, 10), n)
•   noise_1 = rand(Normal(0, 10), n)
•   noise_2 = rand(Normal(0, 10), n)
•   midterm = true_ability + noise_1
•   final = true_ability + noise_2
•   global exams =
•   DataFrame(midterm=midterm, final=final)
• end

```

```

• stan6_5 = "
• data {
•   int N;
•   vector[N] midterm;
•   vector[N] final;
• }
• parameters {
•   real a;
•   real b;
•   real<lower=0> sigma;
• }
• model {
•   vector[N] mu;
•   sigma ~ exponential(1);
•   mu = a + b * midterm;
•   final ~ normal(mu, sigma);
• }";

```

| | parameters | mean | mcse | std |
|---|------------|----------|-------------|----------|
| 1 | "a" | 24.0196 | 0.0393881 | 1.40601 |
| 2 | "b" | 0.525487 | 0.000759792 | 0.02697 |
| 3 | "sigma" | 12.0779 | 0.00617686 | 0.261575 |

```

• let
•   data = (N=nrow(exams),
•   midterm=exams.midterm,
•   final=exams.final)
•   global m6_5s = SampleModel("m6_5s",
•   stan6_5)
•   global rc6_5s = stan_sample(m6_5s; data)
•   success(rc6_5s) && describe(m6_5s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
d.

```

| | parameters | median | mad_sd | mean | std |
|---|------------|--------|--------|--------|------|
| 1 | "a" | 24.05 | 1.377 | 24.02 | 1.40 |
| 2 | "b" | 0.525 | 0.026 | 0.525 | 0.02 |
| 3 | "sigma" | 12.072 | 0.252 | 12.078 | 0.26 |

```

• if success(rc6_5s)
•   post6_5s = read_samples(m6_5s,
•   :dataframe)
•   ms6_5s = model_summary(post6_5s, [:a,
•   :b, :sigma])
end

```

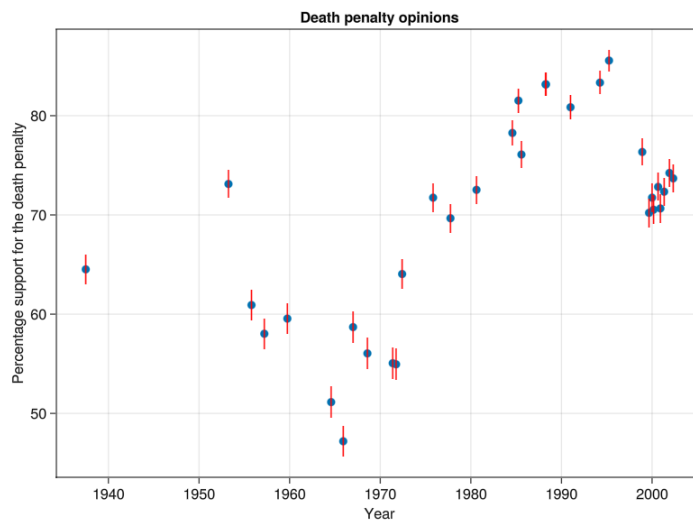
df_poll =

| | poll1 | poll2 | poll3 | poll4 | poll5 |
|----|-------|-------|-------|-------|-------|
| 1 | 2002 | 10.0 | 70.0 | 25.0 | 5.0 |
| 2 | 2002 | 5.0 | 72.0 | 25.0 | 3.0 |
| 3 | 2001 | 10.0 | 68.0 | 26.0 | 6.0 |
| 4 | 2001 | 5.0 | 65.0 | 27.0 | 8.0 |
| 5 | 2001 | 2.0 | 67.0 | 25.0 | 8.0 |
| 6 | 2000 | 8.0 | 67.0 | 28.0 | 5.0 |
| 7 | 2000 | 6.0 | 66.0 | 26.0 | 8.0 |
| 8 | 2000 | 2.0 | 66.0 | 28.0 | 6.0 |
| 9 | 1999 | 5.0 | 71.0 | 22.0 | 7.0 |
| 10 | 1995 | 9.0 | 77.0 | 13.0 | 10.0 |
| ⋮ | more | | | | |
| 32 | 1937 | 12.0 | 60.0 | 33.0 | 7.0 |

```

• df_poll = CSV.read(ros_datadir("Death",
• "polls.csv"), DataFrame)

```



```

begin
    f = Figure()
    ax = Axis(f[1, 1]; title="Death penalty
opinions", xlabel="Year",
ylabel="Percentage support for the
death penalty")
    scatter!(df_poll.year, df_poll.support
.* 100)
    err_lims = [100(sqrt(df_poll.support[i]*
(1-df_poll.support[i])/1000)) for i in
1:nrow(df_poll)]
    errorbars!(df_poll.year, df_poll.support
.* 100, err_lims, color = :red)
    f
end

```

Used in later notebooks.

| | STATE | TOTLDF | DOR | DORAVG | HRS |
|----|--------|--------|-------|--------|-------|
| 1 | "AL" | 296.0 | 33.47 | 32.65 | 11.61 |
| 2 | "AR" | 77.0 | 15.4 | 15.65 | 9.7 |
| 3 | "AZ" | 231.0 | 41.5 | 39.42 | 7.92 |
| 4 | "CA" | 528.0 | 9.21 | 9.14 | 8.8 |
| 5 | "FL" | 851.0 | 30.19 | 30.18 | 10.91 |
| 6 | "GA" | 323.0 | 19.63 | 19.12 | 12.78 |
| 7 | "ID" | 31.0 | 48.48 | 44.16 | 3.55 |
| 8 | "IL" | 238.0 | 11.26 | 10.98 | 8.18 |
| 9 | "IN" | 79.0 | 11.81 | 10.93 | 5.61 |
| 10 | "KY" | 59.0 | 10.67 | 10.24 | 7.03 |
| | : more | | | | |
| 26 | "WY" | 5.0 | 9.98 | 11.63 | 4.58 |

```

• begin
•   death_raw=CSV.read(ros_datadir("Death",
•     "dataforandy.csv"), DataFrame;
•     missingstring="NA")
•   death =
•     death_raw[completecases(death_raw), :]
end

```

```

• let
•   st_abbr = death[:, 1]
•   ex_rate = death[:, 8] ./ 100
•   err_rate = death[:, 7] ./ 100
•   hom_rate = death[:, 5] ./ 100000
•   ds_per_homicide = death[:, 3] ./ 1000
•   ds = death[:, 2]
•   hom = ds ./ ds_per_homicide
•   ex = ex_rate .* ds
•   err = err_rate .* ds
•   pop = hom ./ hom_rate
•   std_err_rate = sqrt.( (err .+ 1) .* (ds
•     .+ 1 .- err) ./ ((ds .+ 2).^2 .* (ds .+
•     3)) )
• end;

```