

See chapter 9 in Regression and Other Stories.

.....

```
• html"""
• <style>
•   main {
•     margin: 0 auto;
•     max-width: 2000px;
•     padding-left: max(160px, 10%);
•     padding-right: max(160px, 10%);
•   }
• </style>
• """
•
```

```
• using Pkg ✓ , DrWatson ✓
```

```
• begin
•   using GLM ✓
•
•   # Specific to ROSStanPluto
•   using StanSample ✓
•
•   # Graphics related
•   using GLMakie ✓
•
•   # Common data files and functions
•   using RegressionAndOtherStories ✓
• end
```

```
Replacing docs for `RegressionAndOtherStories.tr
DataFrame, AbstractString}` in module `Regressio
```

9.1 Propagating uncertainty in inference using posterior simulations.

```
hibbs =
```

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
: more				
16	2012	0.95	52.0	"Obama"


```
• hibbs =  
  CSV.read(ros_datadir("ElectionsEconomy",  
    "hibbs.csv"), DataFrame)
```

```
hibbs_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM,  
vote ~ 1 + growth
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	46.2476	1.62193	28.51	<1e-31
growth	3.06053	0.696274	4.40	0.0001

```
• hibbs_lm = lm(@formula(vote ~ growth),  
  hibbs)
```



```

let
  fig = Figure()
  hibbs.label = string.(hibbs.year)
  xlabel = "Average growth personal
income [%]"
  ylabel = "Incumbent's party vote share"
  let
    title = "Forecasting the election
from the economy"
    ax = Axis(fig[1, 1]; title, xlabel,
ylabel)
    for (ind, yr) in
      enumerate(hibbs.year)
        annotations!("$ (yr)"; position=
(hibbs.growth[ind],
hibbs.vote[ind]), fontsize=10)
    end
  end
  let
    x = LinRange(-1, 4, 100)
    title = "Data and linear fit"
    ax = Axis(fig[1, 2]; title, xlabel,
ylabel)
    scatter!(hibbs.growth, hibbs.vote)
    lines!(x, coef(hibbs_lm)[1] .+
coef(hibbs_lm)[2] .* x;
color=:darkred)
    annotations!("vote = 46.2 + 3.0 *
growth"; position=(0, 41))
  end
  fig
end

```

```

• stan7_1 = "
• data {
•   int<lower=1> N;      // total number of
•   observations
•   vector[N] growth;   // Independent
•   variable: growth
•   vector[N] vote;     // Dependent
•   variable: votes
• }
• parameters {
•   real b;             // Coefficient
•   independent variable
•   real a;             // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(50, 20);
•   b ~ normal(2, 10);
•   sigma ~ exponential(1);
•
•   mu = a + b * growth;
•
•   // likelihood including constants
•   vote ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std	
1	"b"	3.01923	0.016644	0.671776	1
2	"a"	46.3485	0.0384014	1.57594	4
3	"sigma"	3.60832	0.0133485	0.621676	2

```

• let
•   data = (N=nrow(hibbs), vote=hibbs.vote,
•   growth=hibbs.growth)
•   global m7_1s = SampleModel("hibbs",
•   stan7_1)
•   global rc7_1s = stan_sample(m7_1s; data)
•   success(rc7_1s) && describe(m7_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/1
d.

```

2169.02

```

• let
•   ss7_1s = describe(m7_1s)
•   ss7_1s[:sigma, :ess]
• end

```

2169.02

```

• let
•   ss7_1s = describe(m7_1s; showall=true)
•   ss7_1s[:sigma, :ess]
• end


```

	parameters	mean	mcse	
1	"lp__"	-31.9347	0.0343773	1.30
2	"accept_stat__"	0.913995	0.00182292	0.10
3	"stepsize__"	0.418669	0.00720391	0.00
4	"treedepth__"	2.587	0.011743	0.60
5	"n_leapfrog__"	7.42	0.0637245	3.90
6	"divergent__"	0.0	NaN	0.00
7	"energy__"	33.4318	0.0484177	1.70
8	"b"	3.01923	0.016644	0.60
9	"a"	46.3485	0.0384014	1.50
10	"sigma"	3.60832	0.0133485	0.60

```
• describe(m7_1s; showall=true)
```

post7_1s =	b	a	sigma
1	2.4296	46.9814	3.79408
2	2.58025	47.5254	4.10587
3	2.30726	47.2282	4.65373
4	3.07832	47.1264	3.61459
5	2.69067	45.5456	3.6129
6	2.58458	46.0218	3.85655
7	3.09983	45.7698	3.98921
8	2.94073	46.6183	4.04931
9	2.00617	49.353	4.9111
10	1.18727	48.4547	5.04701
⋮ more			
4000	3.28506	46.152	3.04823

```
• post7_1s = success(rc7_1s) &&
  read_samples(m7_1s, :dataframe)
```



```
• trankplot(post7_1s, "b")
```

```
ms7_1s =
```

	parameters	median	mad_sd	mean	st
1	"a"	46.32	1.486	46.348	1.57
2	"b"	3.023	0.654	3.019	0.67
3	"sigma"	3.531	0.602	3.608	0.62

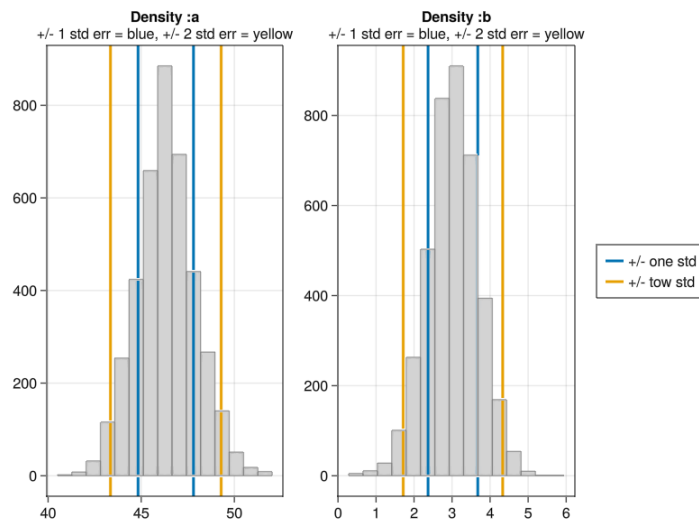
```
• ms7_1s = model_summary(post7_1s, [:a, :b, :sigma])
```

```
sims = 4000×3 Matrix{Float64}:
  2.4296  46.9814  3.79408
  2.58025 47.5254  4.10587
  2.30726 47.2282  4.65373
  3.07832 47.1264  3.61459
  2.69067 45.5456  3.6129
  2.58458 46.0218  3.85655
  3.09983 45.7698  3.98921
  ⋮
  1.76017 48.5191  3.48972
  3.07896 46.1482  3.34962
  3.36978 45.2588  3.34669
  3.21821 45.3887  2.78174
  2.85801 46.2397  4.31696
  3.28506 46.152  3.04823
```

```
• sims = Array(post7_1s)
```

```
1x3 Matrix{Float64}:  
 3.02268 46.3198 3.53123
```

- `median(sims; dims=1)`




```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Density :a",
•           subtitle="+/- 1 std err = blue, +/- 2
•           std err = yellow")
•   hist!(post7_1s.a; bins=15, color =
•           :lightgrey, strokewidth = 1, strokecolor
•           = :grey)
•   one = vlines!([ms7_1s[:a, :median] -
•   ms7_1s[:a, :mad_sd], ms7_1s[:a,
•   :median] + ms7_1s[:a, :mad_sd]];
•   linewidth=3)
•   two = vlines!([ms7_1s[:a, :median] -
•   2ms7_1s[:a, :mad_sd], ms7_1s[:a,
•   :median] + 2ms7_1s[:a, :mad_sd]];
•   linewidth=3)

•   ax = Axis(f[1, 2]; title="Density :b",
•           subtitle="+/- 1 std err = blue, +/- 2
•           std err = yellow")
•   hist!(post7_1s.b; bins=15, color =
•           :lightgrey, strokewidth = 1, strokecolor
•           = :grey)
•   vlines!([ms7_1s[:b, :median] -
•   ms7_1s[:b, :mad_sd], ms7_1s[:b,
•   :median] + ms7_1s[:b, :mad_sd]];
•   linewidth=3)
•   vlines!([ms7_1s[:b, :median] -
•   2ms7_1s[:b, :mad_sd], ms7_1s[:b,
•   :median] + 2ms7_1s[:b, :mad_sd]];
•   linewidth=3)
•   Legend(f[1, 3], [one, two], ["+/- one
•   std", "+/- tow std"])
•   f
• end

```



```

• let
•   growth_range =
•     LinRange(minimum(hibbs.growth),
•       maximum(hibbs.growth), 200)
•   votes = mean.(link(post7_1s, (r,x) ->
•     r.a + x * r.b, growth_range))
•
•   xlabel = "Average growth personal
•     income [%]"
•   ylabel = "Incumbent's party vote share"
•
•   fig = Figure()
•
•   ax = Axis(fig[1, 1]; title="Plot the
•     mcmc draws for :a and :b", xlabel=":a",
•     ylabel=":b")
•   scatter!(post7_1s.a, post7_1s.b;
•     markersize=4)
•
•   xlabel = "Average growth personal
•     income [%]"
•   ylabel="Incumbent's party vote share"
•   ax = Axis(fig[1, 2]; title="Regression
•     line based on 4000 posterior samples",
•     subtitle = "(grey lines are based
•       on first 200 draws of :a and :b)",
•     xlabel, ylabel)
•   for i in 1:100
•     lines!(growth_range, post7_1s.a[i]
•       .+ post7_1s.b[i] .* growth_range,
•       color = :lightgrey)
•   end
•   scatter!(hibbs.growth, hibbs.vote)
•   lines!(growth_range, votes, color =
•     :red)
•   fig

```

end

9.2 Prediction and uncertainty.

	x	y
1	-2.0	50
2	-1.0	44
3	0.0	50
4	1.0	47
5	2.0	56

```
• let
•   x = LinRange(-2, 2, 5)
•   y = [50, 44, 50, 47, 56]
•   global sexratio = DataFrame(x = x, y =
•   y)
• end
```

```

• stan9_1 = "
• data {
•   int<lower=1> N; // total number of
•   observations
•   vector[N] x;    // Independent
•   variable: growth
•   vector[N] y;    // Dependent variable:
•   votes
• }
• parameters {
•   real b;          // Coefficient
•   independent variable
•   real a;          // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(50, 5);
•   b ~ normal(0, 5);
•   sigma ~ uniform(0, 10);
•
•   mu = a + b * x;
•
•   // likelihood including constants
•   y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std	
1	"b"	1.3147	0.0492384	1.70412	-1
2	"a"	49.439	0.036813	2.2286	45
3	"sigma"	5.44629	0.0921171	1.85281	2.

```

• let
•   data = (N = nrow(sexratio), x =
•   sexratio.x, y = sexratio.y)
•   global m9_1s = SampleModel("m9_1s",
•   stan9_1)
•   global rc9_1s = stan_sample(m9_1s; data)
•   success(rc9_1s) && describe(m9_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtttnvs8s2c00000gn/1
d.


```

	parameters	median	mad_sd	mean	st
1	"a"	49.429	1.945	49.439	2.22
2	"b"	1.364	1.501	1.315	1.70
3	"sigma"	5.152	1.966	5.446	1.85

```

• if success(rc9_1s)
•   post9_1s = read_samples(m9_1s,
•   :dataframe)
•   sm9_1s = model_summary(post9_1s, [:a,
•   :b, :sigma])
end

```



```

let
  x_range = LinRange(minimum(sexratio.x),
    maximum(sexratio.x), 200)
  y = mean.(link(post9_1s, (r,x) -> r.a +
    x * r.b, x_range))

  xlabel = "x"
  ylabel = "y"

  fig = Figure()

  ax = Axis(fig[1, 1]; title="Posterior
simulation under default prior",
  xlabel="Intercept, a", ylabel="Slope,
b")
  scatter!(post9_1s.a, post9_1s.b;
  markersize=4)

  ax = Axis(fig[1, 2]; title="Bayes
regression (4000 posterior samples)",
  subtitle = "(grey lines are based
on first 100 draws of a and b)",
  xlabel, ylabel)
  for i in 1:100
    lines!(x_range, post9_1s.a[i] .+
      post9_1s.b[i] .* x_range, color =
      :lightgrey)
  end
  scatter!(sexratio.x, sexratio.y)
  lines!(x_range, y, color = :red)
  fig
end

```

```

• stan9_2 = "
• data {
•   int<lower=1> N; // total number of
•   observations
•   vector[N] x;    // Independent
•   variable: growth
•   vector[N] y;    // Dependent variable:
•   votes
• }
• parameters {
•   real b;          // Coefficient
•   independent variable
•   real a;          // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(48.8, 0.2);
•   b ~ normal(0, 0.2);
•   sigma ~ uniform(0, 10);
•
•   mu = a + b * x;
•
•   // likelihood including constants
•   y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std
1	"b"	0.033602	0.00333889	0.199175
2	"a"	48.8097	0.0033849	0.202378
3	"sigma"	5.14479	0.0342823	1.68717

```

• let
•   data = (N = nrow(sexratio), x =
•   sexratio.x, y = sexratio.y)
•   global m9_2s = SampleModel("m9_2s",
•   stan9_2)
•   global rc9_2s = stan_sample(m9_2s; data)
•   success(rc9_2s) && describe(m9_2s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T
d.

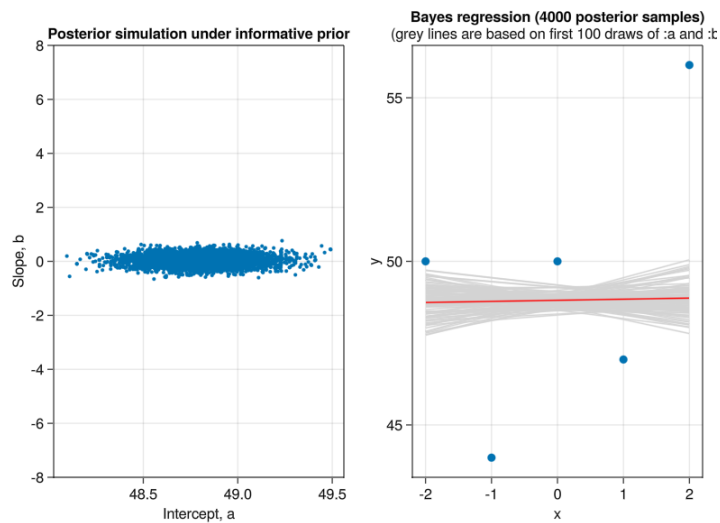
```

	parameters	median	mad_sd	mean	std
1	"a"	48.808	0.203	48.81	0.20
2	"b"	0.029	0.194	0.034	0.19
3	"sigma"	4.844	1.623	5.145	1.68

```

• if success(rc9_2s)
•   post9_2s = read_samples(m9_2s,
•   :dataframe)
•   sm9_2s = model_summary(post9_2s, [:a,
•   :b, :sigma])
end

```

```

let
  x_range = LinRange(minimum(sexratio.x),
    maximum(sexratio.x), 200)
  y = mean.(link(post9_2s, (r,x) -> r.a +
    x * r.b, x_range))

  xlabel = "x"
  ylabel = "y"

  fig = Figure()

  ax = Axis(fig[1, 1]; title="Posterior
simulation under informative prior",
  xlabel="Intercept, a", ylabel="Slope,
  b")
  ylims!(ax, -8, 8)
  scatter!(post9_2s.a, post9_2s.b;
  markersize=4)

  ax = Axis(fig[1, 2]; title="Bayes
regression (4000 posterior samples)",
  subtitle = "(grey lines are based
on first 100 draws of :a and :b)",
  xlabel, ylabel)
  for i in 1:100
    lines!(x_range, post9_2s.a[i] .+
      post9_2s.b[i] .* x_range, color =
      :lightgrey)
  end
  scatter!(sexratio.x, sexratio.y)
  lines!(x_range, y, color = :red)
  fig
end

```

9.3 Prior information and Bayesian synthesis.

Prior based on a previously-fitted model using economic and political condition.

```
• begin
•   theta_hat_prior = 0.524
•   se_prior = 0.041
• end;
```

Survey of 400 people, of whom 190 say they will vote for the Democratic candidate.

```
• begin
•   n = 400
•   y = 190
• end;
```

Data estimate.

```
theta_hat_data = 0.475
```

```
• theta_hat_data = y/n
```

```
se_data = 0.02496873044429772
```

```
• se_data =  $\sqrt{((\underline{y/n}) * (1 - \underline{y/n}) / \underline{n})}$ 
```


Bayes estimate.

```
theta_hat_bayes = 0.48825635323153693
```

```
• theta_hat_bayes =
•   (theta_hat_prior/se_prior2 +
•    theta_hat_data/se_data2) / (1/se_prior2
•    + 1/se_data2)
```

```
se_bayes = 0.02132543263776925
```

```
• se_bayes =  $\sqrt{1 / (1 / \underline{se\_prior}^2 + 1 / \underline{se\_data}^2)}$ 
```




```

• let
•   x = 0.3:0.001:0.7
•   f = Figure()
•   ax = Axis(f[1, 1], title="Prior,
•   likelihood & posterior")
•   prior = lines!(f[1, 1], x, pdf.
•   (Normal(theta_hat_prior, se_prior), x),
•   color=:gray)
•   data = lines!(x, pdf.
•   (Normal(theta_hat_data, se_data),
•   x),color=:darkred)
•   bayes = lines!(x, pdf.
•   (Normal(theta_hat_bayes, se_bayes), x),
•   color=:darkblue)
•   Legend(f[1, 2], [prior, data, bayes],
•   ["Prior", "Data", "Bayesian"])

•   current_figure()
end

```



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1], title="Prior,
•   likelihood & posterior (using
•   `density()`)" )
•   density!(rand(Normal(theta_hat_prior,
•   se_prior), Int(1e6)), lab="prior")
•   density!(rand(Normal(theta_hat_data,
•   se_data), Int(1e6)), lab="likelihood")
•   density!(rand(Normal(theta_hat_bayes,
•   se_bayes), Int(1e6)), lab="bayes")
•   current_figure()
• end

```

9.4 Example of Bayesian inference: beauty and sex ratio.