

In Regression and Other Stories, mcmc is *just* a tool. Hence whether one uses Stan or Turing is not the main focus of the book.

This notebook uses

ElectionsEconomy: hibbs.csv to illustrate how Stan and other tools are used in the Julia *project* ROSStanPluto.jl.

Over time I will expand below the list of topics:

1. Stan (StanSample.jl, ...)
2. Using median and mad to summarize a posterior distribution.
3. ...
4. Model comparison (TBD)
5. DAGs (TBD)
6. Graphs (TBD)
7. ...

See Chapter 1.2, Figure 1.1 in Regression and Other Stories.

Widen the cells.

```

• html"""
• <style>
•     main {
•         margin: 0 auto;
•         max-width: 2000px;
•         padding-left: max(160px, 10%);
•         padding-right: max(160px, 10%);
•     }
• </style>
• """

```

A typical set of Julia packages to include in notebooks.

```

• using Pkg ✓ , DrWatson ✓

```

```

• begin
•     # Specific to this notebook
•     using GLM ✓
•
•     # Specific to ROSStanPluto
•     using StanSample ✓
•
•     # Graphics related
•     using GLMakie ✓
•
•     # Include basic packages
•     using RegressionAndOtherStories ✓
• end

```

```

Replacing docs for `RegressionAndOtherStories.tr
DataFrame, AbstractString}` in module `Regressio

```

Note

All data files are available (as .csv files) in the data subdirectory of package RegressionAndOtherStories.jl.

```

"/Users/rob/.julia/packages/RegressionAndOtherSto
• ros_datadir()

```

```
hibbs =
```

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
: more				
16	2012	0.95	52.0	"Obama"

```
• hibbs =  
  CSV.read(ros_datadir("ElectionsEconomy",  
    "hibbs.csv"), DataFrame)
```

```
hibbs_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM},  
  vote ~ 1 + growth
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	46.2476	1.62193	28.51	<1e-308
growth	3.06053	0.696274	4.40	0.00011

```
• hibbs_lm = lm(@formula(vote ~ growth),  
  hibbs)
```

```
► [-8.99292, 2.66743, 1.0609, 2.20753, -5.89044, 4.27444]
```

```
• residuals(hibbs_lm)
```

```
2.2744434224582912
```

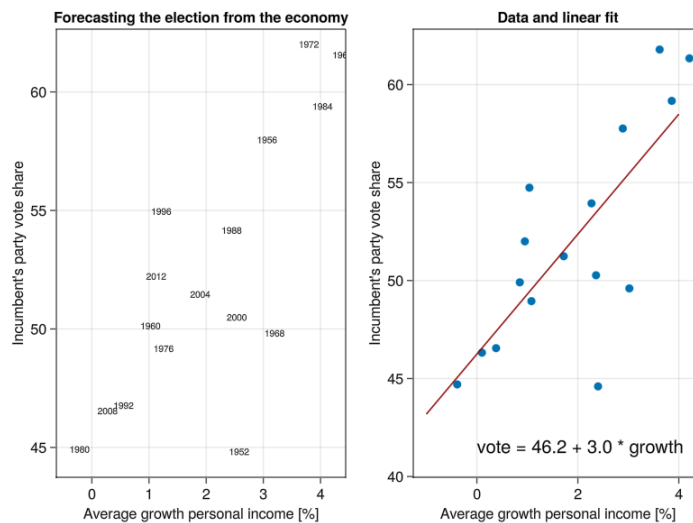
```
• mad(residuals(hibbs_lm))
```

```
3.635681268522063
```

```
• std(residuals(hibbs_lm))
```

```
► [46.2476, 3.06053]
```

```
• coef(hibbs_lm)
```



```

let
  fig = Figure()
  hibbs.label = string.(hibbs.year)
  xlabel = "Average growth personal
income [%]"
  ylabel = "Incumbent's party vote share"
  let
    title = "Forecasting the election
from the economy"
    ax = Axis(fig[1, 1]; title, xlabel,
ylabel)
    for (ind, yr) in
      enumerate(hibbs.year)
        annotations!("$ (yr)"; position=
(hibbs.growth[ind],
hibbs.vote[ind]), textsize=10)
    end
  end
  let
    x = LinRange(-1, 4, 100)
    title = "Data and linear fit"
    ax = Axis(fig[1, 2]; title, xlabel,
ylabel)
    scatter!(hibbs.growth, hibbs.vote)
    lines!(x, coef(hibbs_lm)[1] .+
coef(hibbs_lm)[2] .* x;
color=:darkred)
    annotations!("vote = 46.2 + 3.0 *
growth"; position=(0, 41))
  end
  fig
end

```

Priors used in the Stan model.

```

• stan1_0 = "
• parameters {
•   real b;           // Coefficient
•   independent variable
•   real a;           // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   // priors including constants
•   a ~ normal(50, 20);
•   b ~ normal(2, 10);
•   sigma ~ exponential(1);
• };

```

	parameters	mean	mcse	std	
1	"b"	1.65098	0.177134	10.0993	-1
2	"a"	49.9506	0.356125	20.1822	16
3	"sigma"	1.01953	0.0163088	1.01051	0.

```

• begin
•   m1_0s = SampleModel("hibbs", stan1_0)
•   rc1_0s = stan_sample(m1_0s)
•   success(rc1_0s) && describe(m1_0s)
• end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
d.

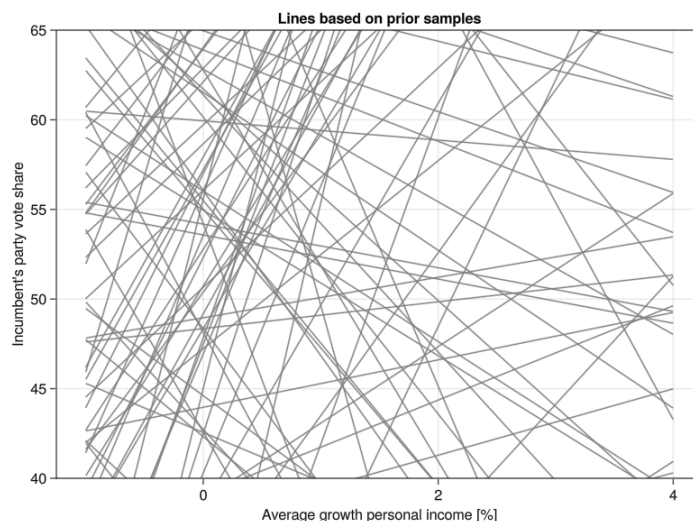
```

	parameters	median	mad_sd	mean	std
1	"a"	49.87	20.042	49.951	20.1
2	"b"	1.477	10.128	1.651	10.0
3	"sigma"	0.71	0.726	1.02	1.01

```

• begin
•   post1_0s = read_samples(m1_0s,
•     :dataframe)
•   ms1_0s = model_summary(post1_0s, [:a,
•     :b, :sigma])
end


```



```

• let
•   fig = Figure()
•   xlabel = "Average growth personal
•     income [%]"
•   ylabel="Incumbent's party vote share"
•   ax = Axis(fig[1, 1]; title="Lines based
•     on prior samples",
•     xlabel, ylabel)
•   ylims!(ax, 40, 65)
•   xrange = LinRange(-1, 4, 200)
•   for i = 1:100
•     lines!(xrange, post1_0s.a[i] .+
•       post1_0s.b[i] .* xrange, color =
•         :grey)
•   end
•   fig
end

```



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Density :a")
•   density!(f[1, 1], post1_0s.a)
•   ax = Axis(f[1, 2]; title="Density :b")
•   density!(f[1, 2], post1_0s.b)
•   ax = Axis(f[1, 3]; title="Density
•   :sigma")
•   density!(f[1, 3], post1_0s.sigma)
•   f
• end

```

Conditioning based on the available data.


```

• stan1_1 = "
• functions {
• }
• data {
•   int<lower=1> N;      // total number of
•   observations
•   vector[N] growth;   // Independent
•   variable: growth
•   vector[N] vote;     // Dependent
•   variable: votes
• }
• parameters {
•   real b;             // Coefficient
•   independent variable
•   real a;             // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•   mu = a + b * growth;
•
•   // priors including constants
•   a ~ normal(50, 20);
•   b ~ normal(2, 10);
•   sigma ~ exponential(1);
•
•   // likelihood including constants
•   vote ~ normal(mu, sigma);
• };

```

	parameters	mean	mcse	std	5%
1	"a"	46.27	0.04	1.54	43.7
2	"b"	3.06	0.02	0.67	1.98
3	"sigma"	3.6	0.02	0.66	2.7

```

• let
•   data = (N=16, vote=hibbs.vote,
•   growth=hibbs.growth)
•   global m1_1s = SampleModel("hibbs",
•   stan1_1)
•   global rc1_1s = stan_sample(m1_1s; data)
•   success(rc1_1s) && describe(m1_1s, [:a,
•   :b, :sigma])
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/7
d.

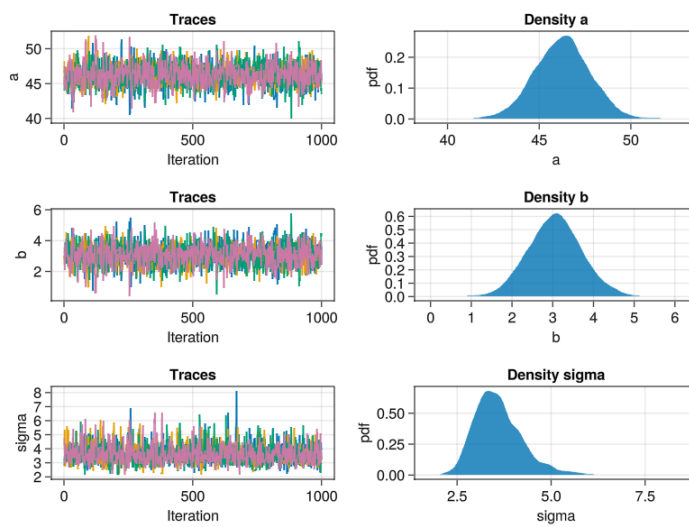
```

	parameters	median	mad_sd	mean	st
1	"a"	46.279	1.488	46.266	1.53
2	"b"	3.061	0.642	3.055	0.66
3	"sigma"	3.516	0.588	3.604	0.68


```

• if success(rc1_1s)
•   post1_1s = read_samples(m1_1s,
•   :dataframe)
•   ms1_1s = model_summary(post1_1s, [:a,
•   :b, :sigma])
end


```



```
• plot_chains(post1_1s, [:a, :b, :sigma])
```



```
• trankplot(post1_1s, "b")
```



```

let
    N = 100
    x = LinRange(-1, 4, N)
    a = rand(Normal(50, 20), N)
    b = rand(Normal(2, 10), N)
    mat1 = zeros(50, 100)
    for i in 1:50
        mat1[i, :] = a[i] .+ b[i] .* x
    end
    ā = ms1_1s[:a, :mean]
    b̄ = ms1_1s[:b, :mean]

    # Maybe could use a 'link' function here
    mat2 = zeros(50, 100)
    for i in 1:50
        mat2[i, :] = post1_1s.a[i] .+
            post1_1s.b[i] .* x
    end

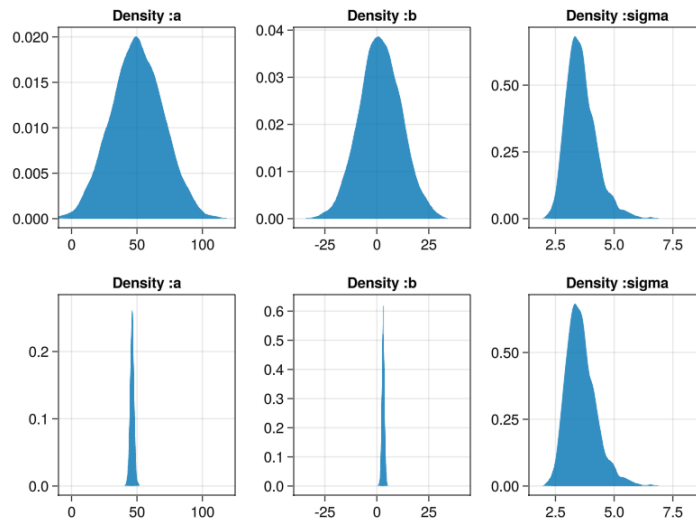
    fig = Figure()
    xlabel = "Average growth personal
income [%]"
    ylabel="Incumbent's party vote share"
    ax = Axis(fig[1, 1]; title="Lines based
on prior samples",
        xlabel, ylabel)
    ylims!(ax, 40, 65)
    series!(fig[1, 1], x, mat1,
        solid_color=:lightgrey)
    ax = Axis(fig[1, 2]; title="Lines based
on posterior samples",
        xlabel, ylabel)
    series!(fig[1, 2], x, mat2,
        solid_color=:lightgrey)
    scatter!(hibbs.growth, hibbs.vote)

```


```
lines!(fig[1, 2], x,  $\bar{a}$  .+  $\bar{b}$  * x, color
= :red)
```

```
fig
```

```
end
```



```
• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Density :a")
•   xlims!(ax, -10, 125)
•   density!(post1_0s.a)
•   ax = Axis(f[1, 2]; title="Density :b")
•   xlims!(ax, -40, 45)
•   density!(post1_0s.b)
•   ax = Axis(f[1, 3]; title="Density
•   :sigma")
•   density!(post1_1s.sigma)
•
•   ax = Axis(f[2, 1]; title="Density :a")
•   density!(post1_1s.a)
•   xlims!(ax, -10, 125)
•   ax = Axis(f[2, 2]; title="Density :b")
•   xlims!(ax, -40, 45)
•   density!(post1_1s.b)
•   ax = Axis(f[2, 3]; title="Density
•   :sigma")
•   density!(post1_1s.sigma)
•   f
end
```



```

• begin
•     fig = Figure()
•     hibbs.label = string.(hibbs.year)
•     xlabel = "Average growth personal
•     income [%]"
•     ylabel="Incumbent's party vote share"
•
•     # Same figure as above
•     let
•         title = "Forecasting the election
•         from the economy"
•         ax = Axis(fig[1, 1]; title, xlabel,
•         ylabel)
•         xlims!(ax, -0.5, 5)
•         for (ind, yr) in
•         enumerate(hibbs.year)
•             annotations!("$ (yr)"; position=
•             (hibbs.growth[ind],
•             hibbs.vote[ind]), fontsize=12)
•         end
•     end
•
•     # Superimpose Stan fit
•     let
•         ā = ms1_1s[:a, :mean]
•         b̄ = ms1_1s[:b, :mean]
•         title = "Compare GLM and Stan
•         fitted lines"
•         axis = (; title, xlabel, ylabel)
•
•         x = LinRange(-1, 4.4, 100)
•         title = "Data and linear fit"
•         ax = Axis(fig[1, 2]; title, xlabel,
•         ylabel)
•         xlims!(ax, -0.5, 5)
•         scatter!(hibbs.growth, hibbs.vote)


```

```

lines!(x, coef(hibbs_lm)[1] .+
coef(hibbs_lm)[2] .* x)
lines!(x,  $\bar{a}$  .+  $\bar{b}$  .* x;
color=:darkred)
annotations!("vote = $(round( $\bar{a}$ ,
digits=1)) + $(round( $\bar{b}$ , digits=0)) \
* growth"; position=(0, 41),
textsize=16)

end
fig
end

```



```

• let
•   N = 10000
•   nt = (
•       a = post1_1s.a,
•       b = post1_1s.b,
•        $\sigma$  = post1_1s.sigma,
•   )
•
•   fig = Figure()
•   for (i, k) in enumerate(keys(nt))
•       ax = Axis(fig[1, i]; title =
•           "Density $k")
•       den = density!(nt[k])
•   end
•   fig
end

```

Compute median and mad.

Alternative computation of mad().

```
► [1.48804, 0.642436, 0.587705]
```

```
• let
•   1.483 .* [
•     median(abs.(post1_1s.a .-
•     median(post1_1s.a))),
•     median(abs.(post1_1s.b .-
•     median(post1_1s.b))),
•     median(abs.(post1_1s.sigma .-
•     median(post1_1s.sigma)))]
•
end
```

```
ms1_1 =
```

	parameters	median	mad_sd	mean	std
1	"a"	46.279	1.488	46.266	1.53
2	"b"	3.061	0.642	3.055	0.66
3	"sigma"	3.516	0.588	3.604	0.65

```
• ms1_1 = model_summary(post1_1s, ["a", "b",
"sigma"])
```

```
0.642
```

```
• ms1_1[:, :mad_sd]
```

```
ss1_1 =
```

	parameters	mean	mcse	std	5%
1	"a"	46.27	0.04	1.54	43.7
2	"b"	3.06	0.02	0.67	1.98
3	"sigma"	3.6	0.02	0.66	2.7

```
• ss1_1 = describe(m1_1s, ["a", "b",
"sigma"]; digits=2)
```

```
1490.84
```

```
• ss1_1["a", "ess"]
```


Quick simulation with median, mad, mean and std of Normal observations.

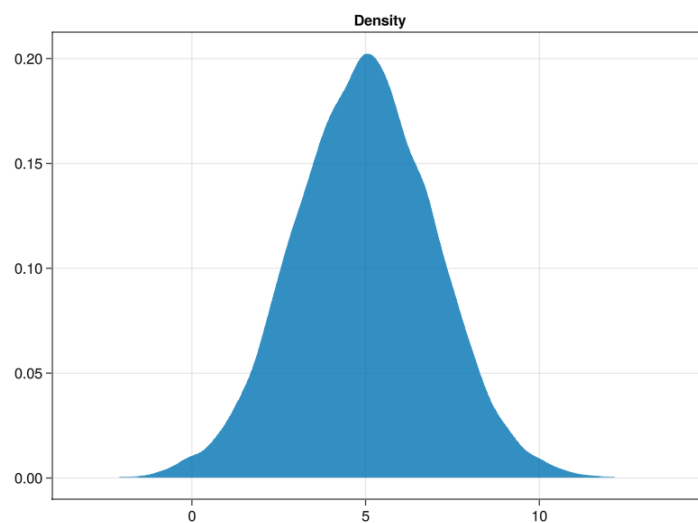
```
nt =  
► (x = [3.99721, 6.45366, 4.15654, 4.08856, 0.1084  
• nt = (x=rand(Normal(5, 2), 10000),)
```

```
► [4.98422, 2.00871, 4.9748, 1.99798]  
• [median(nt.x), mad(nt.x), mean(nt.x),  
std(nt.x)]
```

```
sd_mean = 0.02  
• sd_mean = round(mad(nt.x)/√10000; digits=2)
```

```
1.3548569833911626  
• median(abs.(nt.x .- median(nt.x)))
```

```
2.009252906369094  
• 1.483 * median(abs.(nt.x .- median(nt.x)))
```



```
• let  
• fig = Figure()  
• ax = Axis(fig[1, 1]; title = "Density")  
• den = density!(nt.x)  
• fig  
• end
```

```
► [1.08543, 8.91972]  
• quantile(nt.x, [0.025, 0.975])
```

```
► [3.61705, 6.32982]
```

```
• quantile(nt.x, [0.25, 0.75])
```

A closer look at Stan's summary. Below the full version:

	parameters	mean	mcse	
1	"lp__"	-31.9627	0.037871	1.3
2	"accept_stat__"	0.917142	0.00184506	0.1
3	"stepsize__"	0.41254	0.0281173	0.0
4	"treedepth__"	2.647	0.0165177	0.6
5	"n_leapfrog__"	7.766	0.0928592	4.1
6	"divergent__"	0.0	NaN	0.0
7	"energy__"	33.4793	0.049884	1.8
8	"b"	3.05516	0.0169224	0.6
9	"a"	46.2662	0.039804	1.5
10	"sigma"	3.6036	0.015569	0.6

```
• success(rc1_1s) && describe(m1_1s;  
  showall=true)
```

Usually I use the abbreviated version:

```
ss1_1s =
```

	parameters	mean	mcse	std	5%
1	"b"	3.06	0.02	0.67	1.98
2	"a"	46.27	0.04	1.54	43.7
3	"sigma"	3.6	0.02	0.66	2.7

```
• ss1_1s = success(rc1_1s) && describe(m1_1s,  
  names(post1_1s))
```

	parameters	median	mad_sd	mean	std
1	"a"	46.279	1.488	46.266	1.53
2	"b"	3.061	0.642	3.055	0.66
3	"sigma"	3.516	0.588	3.604	0.65

- `ms1_1s`

1490.84

- `ss1_1s[:a, :ess]`

1.488

- `ms1_1s[:a, :mad_sd]`

Experimental use of BridgeStan.

```

• bernoulli_model = "
• data {
•   int<lower=1> N;
•   int<lower=0,upper=1> y[N];
• }
• parameters {
•   real<lower=0,upper=1> theta;
• }
• model {
•   theta ~ beta(1,1);
•   y ~ bernoulli(theta);
• }
• ";

```

	parameters	mean	mcse	std
1	:lp__	-8.18169	0.0173597	0.752
2	:accept_stat__	0.898038	0.00225164	0.151
3	:stepsize__	1.10995	0.0378972	0.053
4	:treedepth__	1.38625	0.00828256	0.486
5	:n_leapfrog__	2.27	0.0174635	0.972
6	:divergent__	0.0	NaN	0.0
7	:energy__	8.6847	0.0243358	1.041
8	:theta	0.329246	0.0035045	0.131

```

• begin
•   data = Dict("N" => 10, "y" => [0, 1, 0,
•     1, 0, 0, 0, 0, 0, 1])
•   sm = SampleModel("bernoulli",
•     bernoulli_model)
•   rc = stan_sample(sm; data)
•   success(rc) && read_summary(sm)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
dated.

```

```

st =
StanSample.StanTable{Matrix{Float64}} with 4000 rows:
:theta Float64

```

```

• st = success(rc) && read_samples(sm)

```

```

bernoulli_lib =
"/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/

```

```

• bernoulli_lib = joinpath(sm.tmpdir,
  "bernoulli_model.so")

```

	x	q	log_density	gradient
1	0.883837	2.02928	-15.4398	-5.83837

```

• if isfile(bernoulli_lib)
•     blib = Libc.Libdl.dlopen(bernoulli_lib)
•
•     bernoulli_data = joinpath(sm.tmpdir,
• "bernoulli_data_1.json")
•     smb = StanModel(blib, bernoulli_data)
•     x = rand(smb.dims)
•     q = @. log(x / (1 - x))          #
•     unconstrained scale
•     log_density_gradient!(smb, q, jacobian
• = 0)
•     DataFrame(x=x, q=q,
• log_density=smb.log_density,
• gradient=smb.gradient)
• else
•     @info "Shared library
• `bernoulli_model.so` has not been created."
•     @info "Maybe BridgeStan has not been
• installed in $(ENV["CMDSTAN"])?"
• end

```

sim (generic function with 2 methods)

```


• function sim(smb::StanModel, x=0.1:0.1:0.9)
•     y = zeros(length(x))
•     q = zeros(length(x))
•     ld = zeros(length(x))
•     g = zeros(length(x))
•     for (i, p) in enumerate(x)
•         y[i] = p
•         q[i] = @. log(p / (1 - p))          #
•         unconstrained scale
•         log_density_gradient!(smb, q[i],
• jacobian = 0)
•         ld[i] = smb.log_density[1]
•         g[i] = smb.gradient[1]
•     end
•     return DataFrame(x=x, q=q,
• log_density=ld, gradient=g)
• end

```

```
sim_df =
```

	x	q	log_density	gradient
1	0.1	-2.19722	-7.64528	2.0
2	0.2	-1.38629	-6.39032	1.0
3	0.3	-0.847298	-6.10864	0.0
4	0.4	-0.405465	-6.32465	-1.0
5	0.5	0.0	-6.93147	-2.0
6	0.6	0.405465	-7.94651	-3.0
7	0.7	0.847298	-9.49783	-4.0
8	0.8	1.38629	-11.9355	-5.0
9	0.9	2.19722	-16.4342	-6.0

```
• sim_df = sim(smb)
```



```
• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="BridgeStan
•   example")
•   dens = lines!(sim_df.x,
•   sim_df.log_density)
•   gra = lines!(sim_df.x, sim_df.gradient)
•   Legend(f[1, 2], [dens, gra],
•   ["log_density", "gradient"])
•   f
end
```

