

See chapter 1 in Regression and Other Stories.

.....

Widen the cells.

```
• html"""  
• <style>  
•   main {  
•     margin: 0 auto;  
•     max-width: 2000px;  
•     padding-left: max(160px, 10%);  
•     padding-right: max(160px, 10%);  
•   }  
• </style>  
• """
```

A typical set of Julia packages to include in notebooks.

```
• using Pkg ✓ , DrWatson ✓
```

```
• begin  
•   # Specific to this notebook  
•   using GLM ✓  
•  
•   # Specific to ROSStanPluto  
•   using StanSample ✓  
•  
•   # Graphics related  
•   using CairoMakie ✓  
•   using AlgebraOfGraphics ✓  
•  
•   # Include basic packages  
•   using RegressionAndOtherStories ✓  
• end
```

1.1 The three challenges of statistics.

Note

It is not common for me to copy from the book but this particular section deserves an exception!

The three challenges of statistical inference are:

1. Generalizing from sample to population, a problem that is associated with survey sampling but actually arises in nearly every application of statistical inference;
2. Generalizing from treatment to control group, a problem that is associated with causal inference, which is implicitly or explicitly part of the interpretation of most regressions we have seen; and
3. Generalizing from observed measurements to the underlying constructs of interest, as most of the time our data do not record exactly what we would ideally like to study.

All three of these challenges can be framed as problems of prediction (for new people or new items that are not in the sample, future outcomes under different potentially assigned treatments, and underlying constructs of interest, if they could be measured exactly).

1.2 Why learn regression?

```
hibbs =
```

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
: more				
16	2012	0.95	52.0	"Obama"

```
• hibbs =  
  CSV.read(ros_datadir("ElectionsEconomy",  
    "hibbs.csv"), DataFrame)
```

```
hibbs_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM},  
  vote ~ 1 + growth
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	46.2476	1.62193	28.51	<1e-308
growth	3.06053	0.696274	4.40	0.00011

```
• hibbs_lm = lm(@formula(vote ~ growth),  
  hibbs)
```

```
► [-8.99292, 2.66743, 1.0609, 2.20753, -5.89044, 4.27444]
```

```
• residuals(hibbs_lm)
```

```
2.2744434224582912
```


```
• mad(residuals(hibbs_lm))
```

```
3.635681268522063
```

```
• std(residuals(hibbs_lm))
```

```
► [46.2476, 3.06053]
```

```
• coef(hibbs_lm)
```



```

let
  fig = Figure()
  hibbs.label = string.(hibbs.year)
  xlabel = "Average growth personal
income [%]"
  ylabel = "Incumbent's party vote share"
  let
    title = "Forecasting the election
from the economy"
    ax = Axis(fig[1, 1]; title, xlabel,
ylabel)
    for (ind, yr) in
      enumerate(hibbs.year)
        annotations!("$ (yr)"; position=
(hibbs.growth[ind],
hibbs.vote[ind]), textsize=10)
    end
  end
  let
    x = LinRange(-1, 4, 100)
    title = "Data and linear fit"
    ax = Axis(fig[1, 2]; title, xlabel,
ylabel)
    scatter!(hibbs.growth, hibbs.vote)
    lines!(x, coef(hibbs_lm)[1] .+
coef(hibbs_lm)[2] .* x;
color=:darkred)
    annotations!("vote = 46.2 + 3.0 *
growth"; position=(0, 41))
  end
  fig
end

```

1.3 Some examples of regression.

Electric company

	post_test	pre_test	grade	treatment
1	48.9	13.8	1	1
2	70.5	16.5	1	1
3	89.7	18.5	1	1
4	44.2	8.8	1	1
5	77.5	15.3	1	1
6	84.7	15.0	1	1
7	78.9	19.4	1	1
8	86.8	15.0	1	1
9	60.8	11.8	1	1
10	75.7	16.4	1	1
⋮ more				
192	110.0	102.6	4	0

```
• begin
•   electric =
•   CSV.read(ros_datadir("ElectricCompany",
•   "electric.csv"), DataFrame)
•   electric = electric[:, [:post_test,
•   :pre_test, :grade, :treatment]]
•   electric.grade =
•   categorical(electric.grade)
•   electric.treatment =
•   categorical(electric.treatment)
•   electric
• end
```

A quick look at the overall values of pre_test and post_test.


	variable	mean	min	median	max
1	:post_test	97.1495	44.2	102.3	122.0
2	:pre_test	72.2245	8.8	80.75	119.8
3	:grade	nothing	1	nothing	4
4	:treatment	nothing	0	nothing	1

```
• describe(electric)
```


true

```
• all(completeness(electric)) == true
```

**Post-test density for each grade
conditioned on treatment.**



```
• let
•   f = Figure()
•   axis = (; width = 150, height = 150)
•   el = data(electric) *
•   mapping(:post_test, col=:grade,
•   color=:treatment)
•   plt = el *
•   AlgebraOfGraphics.histogram(;bins=20) *
•   mapping(row=:treatment)
•   draw!(f[1, 1], plt; axis)
•   f
end
```




```

let
  f = Figure()
  axis = (; width = 150, height = 150)
  el = data(electric) *
  mapping(:post_test, col=:grade,
  color=:treatment)
  plt = el * AlgebraOfGraphics.density()
  * mapping(row=:treatment)
  draw!(f[1, 1], plt; axis)
f
end

```

Note


In above cell, as `density()` is exported by both GLMakie and AlgebraOfGraphics, it needs to be qualified.



```

let
  f = Figure()
  el = data(electric) *
  mapping(:post_test, col=:grade)
  plt = el * AlgebraOfGraphics.density()
  * mapping(color=:treatment)
  draw!(f[1, 1], plt)
f
end

```



```

let
  plt = data(electric) * visual(Violin) *
  mapping(:grade, :post_test,
  dodge=:treatment, color=:treatment)
  draw(plt)
end

```

Peacekeeping

peace =

	war	cfdate	faildate
1	"Afghanistan-Mujahideen"	8150	8257
2	"Afghanistan-Taliban"	8466	8505
3	"Algeria-FIS/AIS"	10149	12783
4	"Angola"	7820	8319
5	"Angola"	9089	10564
6	"Azerbaijan-N.K."	8643	8678
7	"Azerbaijan-N.K."	8901	12783
8	"Bangladesh-CHT"	8248	12783
9	"Myanmar-Karen"	8153	9282
10	"Myanmar-Karen"	9296	9907
: more			
96	"Yugoslavia-Kosovo"	10751	12783

```
• peace =  
  CSV.read(ros_datadir("PeaceKeeping",  
    "peacekeeping.csv"), missingstring="NA",  
    DataFrame)
```

	variable	mean	min
1	:war	nothing	"Afghanistan-Mujah:
2	:cfdate	8925.1	6985
3	:faildate	10795.8	7074
4	:peacekeepers	0.354167	0
5	:badness	-8.15228	-12.26
6	:delay	5.12177	0.04
7	:censored	0.416667	0

```
• describe(peace)
```

A quick look at this Dates stuff!

8150

- `peace.cfdate[1]`

1992-04-25T00:00:00

- `DateTime(1992, 4, 25)`

107 days

- `Date(1992, 8, 10) - Date(1992, 4, 25)`

1970-01-01

- `Date(1970,1,1)`

1992-04-25

- `Date(1970,1,1) + Dates.Day(8150)`

8150 days

- `Date(1992, 4, 25) - Date(1970, 1, 1)`

107

- `peace.faildate[1] - peace.cfdate[1]`

- `begin`
- `pks_df = peace[peace.peacekeepers .==`
- `1, [:cfdate, :faildate]]`
- `nopks_df = peace[peace.peacekeepers .==`
- `0, [:cfdate, :faildate]]`
- `end;`

0.4166666666666667

- `mean(peace.censored)`

64

- `length(unique(peace.war))`

0.5588235294117647

- `mean(peace[peace.peacekeepers .== 1,`
- `:censored])`

0.3387096774193548

- `mean(peace[peace.peacekeepers .== 0,`
- `:censored])`

1.382

- `mean(peace[peace.peacekeepers .== 1 .&& peace.censored .== 0, :delay])`

1.5153658536585364


- `mean(peace[peace.peacekeepers .== 0 .&& peace.censored .== 0, :delay])`

1.05

- `median(peace[peace.peacekeepers .== 1 .&& peace.censored .== 0, :delay])`

0.59

- `median(peace[peace.peacekeepers .== 0 .&& peace.censored .== 0, :delay])`



```

let
    f = Figure()
    pks = peace[peace.peacekeepers == 1
    .&& peace.censored == 0, :]
    nopks = peace[peace.peacekeepers == 0
    .&& peace.censored == 0, :]

    for i in 1:2
        title = i == 1 ? "Peacekeepers" :
            "No peacekeepers"

        ax = Axis(f[i, 1]; title,
            xlabel="Years until return to war",
            ylabel = "Frequency",
            yminorticksvisible = true,
            yminorgridvisible = true,
            yminorticks = IntervalsBetween(8))

        xlims!(ax, [0, 8])
        hist!(i == 1 ? pks.delay :
            nopks.delay)
    end
end
f
end

```

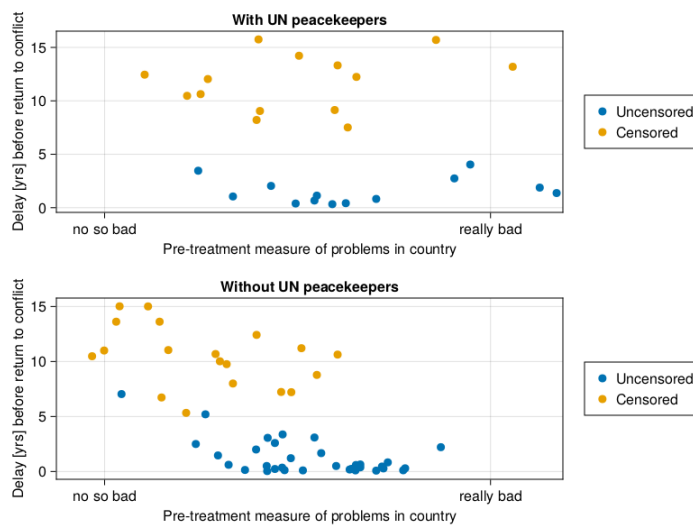
Note

Censored means conflict had not returned until end of observation period (2004).

```

• begin
•   # Filter out missing badness rows.
•   pb = peace[peace.badness .!= missing,
•   :];
•
•   # Delays until return to war for
•   uncensored, peacekeeper cases
•   pks_uc = pb[pb.peacekeepers .== 1 .&&
•   pb.censored .== 0, :delay]
•   # Delays until return to war for
•   censored, peacekeeper cases
•   pks_c = pb[pb.peacekeepers .== 1 .&&
•   pb.censored .== 1, :delay]
•
•   # No peacekeepr cases.
•   nopks_uc = pb[pb.peacekeepers .== 0 .&&
•   pb.censored .== 0, :delay]
•   nopks_c = pb[pb.peacekeepers .== 0 .&&
•   pb.censored .== 1, :delay]
•
•   # Crude measure (:badness) used for
•   assessing situation
•   badness_pks_uc = pb[pb.peacekeepers .==
•   1 .&& pb.censored .== 0,
•   :badness]
•   badness_pks_c = pb[pb.peacekeepers .== 1
•   .&& pb.censored .== 1,
•   :badness]
•   badness_nopks_uc = pb[pb.peacekeepers
•   .== 0 .&& pb.censored .== 0,
•   :badness]
•   badness_nopks_c = pb[pb.peacekeepers
•   .== 0 .&& pb.censored .== 1,
•   :badness]
• end;

```



```

• begin
•     f = Figure()
•     ax = Axis(f[1, 1], title = "With UN
•         peacekeepers",
•         xlabel = "Pre-treatment measure of
•             problems in country",
•         ylabel = "Delay [yrs] before return
•             to conflict")
•     sca1 = scatter!(badness_pks_uc, pks_uc)
•     sca2 = scatter!(badness_pks_c, pks_c)
•     xlims!(ax, [-13, -2.5])
•     Legend(f[1, 2], [sca1, sca2],
•         ["Uncensored", "Censored"])
•     ax.xticks = ([-12, -4], ["no so bad",
•         "really bad"])
•
•
•     ax = Axis(f[2, 1], title = "Without UN
•         peacekeepers",
•         xlabel = "Pre-treatment measure of
•             problems in country",
•         ylabel = "Delay [yrs] before return
•             to conflict")
•     sca1 = scatter!(badness_nopks_uc,
•         nopks_uc)
•     sca2 = scatter!(badness_nopks_c,
•         nopks_c)
•     xlims!(ax, [-13, -2.5])
•     Legend(f[2, 2], [sca1, sca2],
•         ["Uncensored", "Censored"])
•     ax.xticks = ([-12, -4], ["no so bad",
•         "really bad"])
•
•     f
• end

```

1.4 Challenges in building, understanding, and interpreting regression.

Simple causal

Note

In models like below I usually prefer to create 2 separate Stan Language models, one for the continuous case and another for the binary case. But they can be combined in a single model as shown below. I'm using this example to show one way to handle vectors returned from Stan's cmdstan.

```
• stan1_4_1 = "  
• data {  
•   int N;  
•   vector[N] x;  
•   vector[N] x_binary;  
•   vector[N] y;  
• }  
• parameters {  
•   vector[2] a;  
•   vector[2] b;  
•   vector<lower=0>[2] sigma;  
• }  
• model {  
•   // Priors  
•   a ~ normal(10, 10);  
•   b ~ normal(10, 10);  
•   sigma ~ exponential(1);  
•   // Likelihood  
•   y ~ normal(a[1] + b[1] * x, sigma[1]);  
•   y ~ normal(a[2] + b[2] * x_binary,  
•   sigma[2]);  
• }  
• ";
```


Note

Aki Vehtari did not include a seed number in his code.

```
• begin
•   Random.seed!(123)
•   n = 50
•   x = rand(Uniform(1, 5), n)
•   x_binary = [x[i] < 3 ? 0 : 1 for i in
•   1:n]
•   y = [rand(Normal(10 + 3x[i], 3), 1)[1]
•   for i in 1:n]
end;
```

	parameters	mean	mcse	std	5%
1	"a[1]"	9.4	0.029	1.4	7.1
2	"a[2]"	16.0	0.013	0.69	15.0
3	"b[1]"	3.2	0.009	0.44	2.5
4	"b[2]"	7.0	0.02	1.0	5.3
5	"sigma[1]"	3.5	0.0056	0.34	3.0
6	"sigma[2]"	3.7	0.0062	0.37	3.1

```
• let
•   data = (N = n, x = x, x_binary =
•   x_binary, y = y)
•   global m1_4_1s = SampleModel("m1_4_1s",
•   stan1_4_1);
•   global rc1_4_1s = stan_sample(m1_4_1s;
•   data)
•   success(rc1_4_1s) && describe(m1_4_1s)
end
```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/1
ted.

Note

This is a good point to take a quick look at Pluto cell metadata: the top left eye symbol and the top right 3-dots in a circle glyph (both only visible when the cursor is in the input cell). Both are used quite often in these notebooks. Try them out!

The output of above method of the function `model_summary(::SampleModel)`, called directly on a `SampleModel`, is different from method `model_summary(::DataFrame)`, typically used later on. Above table shows important mcmc diagnostic columns like `n_eff` and `r_hat`.

If Stan parameters are vectors (as in this example), `cmdstan` returns those using ' notation, e.g. `a.1`, `a.2`, ...

	parameters	median	mad_sd	mean	st
1	"a.1"	9.402	1.396	9.387	1.41
2	"a.2"	16.121	0.674	16.116	0.69
3	"b.1"	3.241	0.444	3.249	0.43
4	"b.2"	7.001	1.027	7.006	1.03
5	"sigma.1"	3.423	0.324	3.452	0.34
6	"sigma.2"	3.656	0.35	3.69	0.36

```
• if success(rc1_4_1s)
•   post1_4_1s = read_samples(m1_4_1s,
•     :dataframe)
•   model_summary(post1_4_1s,
•     names(post1_4_1s))
end
```

With vector parameters `read_samples()` can create a nested `DataFrame`:

```
nd1_4_1s =
```

	a	b
1	▶ [8.24802, 16.424]	▶ [3.68005, 7.67235]
2	▶ [8.83781, 16.203]	▶ [3.34579, 5.17041]
3	▶ [9.84559, 16.0788]	▶ [3.35379, 5.76383]
4	▶ [11.0378, 16.134]	▶ [3.05522, 6.06792]
5	▶ [10.8772, 15.807]	▶ [3.00888, 5.84699]
6	▶ [10.6881, 16.2392]	▶ [2.72581, 6.47463]
7	▶ [7.44081, 16.0745]	▶ [3.96654, 6.36506]
8	▶ [10.062, 15.7894]	▶ [2.99653, 7.15061]
9	▶ [9.95532, 15.5407]	▶ [3.10871, 6.93601]
10	▶ [11.3043, 16.2698]	▶ [2.75944, 5.98389]
⋮ more		
4000	▶ [10.5984, 17.5858]	▶ [2.70611, 4.93714]

```
• nd1_4_1s = read_samples(m1_4_1s,  
:nesteddataframe)
```

```
ms1_4_1s =
```

	parameters	median	mad_sd	mean	st
1	"a.1"	9.402	1.396	9.387	1.41
2	"a.2"	16.121	0.674	16.116	0.69
3	"b.1"	3.241	0.444	3.249	0.43
4	"b.2"	7.001	1.027	7.006	1.03
5	"sigma.1"	3.423	0.324	3.452	0.34
6	"sigma.2"	3.656	0.35	3.69	0.36

```
• ms1_4_1s = success(rc1_4_1s) &&  
model_summary(post1_4_1s, names(post1_4_1s))
```

```
1.027
```

```
• ms1_4_1s["b.2", "mad_sd"]
```

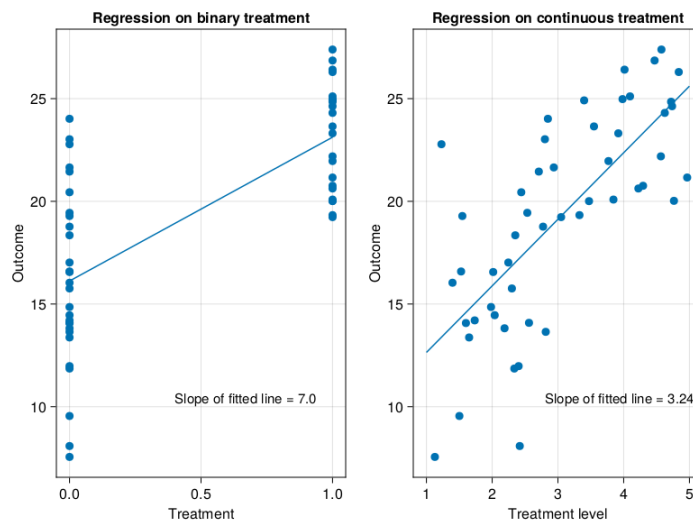
Nested dataframes are handy to obtain a matrix of say the b values:

```
4000x2 Matrix{Float64}:  
 3.68005  7.67235  
 3.34579  5.17041  
 3.35379  5.76383  
 3.05522  6.06792  
 3.00888  5.84699  
 2.72581  6.47463  
 3.96654  6.36506  
  ⋮  
 3.24994  7.27068  
 2.90944  5.6818  
 3.25186  7.30268  
 3.82635  7.82361  
 3.5807   8.38818  
 2.70611  4.93714
```

```
• array(nd1_4_1s, :b)
```

```
4000x2 Matrix{Float64}:  
 3.68005  7.67235  
 3.34579  5.17041  
 3.35379  5.76383  
 3.05522  6.06792  
 3.00888  5.84699  
 2.72581  6.47463  
 3.96654  6.36506  
  ⋮  
 3.24994  7.27068  
 2.90944  5.6818  
 3.25186  7.30268  
 3.82635  7.82361  
 3.5807   8.38818  
 2.70611  4.93714
```

```
• Array(post1_4_1s[:, ["b.1", "b.2"]])
```



```

let
  x1 = 1.0:0.01:5.0
  f = Figure()
  medians = ms1_4_1s[:, "median"]
  ax = Axis(f[1, 2], title = "Regression
on continuous treatment",
    xlabel = "Treatment level", ylabel
    = "Outcome")
  sca1 = scatter!(x, y)
  annotations!("Slope of fitted line =
$(round(medians[3], digits=2))",
    position = (2.8, 10), fontsize=15)
  lin1 = lines!(x1, medians[1] .+
  medians[3] * x1)

  x2 = 0.0:0.01:1.0
  ax = Axis(f[1, 1], title="Regression on
binary treatment",
    xlabel = "Treatment", ylabel =
    "Outcome")
  sca1 = scatter!(x_binary, y)
  lin1 = lines!(x2, medians[2] .+
  medians[4] * x2)
  annotations!("Slope of fitted line =
$(round(medians[4], digits=2))",
    position = (0.4, 10), fontsize=15)
f
end

```

```

• stan1_4_2 = "
• data {
•   int N;
•   vector[N] x;
•   vector[N] y;
• }
• parameters {
•   vector[2] a;
•   real b;
•   real b_exp;
•   vector<lower=0>[2] sigma;
• }
• model {
•   // Priors
•   a ~ normal(10, 5);
•   b ~ normal(0, 5);
•   b_exp ~ normal(5, 5);
•   sigma ~ exponential(1);
•   // Likelihood
•   vector[N] mu;
•   for ( i in 1:N )
•     mu[i] = a[2] + b_exp * exp(-x[i]);
•   y ~ normal(mu, sigma[2]);
•   y ~ normal(a[1] + b * x, sigma[1]);
• }
• "
• ;

```

	parameters	mean	mcse	std	5%
1	"a[1]"	13.0	0.02	0.98	11.0
2	"a[2]"	5.9	0.007	0.38	5.3
3	"b"	-1.74	0.01	0.3	-2.2
4	"b_exp"	17.94	0.06	3.02	12.9
5	"sigma[1]"	2.3	0.0038	0.23	1.9
6	"sigma[2]"	2.2	0.0038	0.23	1.8

```

• let
•   #Random.seed!(1533)
•   n1 = 50
•   x1 = rand(Uniform(1, 5), n1)
•   y1 = [rand(Normal(5 + 30exp(-x1[i]),
•   2), 1)[1] for i in 1:n]
•   data = (N = n1, x = x1, y = y1)
•   global m1_4_2s = SampleModel("m1.4_2s",
•   stan1_4_2);
•   global rc1_4_2s = stan_sample(m1_4_2s;
•   data)
•   success(rc1_4_2s) && describe(m1_4_2s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/7
ted.

```

	parameters	median	mad_sd	mean	std
1	"a.1"	12.758	0.95	12.747	0.98
2	"a.2"	5.919	0.386	5.927	0.38
3	"b"	-1.743	0.285	-1.743	0.29
4	"b_exp"	17.931	2.967	17.936	3.02
5	"sigma.1"	2.266	0.225	2.281	0.23
6	"sigma.2"	2.17	0.224	2.186	0.22

```

• if success(rc1_4_2s)
•   post1_4_2s = read_samples(m1_4_2s,
•   :dataframe)
•   ms1_4_2s = model_summary(post1_4_2s,
•   ["a.1", "a.2", "b", "b_exp", "sigma.1",
•   "sigma.2"])
end

```


nd1_4_2s =

	b	b_exp	a
1	-2.00479	10.8794	▶ [13.6309, 6.32045]
2	-1.48959	17.7817	▶ [11.7728, 6.34427]
3	-1.42222	20.2764	▶ [11.4523, 5.56977]
4	-2.08127	14.8388	▶ [13.6406, 5.95356]
5	-1.75675	17.4339	▶ [13.041, 6.42338]
6	-1.00498	18.4521	▶ [10.4521, 6.09275]
7	-2.09687	18.1361	▶ [13.6547, 5.40958]
8	-1.61418	16.5616	▶ [13.3017, 6.63865]
9	-1.70031	16.8807	▶ [13.174, 6.87947]
10	-2.02728	15.8727	▶ [13.6482, 5.19274]
⋮ more			
4000	-1.97497	16.1358	▶ [13.7143, 5.49498]

- `nd1_4_2s = read_samples(m1_4_2s,
:nesteddataframe)`

4000×2 Matrix{Float64}:

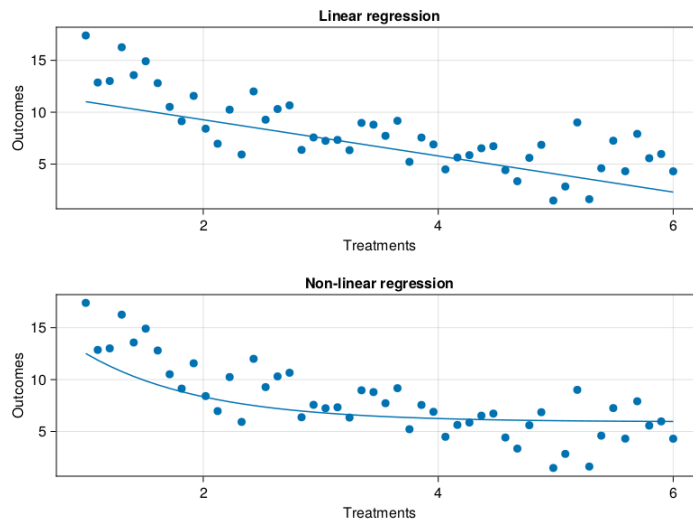
```
13.6309  6.32045
11.7728  6.34427
11.4523  5.56977
13.6406  5.95356
13.041   6.42338
10.4521  6.09275
13.6547  5.40958
⋮
11.4669  6.19831
13.0419  5.50892
13.1873  5.30649
14.6559  6.19775
10.8018  6.18113
13.7143  5.49498
```

- `array(nd1_4_2s, :a)`

- `\hat{a}_1 , \hat{a}_2 , \hat{b} , \hat{b}_{exp} , $\hat{\sigma}_1$, $\hat{\sigma}_2$ = [ms1_4_2s[p,
"median"] for p in ["a.1", "a.2", "b",
"b_exp", "sigma.1", "sigma.2"]];`

5.919

- \hat{a}_2



```
let
  x1 = LinRange(1, 6, 50)
  y1 = [rand(Normal(5 + 30exp(-x1[i]),
    2), 1)[1] for i in 1:length(x1)]

  f = Figure()
  ax = Axis(f[1, 1], title = "Linear
    regression",
    xlabel = "Treatments", ylabel =
      "Outcomes")
  scatter!(x1, y1)
  lines!(x1,  $\hat{a}_1$  .+  $\hat{b}$  .* x1)

  ax = Axis(f[2, 1], title = "Non-linear
    regression",
    xlabel = "Treatments", ylabel =
      "Outcomes")
  scatter!(x1, y1)
  lines!(x1,  $\hat{a}_2$  .+  $\hat{b}_{exp}$  .* exp.(-x1))
f
end
```

	xx	z	yy
1	3.1425	0	37.5294
2	0.0335661	1	30.0628
3	1.60408	0	28.1095
4	0.478735	1	31.3638
5	2.65874	0	35.7382
6	1.02705	1	36.0009
7	1.28799	0	24.3213
8	0.052966	1	29.5242
9	0.543994	0	25.4181
10	0.0304007	1	25.1863
: more			
100	0.00156342	1	28.8751

```

• begin
•   Random.seed!(12573)
•   n2 = 100
•   z = repeat([0, 1]; outer=50)
•   df1_8 = DataFrame()
•   df1_8.xx = [(z[i] == 0 ? rand(Normal(0,
•   1.2), 1).^2 : rand(Normal(0, 0.8),
•   1).^2)[1] for i in 1:n2]
•   df1_8.z = z
•   df1_8.yy = [rand(Normal(20 .+
•   5df1_8.xx[i] .+ 10df1_8.z[i], 3), 1)[1]
•   for i in 1:n2]
•   df1_8
• end

```

```
lm1_8 =
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
yy ~ 1 + xx + z
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	20.1093	0.529823	37.95	<1e-16
xx	4.97503	0.213492	23.30	<1e-16
z	9.625	0.604978	15.91	<1e-16

```
• lm1_8 = lm(@formula(yy ~ xx + z), df1_8)
```

```
lm1_8_0 =
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
yy ~ 1 + xx
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	20.0337	0.544062	36.82	<1e-16
xx	5.01957	0.226965	22.12	<1e-16

```
• lm1_8_0 = lm(@formula(yy ~ xx),
df1_8[df1_8.z .== 0, :])
```


```
lm1_8_1 =
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
yy ~ 1 + xx
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	29.8841	0.49051	60.92	<1e-16
xx	4.66553	0.609796	7.65	<1e-16

```
• lm1_8_1 = lm(@formula(yy ~ xx),
df1_8[df1_8.z .== 1, :])
```



```

• let
•    $\hat{a}_1$ ,  $\hat{b}_1$  = coef(lm1_8_0)
•    $\hat{a}_2$ ,  $\hat{b}_2$  = coef(lm1_8_1)
•   x = LinRange(0, maximum(df1_8.xx), 40)
•
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Figure 1.8")
•   scatter!(df1_8.xx[df1_8.z .== 0],
•   df1_8.yy[df1_8.z .== 0])
•   scatter!(df1_8.xx[df1_8.z .== 1],
•   df1_8.yy[df1_8.z .== 1])
•   lines!(x,  $\hat{a}_1$  .+  $\hat{b}_1$  * x, label =
•   "Control")
•   lines!(x,  $\hat{a}_2$  .+  $\hat{b}_2$  * x, label =
•   "Treated")
•   axislegend(; position=(right, bottom))
•   current_figure()
• end

```

1.5 Classical and Bayesian inference.

1.6 Computing least-squares and Bayesian regression.

1.8 Exercises.

Helicopters

```
helicopters =
```

	Helicopter_ID	width_cm	length_cm	time_s
1	1	4.6	8.2	1.64
2	1	4.6	8.2	1.74
3	1	4.6	8.2	1.68
4	1	4.6	8.2	1.62
5	1	4.6	8.2	1.68
6	1	4.6	8.2	1.7
7	1	4.6	8.2	1.62
8	1	4.6	8.2	1.66
9	1	4.6	8.2	1.69
10	1	4.6	8.2	1.62
⋮ more				
20	2	4.6	8.2	1.61

```
• helicopters =  
  CSV.read(ros_datadir("Helicopters",  
    "helicopters.csv"), DataFrame)
```

Simulate 40 helicopters.

	width_cm	length_cm	time_sec
1	3.34899	0.821947	0.505001
2	2.78591	6.26776	1.33218
3	6.20855	10.1125	1.52182
4	8.01419	9.74486	1.53831
5	2.82847	8.63838	1.13026
6	4.21747	15.7541	1.86944
7	5.37202	9.76108	1.61338
8	6.95056	12.7485	1.79877
9	5.29706	16.2845	2.15528
10	1.71303	9.24657	1.36999
⋮	more		
40	6.74333	10.5377	1.76125

```

• begin
•   helis = DataFrame(width_cm =
•     rand(Normal(5, 2), 40), length_cm =
•     rand(Normal(10, 4), 40))
•   helis.time_sec = 0.5 .+ 0.04 .*
•   helis.width_cm .+ 0.08 .*
•   helis.length_cm .+ 0.1 .*
•   rand(Normal(0, 1), 40)
•   helis
• end

```

```

• stan1_5 = "
• data {
•   int N;
•   vector[N] w;
•   vector[N] l;
•   vector[N] y;
• }
• parameters {
•   real a;
•   real b;
•   real c;
•   real<lower=0> sigma;
• }
• model {
•   // Priors
•   a ~ normal(10, 5);
•   b ~ normal(0, 5);
•   sigma ~ exponential(1);
•
•   // Likelihood time on width
•   vector[N] mu;
•   for ( i in 1:N )
•     mu[i] = a + b * w[i] + c * l[i];
•   y ~ normal(mu, sigma);
• }
• ";

```


	parameters	mean	mcse	std
1	"a"	0.471211	0.00116792	0.05343
2	"b"	0.0400364	0.000170083	0.00764
3	"c"	0.0843282	8.34419e-5	0.00418
4	"sigma"	0.103305	0.000284722	0.01245

```

• let
•   data = (N = nrow(helis), y =
•           helis.time_sec, w = helis.width_cm, l =
•           helis.length_cm)
•   global m1_5s = SampleModel("m1.5s",
•                               stan1_5);
•   global rc1_5s = stan_sample(m1_5s; data)
•   success(rc1_5s) && describe(m1_5s)
end

```

```

/
c
Informational Message: The current Metropolis
Exception: normal_lpdf: Scale parameter is 0,
r/folders/l7/pr04h0650q5dvqttnvs8s2c00000gn/T/
3, column 1 to column 23)
If this warning occurs sporadically, such as f
types like covariance matrices, then the sampl
but if this warning occurs often then your mod
conditioned or misspecified.

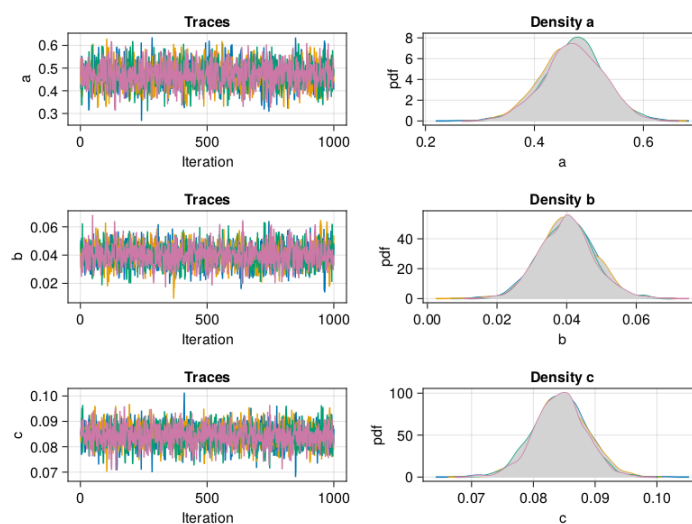
```

	parameters	median	mad_sd	mean	std
1	"a"	0.4724	0.0532	0.4712	0.0532
2	"b"	0.04	0.0074	0.04	0.0074
3	"c"	0.0843	0.0041	0.0843	0.0041
4	"sigma"	0.1022	0.0123	0.1033	0.0123

```

• if success(rc1_5s)
•   post1_5s = read_samples(m1_5s,
•     :dataframe)
•   model_summary(post1_5s, [:a, :b, :c,
•     :sigma]; digits=4)
end


```




```

• plot_chains(post1_5s, [:a, :b, :c])

```



- `plot_chains(post1_5s, [:sigma])`



```

• let
•   w_range = LinRange(1.0, 8.0, 100)
•   w_times = mean.(link(post1_5s, (r, w) -
•     > r.a + r.c + r.b * w, w_range))
•   l_range = LinRange(6.0, 15.0, 100)
•   l_times = mean.(link(post1_5s, (r, l) -
•     > r.a + r.b + r.c * l, l_range))
•
•   f = Figure()
•   ax = Axis(f[1, 1], title = "Time in the
•     air on width and length",
•     xlabel = "Width/Length", ylabel =
•       "Time in the air")
•
•   lines!(w_range, w_times; label="Width")
•   lines!(l_range, l_times; label="Length")
•
•   f[1, 2] = Legend(f, ax, "Regression
•     lines", framevisible = false)
•
•   current_figure()
• end

```

```
lnk1_5s =
```

```
▶ [[0.894687, 0.881022, 0.968542, 0.959784, 0.984687,
```

```
• lnk1_5s = link(post1_5s, (r, l) -> r.a + r.b
•   + r.c * l, [5, 10, 12])
```

```
▶ [0.933737, 1.35465, 1.52305]
```

```
• median.(lnk1_5s)
```

```
► [0.038196, 0.03567, 0.0378395]
```

```
• mad.(lnk1_5s)
```

```
► [0.932888, 1.35453, 1.52319]
```

```
• mean.(link(post1_5s, (r, l) -> r.a + r.b +  
r.c * l, [5, 10, 12]))
```

	a	b	c	sigma
1	0.424878	0.0489688	0.084168	0.106223
2	0.406936	0.0504471	0.0847278	0.096517
3	0.516133	0.0306989	0.084342	0.100765
4	0.531751	0.0441585	0.0767748	0.101889
5	0.554138	0.0364994	0.0788079	0.107534
6	0.535185	0.0398682	0.0791222	0.103828
7	0.495928	0.0451504	0.0801624	0.103149
8	0.458166	0.0361234	0.0874675	0.105228
9	0.450037	0.0333757	0.0870754	0.106536
10	0.451163	0.0434026	0.0841383	0.096611
⋮ more				
4000	0.506864	0.0397577	0.0804016	0.117443

```
• read_samples(m1_5s, :nesteddataframe)
```

```
No nested columns found.
```