

See chapter 7 in Regression and Other Stories.

.....

Widen the notebook.

```
• html"""
• <style>
•   main {
•     margin: 0 auto;
•     max-width: 2000px;
•     padding-left: max(160px, 10%);
•     padding-right: max(160px, 10%);
•   }
• </style>
• """
```

```
• using Pkg ✓ , DrWatson ✓
```

A typical set of Julia packages to
include in notebooks.

```
• begin
•   # Specific to this notebook
•   using GLM ✓
•
•   # Specific to ROSStanPluto
•   using StanSample ✓
•
•   # Graphics related
•   using GLMakie ✓
•
•   # Common data files and functions
•   using RegressionAndOtherStories ✓
• end
```

7.1 Example: Predicting presidential vote from the economy.

`hdi =`

	rank	state	hdi	canada
1	1	"Connecticut"	0.962	2
2	2	"Massachusetts"	0.961	2
3	3	"New Jersey"	0.961	2
4	4	"Washington, D.C."	0.96	4
5	5	"Maryland"	0.96	3
6	6	"Hawaii"	0.959	2
7	7	"New York"	0.959	1
8	8	"New Hampshire"	0.958	1
9	9	"Minnesota"	0.958	1
10	10	"Rhode Island"	0.958	3
	: more			
51	51	"Mississippi"	0.799	5

```
• hdi = CSV.read(ros_datadir("HDI",  
"hdi.csv"), DataFrame)
```

```
hibbs =
```

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
⋮ more				
16	2012	0.95	52.0	"Obama"

```
• hibbs =  
  CSV.read(ros_datadir("ElectionsEconomy",  
    "hibbs.csv"), DataFrame)
```

```
hibbs_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM},  
  vote ~ 1 + growth
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	46.2476	1.62193	28.51	<1e-30
growth	3.06053	0.696274	4.40	0.0004

```
• hibbs_lm = lm(@formula(vote ~ growth),  
  hibbs)
```

```
► [-8.99292, 2.66743, 1.0609, 2.20753, -5.89044, 4.27444]
```

```
• residuals(hibbs_lm)
```

```
2.2744434224582912
```

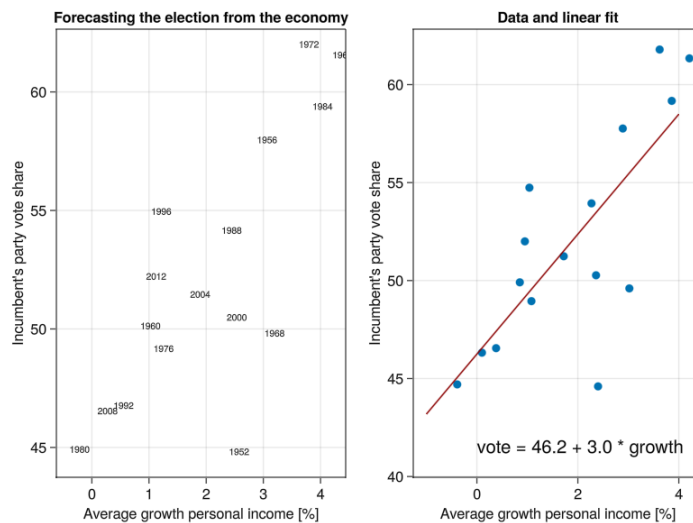
```
• mad(residuals(hibbs_lm))
```

```
3.635681268522063
```

```
• std(residuals(hibbs_lm))
```

```
► [46.2476, 3.06053]
```

```
• coef(hibbs_lm)
```



```

let
  fig = Figure()
  hibbs.label = string.(hibbs.year)
  xlabel = "Average growth personal
income [%]"
  ylabel = "Incumbent's party vote share"
  let
    title = "Forecasting the election
from the economy"
    ax = Axis(fig[1, 1]; title, xlabel,
ylabel)
    for (ind, yr) in
      enumerate(hibbs.year)
        annotations!("$ (yr)"; position=
(hibbs.growth[ind],
hibbs.vote[ind]), textsize=10)
    end
  end
let
  x = LinRange(-1, 4, 100)
  title = "Data and linear fit"
  ax = Axis(fig[1, 2]; title, xlabel,
ylabel)
  scatter!(hibbs.growth, hibbs.vote)
  lines!(x, coef(hibbs_lm)[1] .+
coef(hibbs_lm)[2] .* x;
color=:darkred)
  annotations!("vote = 46.2 + 3.0 *
growth"; position=(0, 41))
end
fig
end

```

```

• stan7_1 = "
• data {
•   int<lower=1> N;      // total number of
•   observations
•   vector[N] growth;   // Independent
•   variable: growth
•   vector[N] vote;     // Dependent
•   variable: votes
• }
• parameters {
•   real b;              // Coefficient
•   independent variable
•   real a;              // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(50, 20);
•   b ~ normal(2, 10);
•   sigma ~ exponential(1);
•
•   mu = a + b * growth;
•
•   // likelihood including constants
•   vote ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std	
1	"b"	3.03591	0.0167247	0.652008	1
2	"a"	46.2957	0.0392767	1.52911	4
3	"sigma"	3.57778	0.0136109	0.602732	2

```

• let
•   data = (N=nrow(hibbs), vote=hibbs.vote,
•         growth=hibbs.growth)
•   global m7_1s = SampleModel("hibbs",
•         stan7_1)
•   global rc7_1s = stan_sample(m7_1s; data)
•         success(rc7_1s) && describe(m7_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/1
d.

```

Note


Sometimes I hide or show the output logs. To show them, click on the little circle with 3 dots visible in the top right of the input cell if the cursor is in there. Try it!

	parameters	median	mad_sd	mean	std
1	"a"	46.308	1.479	46.296	1.52
2	"b"	3.047	0.633	3.036	0.65
3	"sigma"	3.496	0.566	3.578	0.60


```

• if success(rc7_1s)
•   post7_1s = success(rc7_1s) &&
•   read_samples(m7_1s, :dataframe)
•   ms7_1s = model_summary(post7_1s, [:a,
•   :b, :sigma])
end

```



- `trankplot(post7_1s, "b")`



```


• let
•   growth_range =
•     LinRange(minimum(hibbs.growth),
•       maximum(hibbs.growth), 200)
•   votes = mean.(link(post7_1s, (r,x) ->
•     r.a + x * r.b, growth_range))
•
•   fig = Figure()
•   xlabel = "Average growth personal
•     income [%]"
•   ylabel="Incumbent's party vote share"
•   ax = Axis(fig[1, 1]; title="Regression
•     line based on 4000 posterior samples",
•     subtitle = "(grey lines are based
•       on first 200 draws of :a and :b)",
•     xlabel, ylabel)
•   for i in 1:200
•     lines!(growth_range, post7_1s.a[i]
•       .+ post7_1s.b[i] .* growth_range,
•       color = :lightgrey)
•   end
•   scatter!(hibbs.growth, hibbs.vote)
•   lines!(growth_range, votes, color =
•     :red)
•   fig
• end

```

0.7385523916379624

```
• let
•   println(46.3 + 3 * 2.0) # 52.3,  $\sigma = 3.6$ 
•   (from ms7_1s above)
•   probability_of_Clinton_winning = 1 -
•   cdf(Normal(52.3, 3.6), 50)
end
```

52.3



```
• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title = "")
•   x_range = LinRange(30, 70, 100)
•   y = pdf.(Normal(52.3, 3.6), x_range)
•   lines!(x_range, y)
•
•   x1 = range(50, 70; length=200)
•   band!(x1, fill(0, length(x1)), pdf.
•   (Normal(52.3, 3.6), x1);
•   color = (:grey, 0.75), label =
•   "Label")
•
•   annotations!("Predicted\n74% change\nof
•   Clinton victory", position=(51, 0.02),
•   fontsize=13)
•   f
end
```

7.2 Checking the model-fitting procedure using simulation.

	parameters	mean	mcse	std	
1	"b"	2.94595	0.0163709	0.633977	1
2	"a"	45.3194	0.0365307	1.47054	2
3	"sigma"	3.39667	0.0139647	0.58763	2

```
• let
•   a = 46.3
•   b = 3.0
•   sigma = 3.9
•   x = hibbs.growth
•   n = length(x)
•
•   y = a .+ b .* x + rand(Normal(0,
•   sigma), n)
•   fake = DataFrame(x = x, y = y)
•
•   data = (N=nrow(fake), vote=fake.y,
•   growth=fake.x)
•   global m7_2s = SampleModel("fake",
•   stan7_1)
•   global rc7_2s = stan_sample(m7_2s; data)
•   success(rc7_2s) && describe(m7_2s)
end
```

```
/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T
d.
```

	parameters	median	mad_sd	mean	std
1	"b"	2.957	0.603	2.946	0.63
2	"a"	45.336	1.403	45.319	1.47
3	"sigma"	3.318	0.555	3.397	0.58

```

• if success(rc7_2s)
•   post7_2s = read_samples(m7_2s,
•   :dataframe)
•   ms7_2s = model_summary(post7_2s,
•   names(post7_2s))
end

```

	parameters	median	mad_sd	mean	std
1	"a"	46.308	1.479	46.296	1.52
2	"b"	3.047	0.633	3.036	0.65
3	"sigma"	3.496	0.566	3.578	0.60

```

• ms7_1s

```

sim (generic function with 1 method)

```
• function sim(sm::SampleModel)
•   a = 46.3
•   b = 3.0
•   sigma = 3.9
•   x = hibbs.growth
•   n = length(x)
•
•   y = a .+ b .* x + rand(Normal(0,
•   sigma), n)
•   println(mean(y))
•   data_sim = (N=n, vote=y, growth=x)
•   rc = stan_sample(sm; data=data_sim)
•   post = read_samples(sm, :dataframe)
•   ms = model_summary(post, Symbol.([:a,
•   :b, :sigma]))
•   b̂ = ms[:b, :mean]
•   b_se = ms[:b, :std]
•
•   (
•     b̂ = b̂,
•     b_se = b_se,
•     cover_68 = Int(abs(b - b̂) < b_se),
•     cover_95 = Int(abs(b - b̂) < 2b_se)
•   )
• end
```

```
• m7_2_1s = SampleModel("fake_sim", stan7_1);
```

```
/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
ated.
```

► (\hat{b} = 2.606, b_se = 0.548, cover_68 = 1, cover_95 = 1)

```
• sim(m7_2_1s)
```

```
50.21592438279586 ⓘ
```

false

```
• isdefined(Main, :StanSample)
```

	parameters	median	mad_sd	mean	std
1	"a"	46.308	1.479	46.296	1.52
2	"b"	3.047	0.633	3.036	0.65
3	"sigma"	3.496	0.566	3.578	0.60

- `ms7_1s`

Or use the underlying DataFrame directly.

1.479

- `ms7_1s["a", "mad_sd"]`

1.479

- `ms7_1s[:a, :mad_sd]`

► [1.479, 0.633, 0.566]

- `ms7_1s[:, :mad_sd]`

- `ms7_1s[:c, :mad_sd]`

Parameter `c` is not in ["a", "b", "sigma"].

- `ms7_1s[:a, :mad]`

Statistic `mad` is not in ["parameters", "median"]

DataFrameRow (2 columns)

	median	mad_sd
	Float64	Float64
3	3.496	0.566

- `ms7_1s[3, [:median, :mad_sd]]`

String

- `eltype(ms7_1s.parameters)`

	variable	mean	min	median	max
1	: \hat{b}	3.117	2.017	3.277	3.696
2	:b_se	0.704	0.587	0.704	0.824
3	:cover_68	0.8	0	1.0	1
4	:cover_95	1.0	1	1.0	1

```

• let
•   n_fake = 10 # 1000
•   df = DataFrame()
•   cover_68 = Float64[]
•   cover_95 = Float64[]
•   m7_2_1s = SampleModel("fake_sim_1",
•   stan7_1)
•
•   for i in 1:n_fake
•       res = sim(m7_2_1s)
•       append!(df, DataFrame(;res...))
•   end
•   describe(df)
end

```

Note

In above cell, I have hidden the logs. To show them, click on the little circle with 3 dots.

7.3 Formulating comparisons as regression models.

```

• stan7_3 = "
• data {
•   int N;
•   vector[N] y;
• }
• parameters {
•   real a;
•   real sigma;
• }
• model {
•   y ~ normal(a, sigma);
• }
• ";

```

► [3.305, 1.12992]

```

• begin
•   r_0 = [-0.3, 4.1, -4.9, 3.3, 6.4, 7.2,
•         10.7, -4.6, 4.7, 6.0, 1.1, -6.7, 10.2, 9.7,
•         5.6,
•         1.7, 1.3, 6.2, -2.1, 6.5]
•   [mean(r_0), std(r_0)/sqrt(length(r_0))]
• end

```

► (diff = 4.89914, se_0 = 1.12992, se_1 = 0.89368

```

• begin
•   Random.seed!(3)
•   n_0 = 20
•   y_0 = r_0
•   fake_0 = DataFrame(y_0 = r_0)
•   data_0 = (N = nrow(fake_0), y =
•             fake_0.y_0)
•
•   n_1 = 30
•   y_1 = rand(Normal(8.0, 5.0), n_1)
•   data_1 = (N = n_1, y = y_1)
•
•   se_0 = std(y_0)/sqrt(n_0)
•   se_1 = std(y_1)/sqrt(n_1)
•
•   (diff=mean(y_1)-mean(y_0), se_0=se_0,
•    se_1=se_1, se=sqrt(se_0^2 + se_1^2))
• end

```


	parameters	mean	mcse	std
1	"a"	3.30376	0.0264402	1.26012
2	"sigma"	5.42711	0.0221087	0.954644

```

• begin
•   m7_3_0s = SampleModel("fake_0", stan7_3)
•   rc7_3_0s = stan_sample(m7_3_0s;
•   data=data_0)
•   success(rc7_3_0s) && describe(m7_3_0s)
end

```

	parameters	mean	mcse	std
1	"a"	8.1841	0.0203487	0.957577
2	"sigma"	5.12561	0.014233	0.709505

```

• begin
•   m7_3_1s = SampleModel("fake_1", stan7_3)
•   rc7_3_1s = stan_sample(m7_3_1s;
•   data=data_1)
•   success(rc7_3_1s) && describe(m7_3_1s)
end

```

Note

In above cells, the logs are hidden.

	parameters	median	mad_sd	mean	std
1	"a"	3.293	1.23	3.304	1.26
2	"sigma"	5.302	0.873	5.427	0.95

```

• if success(rc7_3_0s)
•   post7_3_0s = read_samples(m7_3_0s,
•   :dataframe)
•   sm7_3_0s = model_summary(post7_3_0s,
•   [:a, :sigma])
end

```

	parameters	median	mad_sd	mean	std
1	"a"	8.187	0.92	8.184	0.95
2	"sigma"	5.041	0.684	5.126	0.71

```

• if success(rc7_3_1s)
•   post7_3_1s = read_samples(m7_3_1s,
•   :dataframe)
•   sm7_3_1s = model_summary(post7_3_1s,
•   [:a, :sigma])
end

```

```

• stan7_3_2 = "
• data {
•   int N;
•   vector[N] y;
•   vector[N] x;
• }
• parameters {
•   real a;
•   real b;
•   real sigma;
• }
• model {
•   vector[N] mu;
•   mu = a + b * x;
•   y ~ normal(mu, sigma);
• }
• ";

```

	parameters	mean	mcse	std	5%
1	"a"	3.32	0.03	1.12	1.43
2	"b"	4.89	0.03	1.44	2.56
3	"sigma"	5.07	0.01	0.53	4.28

```

• let
•   n = n0 + n1
•   y = vcat(y0, y1)
•   x = vcat(zeros(Int, n0), ones(Int, n1))
•   global fake = DataFrame(x=x, y=y)
•   data = (N = n, x = x, y = y)
•   global m7_3_2s = SampleModel("fake_2",
•   stan7_3_2)
•   global rc7_3_2s = stan_sample(m7_3_2s;
•   data)
•   success(rc7_3_2s) && describe(m7_3_2s,
•   [:a, :b, :sigma])
end

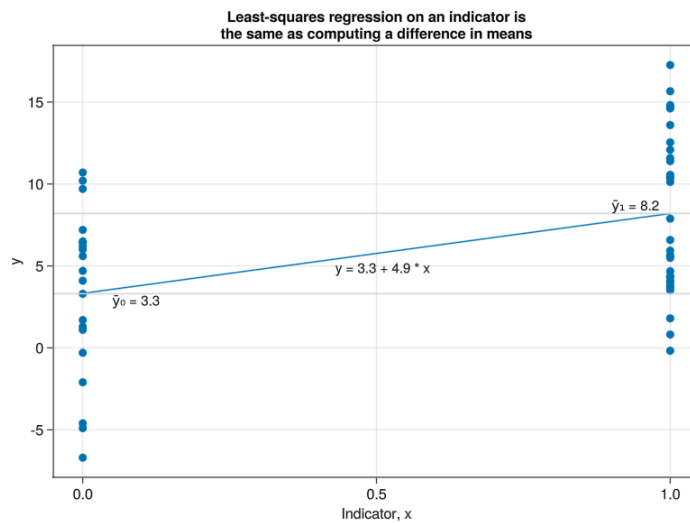
```

	parameters	median	mad_sd	mean	std
1	"a"	3.327	1.117	3.317	1.11
2	"b"	4.871	1.45	4.892	1.44
3	"sigma"	5.026	0.531	5.073	0.53

```

• if success(rc7_3_2s)
•   post7_3_2s = read_samples(m7_3_2s,
•   :dataframe)
•   sm7_3_2s = model_summary(post7_3_2s,
•   [:a, :b, :sigma])
end

```



```

let
  f = Figure()
  ax = Axis(f[1, 1]; title="Least-squares
    regression on an indicator is\nthe same
    as computing a difference in means",
    xlabel="Indicator, x", ylabel="y")
  x_range = LinRange(0, 1, 100)
  â = sm7_3_2s[:a, :median]
  b̂ = sm7_3_2s[:b, :median]
  y = â .+ b̂ .* x_range
  lines!(x_range, y)
  x = vcat(zeros{Int, n0}, ones{Int, n1})
  scatter!(fake.x, fake.y)
  ȳ₀ = mean(y₀)
  ȳ₁ = mean(y₁)
  hlines!(ax, [ȳ₀, ȳ₁]; color=:lightgrey)
  annotations!("ȳ₀ = $(round(ȳ₀,
    digits=1))", position=(0.05, 2.4),
    fontsize=15)
  annotations!("ȳ₁ = $(round(ȳ₁,
    digits=1))", position=(0.9, 8.2),
    fontsize=15)
  annotations!("y = $(round(â, digits=1))
    + $(round(b̂, digits=1)) * x", position=
    (0.43, 4.4), fontsize=15)
  f
end

```