

See chapter 8 in Regression and Other Stories.

Widen the notebook.

```
• html"""  
• <style>  
•   main {  
•     margin: 0 auto;  
•     max-width: 2000px;  
•     padding-left: max(160px, 10%);  
•     padding-right: max(160px, 10%);  
•   }  
• </style>  
• """
```

```
• using Pkg ✓ , DrWatson ✓
```

A typical set of Julia packages to include in notebooks.

```
• begin  
•   # Specific to this notebook  
•   using GLM ✓  
•   using Optim ✓  
•   
•   # Specific to ROSStanPluto  
•   using StanSample ✓  
•   using StanOptimize ✓  
•   
•   # Graphics related  
•   using GLMakie ✓  
•   
•   # Common data files and functions  
•   using RegressionAndOtherStories ✓  
• end
```

8.1 Least squares, maximum likelihood, and Bayesian inference.

	x	y	ϵ	error
1	0.0	45.9177	-0.282333	-0.282333
2	0.0251256	48.4013	2.12591	2.12591
3	0.0502513	43.1233	-3.22741	-3.22741
4	0.0753769	56.2541	9.82797	9.82797
5	0.100503	51.161	4.6595	4.6595
6	0.125628	47.6471	1.07026	1.07026
7	0.150754	53.652	6.99973	6.99973
8	0.175879	43.4236	-3.30408	-3.30408
9	0.201005	42.632	-4.17101	-4.17101
10	0.226131	45.5619	-1.31654	-1.31654
⋮ more				
200	5.0	57.0887	-4.11126	-4.11126

```

• let
•   Random.seed!(1)
•   a = 46.2
•   b = 3.0
•   sigma = 4.0
•   x = LinRange(0, 5, 200)
•   ϵ = rand(Normal(0, sigma), length(x))
•   y = a .+ b .* x .+ ϵ
•
•   # DataFrame used to collect differen
•   estimates, shown later on.
•
•   global estimate_comparison = DataFrame()
•   estimate_comparison.parameters = [:a,
•   :b, :sigma]
•
•   global sim = DataFrame(x = x, y = y, ϵ
•   = ϵ, error = y .- (a .+ b .* x))
end

```

```

• stan8_1 = "
• data {
•   int<lower=1> N;      // total number of
•   observations
•   vector[N] x;        // Independent
•   variable: growth
•   vector[N] y;        // Dependent
•   variable: votes
• }
• parameters {
•   real b;              // Coefficient
•   independent variable
•   real a;              // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(1, 5);
•   b ~ normal(1, 5);
•   sigma ~ exponential(1);
•
•   mu = a + b * x;
•
•   // likelihood including constants
•   y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std
1	"b"	3.2503	0.00490206	0.208701
2	"a"	45.6201	0.0141328	0.599706
3	"sigma"	4.37646	0.00496148	0.22432

```

• let
•   data = (N = nrow(sim), x = sim.x, y =
•   sim.y)
•   global m8_1s = SampleModel("m8_1s",
•   stan8_1)
•   global rc8_1s = stan_sample(m8_1s; data)
•   success(rc8_1s) && describe(m8_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T
d.

```

	parameters	median	mad_sd	mean	std
1	"a"	45.628	0.595	45.62	0.6
2	"b"	3.251	0.206	3.25	0.20
3	"sigma"	4.368	0.22	4.376	0.22

```

• if success(rc8_1s)
•   post8_1s = read_samples(m8_1s,
•   :dataframe)
•   ms8_1s = model_summary(post8_1s, [:a,
•   :b, :sigma])
•   estimate_comparison[!, :m8_1s] =
•   [Vector(i) for i in eachrow(ms8_1s[:,
•   :median, :mad_sd])]]
•   ms8_1s
end

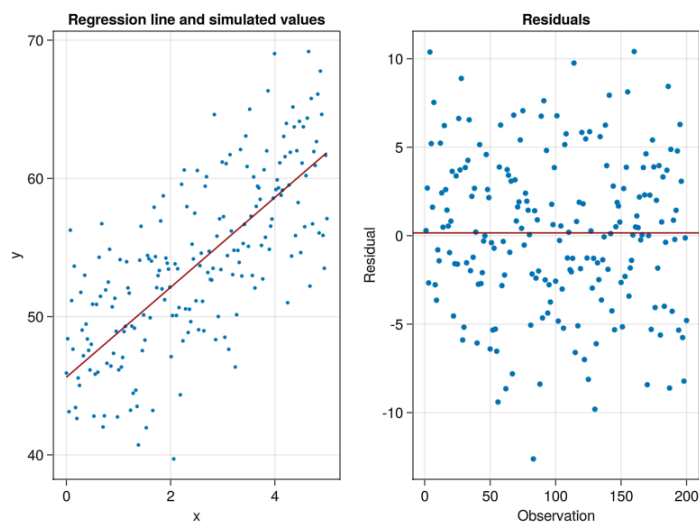
```

	x	y	ϵ	error
1	0.0	45.9177	-0.282333	-0.282333
2	0.0251256	48.4013	2.12591	2.12591
3	0.0502513	43.1233	-3.22741	-3.22741
4	0.0753769	56.2541	9.82797	9.82797
5	0.100503	51.161	4.6595	4.6595
6	0.125628	47.6471	1.07026	1.07026
7	0.150754	53.652	6.99973	6.99973
8	0.175879	43.4236	-3.30408	-3.30408
9	0.201005	42.632	-4.17101	-4.17101
10	0.226131	45.5619	-1.31654	-1.31654
⋮ more				
200	5.0	57.0887	-4.11126	-4.11126

```

• let
•    $\hat{a}$  = ms8_1s[:a, :median]
•    $\hat{b}$  = ms8_1s[:b, :median]
•   sim.residual = sim.y .- ( $\hat{a}$  .+  $\hat{b}$  .*
•   sim.x)
•   sim
• end

```




```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Regression
•     line and simulated values", xlabel="x",
•     ylabel="y")
•   x_range = LinRange(minimum(sim.x),
•     maximum(sim.x), 200)
•   y_res = mean.(link(post8_1s, (r,x) ->
•     r.a + x * r.b, x_range))
•   scatter!(sim.x, sim.y; markersize=4)
•   lines!(x_range, y_res; color=:darkred)
•
•   ax = Axis(f[1, 2]; title="Residuals",
•     xlabel="Observation", ylabel="Residual")
•   scatter!(sim.residual; markersize=6)
•   hlines!(ax, mean(sim.residual);
•     color=:darkred)
•   f
• end

```

```
RSS = 3815.9793429532633
```

```
• RSS = sum(sim.residual .^ 2)
```



```

• let
•    $\hat{a}$  = ms8_1s[:a, :mean]
•    $\hat{b}$  = ms8_1s[:b, :mean]
•
•   f = Figure()
•   ax = Axis(f[1, 1]; title="RSS as a
•   function of a", xlabel="a",
•   ylabel="RSS")
•   a_range = LinRange(43, 48, 100)
•   r = [sum((sim.y .- (k .+  $\hat{b}$  .* sim.x))
•   .^ 2) for k in a_range]
•   lines!(a_range, r)
•   annotations!("$(( $\hat{a}$  =  $\hat{a}$ ,  $\hat{b}$  =  $\hat{b}$ ))",
•   position=(44.5, 4500), fontsize=15)
•   ax = Axis(f[1, 2]; title="RSS as a
•   function of b", xlabel="b",
•   ylabel="RSS")
•   b_range = LinRange(2.1, 4.4, 100)
•   r = [sum((sim.y .- ( $\hat{a}$  .+ k .* sim.x))
•   .^ 2) for k in b_range]
•   lines!(b_range, r)
•   annotations!("$(( $\hat{a}$  =  $\hat{a}$ ,  $\hat{b}$  =  $\hat{b}$ ))",
•   position=(2.58, 4800), fontsize=15)
•   f
• end

```

Least squares

► (46.2831, 3.05172)

```
• let
•   global lsq = [0.0 missing; 0.0 missing;
•   0.0 missing]
•   df = DataFrame(ones = ones(nrow(sim)), x
•   = sim.x)
•   X = Array(df)
•   Xt = transpose(X)
•   â, b̂ = (Xt * X)^-1 * Xt * sim.y
•   lsq[1, 1] = â
•   lsq[2, 1] = b̂
•   â, b̂
end
```

► (â = 46.2831, b̂ = 3.05172)

```
• let
•   b̂ = sum((sim.x .- mean(sim.x)) .*
•   sim.y) / sum(((sim.x .- mean(sim.x)) .^
•   2))
•   â = mean(sim.y) - b̂ * mean(sim.x)
•   (â = â, b̂ = b̂)
end
```

4.390059560439744

```
• let
•   ô = sqrt(sum(sim.residual .^
•   2)/(nrow(sim) - 2))
•   lsq[3, 1] = ô
•   estimate_comparison[!, :least_squares]
•   = [Vector(i) for i in eachrow(lsq)]
•   ô
end
```

Maximum likelihood

loglik (generic function with 1 method)

```
• function loglik(x)
•   ll = 0.0
•   ll += log(pdf(Normal(50, 20), x[1]))
•   ll += log(pdf(Normal(2, 10), x[2]))
•   ll += log(pdf(Exponential(1), x[3]))
•   for i in 1:nrow(sim)
•       ll += sum(logpdf.(Normal(x[1] .+
•       x[2] .* sim.x[i], x[3]), sim.y[i]))
•   end
•   -ll
end
```



```
0.1353352832366127
```

```
• pdf(Exponential(1), 2.0)
```

```
► [170.0, 10.0, 2.0]
```

```
• begin  
•   lower = [0.0, 0.0, 0.0]  
•   upper = [250.0, 50.0, 10.0]  
•   x0 = [170.0, 10.0, 2.0]  
• end
```

```
res =  
* Status: success (objective increased between :  
  
* Candidate solution  
  Final objective value:      5.895739e+02  
  
* Found with  
  Algorithm:      Fminbox with L-BFGS  
  
* Convergence measures  
   $|x - x'|$  = 2.27e-08  $\nless 0.0e+00$   
   $|x - x'|/|x'|$  = 4.88e-10  $\nless 0.0e+00$   
   $|f(x) - f(x')|$  = 0.00e+00  $\nless 0.0e+00$   
   $|f(x) - f(x')|/|f(x')|$  = 0.00e+00  $\nless 0.0e+00$   
   $|g(x)|$  = 6.16e-09  $\nless 1.0e-08$   
  
* Work counters  
  Seconds run:      1 (vs limit Inf)  
  Iterations:      5  
  f(x) calls:      120  
   $\nabla f(x)$  calls:    120  
  
• res = optimize(loglik, lower, upper, x0)
```

```
► [46.2877, 3.05023, 4.30947]
```

```
• let  
•   mle = Optim.minimizer(res)  
•   lsq[:, 1] = mle  
•   estimate_comparison[!, :mle] =  
•   [Vector(i) for i in eachrow(lsq)]  
•   mle  
• end
```

MLE estimate (using StanOptimize and 4 chains)

	a	b	sigma
1	45.6309	3.24582	4.32164
2	45.6312	3.24574	4.32149
3	45.6313	3.24578	4.32148
4	45.6313	3.24577	4.32148

```

• let
•   data = (N=nrow(sim), y=sim.y, x=sim.x)
•   o8_1s = OptimizeModel("m8_1s", stan8_1)
•   rc8_1s = stan_optimize(o8_1s; data)
•   result = success(rc8_1s) &&
•   read_optimize(o8_1s)
•   global o8_1_df = DataFrame()
•   for p in ["a", "b", "sigma"]
•       o8_1_df[!, p] = result[1][p]
•   end
•   o8_1_df
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T
d.

```

Compare the four results.

	parameters	m8_1s	least_squ
1	:a	▶ [45.628, 0.595]	▶ [46.2831, m
2	:b	▶ [3.251, 0.206]	▶ [3.05172, m
3	:sigma	▶ [4.368, 0.22]	▶ [4.39006, m


```

• let
•   lsq[:, 1] = mean(Array(o8_1_df); dims=1)
•   estimate_comparison[!, :o8_1s] =
•   [Vector(i) for i in eachrow(lsq)]
•   estimate_comparison
end

```

590.2799283100538

- `loglik([45.6, 3.25, 4.4])`



```
let
  f = Figure()
  ax = Axis(f[1, 1])
  lines!(30:0.1:60, [-loglik([a, 3.25,
  4.4]) for a in 30:0.1:60])
  ax = Axis(f[1, 2])
  lines!(0:0.1:5, [-loglik([46.5, b,
  4.4]) for b in 0:0.1:5])
  f
end
```

600.0086334504888

- `loglik([45, 3, 4.4])`

2×200 Matrix{Float64}:

1.0	1.0201	1.0402	1.0603	1.0804	...
46.6171	43.4073	42.8262	48.5102	51.4102	

```
let
  using StatsAPI ✓
  Random.seed!(123)
  a = 46.2
  b = 3.0
  sigma = 4.0
  x = LinRange(1, 5, 200)
  ε = rand(Normal(0, sigma), length(x))
  y = a .+ b .* x .+ ε
  global obs = Matrix(hcat(x, y)')
end
```

```
distr8_1 =
FullNormal(
dim: 2
μ: [2.9999999999999996, 55.07524018977074]
Σ: [1.3467336683417086 3.9331584800491735; 3.9331584800491735 55.02344421495838]
)
```


```
• distr8_1 = fit_mle(MvNormal, obs)
```

```
2×1 Matrix{Float64}:
 3.000724768186341
 55.02344421495838
```

```
• mean(rand(distr8_1, 1000); dims=2)
```

```
-3.389758334022121
```

```
• loglikelihood(distr8_1, [3, 55])
```



```
• let
•     a = collect(LinRange(30, 80, 50))
•     b = collect(LinRange(0, 8, 50))
•     global z = [loglikelihood(distr8_1, [b,
•     a]) for a in a, b in b]
•     m, i = findmax(z)
•     maxz = [a[i[1]], b[i[1]], z[i]]
•     println(maxz)
•     wireframe(a, b, z, axis=(type=Axis3,))
end
```


```
[55.51020408163265, 4.081632653061225, -3.39401]
```

```
my_μ = ▶ [45.62, 3.25]
```

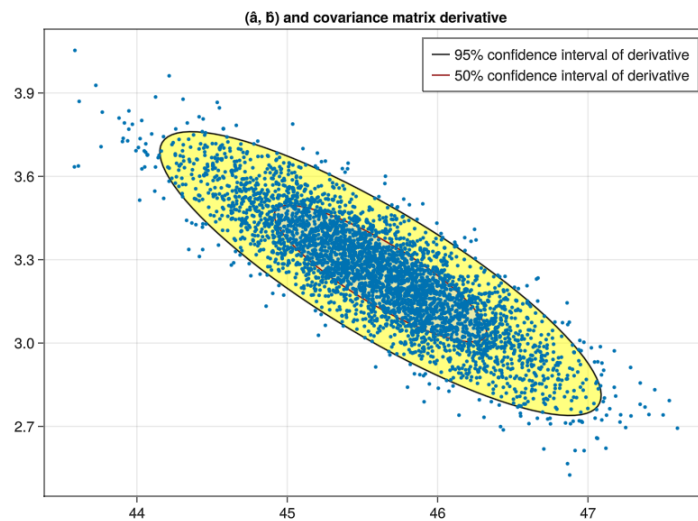
```
• my_μ = [ms8_1s["a", "mean"], ms8_1s["b",
"mean"]]
```

```
my_Σ = 2×2 Matrix{Float64}:
  0.359647 -0.107326
 -0.107326  0.043556
```

- `my_Σ = cov([post8_1s.a post8_1s.b])`



```
let
  f = Figure()
  ax = Axis(f[1, 1]; title="(â, b) and
  covariance matrix derivative")
  lines!(getellipsepoints(my_μ, my_Σ)...,
  label="95% confidence interval of
  derivative", color=:black)
  lines!(getellipsepoints(my_μ, my_Σ,
  0.5)..., label="50% confidence interval
  of derivative", color=:darkred)
  scatter!(post8_1s.a, post8_1s.b;
  markersize=4)
  axislegend(position=:rt)
f
end
```

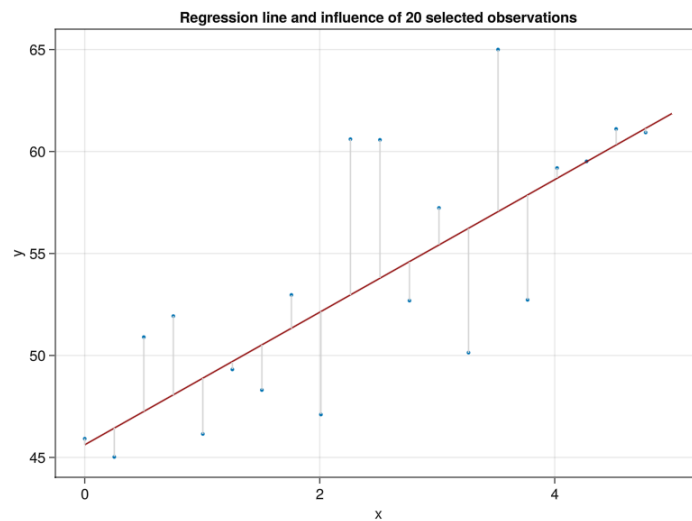


```

let
  f = Figure()
  ax = Axis(f[1, 1]; title="(â, b) and
  covariance matrix derivative")
  poly!(Point2f.
    (zip(getellipsepoints(my_μ, my_Σ)...));
    color=(:yellow, 0.5))
  poly!(Point2f.
    (zip(getellipsepoints(my_μ, my_Σ,
    0.50)...)); color=(:lightgrey, 0.5))
  lines!(getellipsepoints(my_μ, my_Σ)...,
    label="95% confidence interval of
    derivative", color=:black)
  lines!(getellipsepoints(my_μ, my_Σ,
    0.5)..., label="50% confidence interval
    of derivative", color=:darkred)
  scatter!(post8_1s.a, post8_1s.b;
    markersize=4)
  axislegend(position=:rt)
f
end

```

8.2 Influence of individual points in a fitted regression.



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1]; title="Regression
•   line and influence of 20 selected
•   observations", xlabel="x", ylabel="y")
•   x_range = LinRange(minimum(sim.x),
•   maximum(sim.x), 200)
•   y_res = mean.(link(post8_1s, (r,x) ->
•   r.a + x * r.b, x_range))
•   select_obs = 1:10:200
•   scatter!(sim.x[select_obs],
•   sim.y[select_obs]; markersize=4)
•   lines!(x_range, y_res; color=:darkred)
•   for ind in select_obs
•       ymin = min(sim.y[ind], y_res[ind])
•       ymax = max(sim.y[ind], y_res[ind])
•       lines!([sim.x[ind], sim.x[ind]],
•       [ymin, ymax]; color=:lightgrey)
•   end
•   f
• end

```

8.3 Least squares slope as a weighted average of slopes of pairs.

```

▶ (weighted_slopes = 3.05172, least_squares = [3.
• let
•   s1 = sum([(sim.x[i]-sim.x[j]) *
•     (sim.y[i]-sim.y[j]) for i in
•     1:length(sim.x), j in 1:length(sim.y)])
•   s2 = sum([(sim.x[i]-sim.x[j])^2 for i in
•     1:length(sim.x), j in 1:length(sim.y)])
•   (weighted_slopes = round(s1/s2;
•     digits=5),
•     least_squares=estimate_comparison[2,
•       :least_squares])
end

```

8.4 Comparing two fitting functions: glm and stan_sample.

```

• stan8_2 = "
• data {
•   int<lower=1> N;      // total number of
•   observations
•   vector[N] x;        // Independent
•   variable: growth
•   vector[N] y;        // Dependent
•   variable: votes
• }
• parameters {
•   real b;              // Coefficient
•   independent variable
•   real a;              // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(0, 50);
•   b ~ normal(0, 50);
•   sigma ~ uniform(0, 50);
•
•   mu = a + b * x;
•
•   // likelihood including constants
•   y ~ normal(mu, sigma);
• }";

```


	parameters	mean	mcse	std
1	"b"	5.05481	0.035551	1.26818
2	"a"	-13.4538	0.209502	7.64578
3	"sigma"	11.1433	0.091969	3.43567

```

• let
•   x = LinRange(1, 10, 10)
•   y = [1, 1, 2, 3, 5, 8, 13, 21, 34, 55]
•   global fake = DataFrame(x = x, y = y)
•   data = (N = nrow(fake), x = fake.x, y =
•         fake.y)
•   global m8_2s = SampleModel("m8_2s",
•         stan8_2)
•   global rc8_2s = stan_sample(m8_2s; data)
•         success(rc8_2s) && describe(m8_2s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
d.

```

	parameters	median	mad_sd	mean	std
1	"a"	-13.63	7.192	-13.454	7.64
2	"b"	5.097	1.176	5.055	1.26
3	"sigma"	10.44	2.645	11.143	3.43

```

• if success(rc8_2s)
•   post8_2s = read_samples(m8_2s,
•         :dataframe)
•   ms8_2s = model_summary(post8_2s, [:a,
•         :b, :sigma])
end

```

```

▶ [2.45738, 7.50822]

```

```

• quantile(post8_2s.b, [0.025, 0.975])

```

```

▶ [2.96928, 7.04544]

```

```

• quantile(post8_2s.b, [0.05, 0.95])

```

```
fake_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM
```

```
y ~ 1 + x
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	-13.8667	6.32766	-2.19	0.031
x	5.12121	1.01979	5.02	0.000

```
• fake_lm = lm(@formula(y ~ x), fake)
```