

See chapter 9 in Regression and Other Stories.

.....

```
• html"""
• <style>
•   main {
•     margin: 0 auto;
•     max-width: 2000px;
•     padding-left: max(160px, 10%);
•     padding-right: max(160px, 10%);
•   }
• </style>
• """
•
```

```
• using Pkg ✓ , DrWatson ✓
```

```
• begin
•   using GLM ✓
•
•   # Specific to ROSStanPluto
•   using StanSample ✓
•
•   # Graphics related
•   using GLMakie ✓
•
•   # Common data files and functions
•   using RegressionAndOtherStories ✓
• end
```

9.1 Propagating uncertainty in inference using posterior simulations.

```
hibbs =
```

	year	growth	vote	inc_party_candidate
1	1952	2.4	44.6	"Stevenson"
2	1956	2.89	57.76	"Eisenhower"
3	1960	0.85	49.91	"Nixon"
4	1964	4.21	61.34	"Johnson"
5	1968	3.02	49.6	"Humphrey"
6	1972	3.62	61.79	"Nixon"
7	1976	1.08	48.95	"Ford"
8	1980	-0.39	44.7	"Carter"
9	1984	3.86	59.17	"Reagan"
10	1988	2.27	53.94	"Bush, Sr."
⋮ more				
16	2012	0.95	52.0	"Obama"

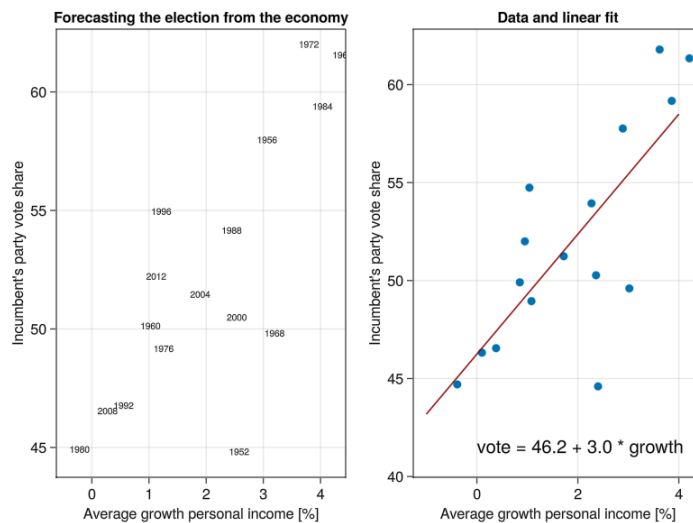
```
• hibbs =  
  CSV.read(ros_datadir("ElectionsEconomy",  
    "hibbs.csv"), DataFrame)
```

```
hibbs_lm =  
StatsModels.TableRegressionModel{LinearModel{GLM},  
  vote ~ 1 + growth
```

Coefficients:

	Coef.	Std. Error	t	Pr(> t)
(Intercept)	46.2476	1.62193	28.51	<1e-31
growth	3.06053	0.696274	4.40	0.0001

```
• hibbs_lm = lm(@formula(vote ~ growth),  
  hibbs)
```



```

let
  fig = Figure()
  hibbs.label = string.(hibbs.year)
  xlabel = "Average growth personal
income [%]"
  ylabel = "Incumbent's party vote share"
  let
    title = "Forecasting the election
from the economy"
    ax = Axis(fig[1, 1]; title, xlabel,
ylabel)
    for (ind, yr) in
      enumerate(hibbs.year)
        annotations!("$ (yr)"; position=
(hibbs.growth[ind],
hibbs.vote[ind]), textsize=10)
    end
  end
let
  x = LinRange(-1, 4, 100)
  title = "Data and linear fit"
  ax = Axis(fig[1, 2]; title, xlabel,
ylabel)
  scatter!(hibbs.growth, hibbs.vote)
  lines!(x, coef(hibbs_lm)[1] .+
coef(hibbs_lm)[2] .* x;
color=:darkred)
  annotations!("vote = 46.2 + 3.0 *
growth"; position=(0, 41))
end
fig
end

```

```

• stan7_1 = "
• data {
•   int<lower=1> N;      // total number of
•   observations
•   vector[N] growth;   // Independent
•   variable: growth
•   vector[N] vote;     // Dependent
•   variable: votes
• }
• parameters {
•   real b;              // Coefficient
•   independent variable
•   real a;              // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(50, 20);
•   b ~ normal(2, 10);
•   sigma ~ exponential(1);
•
•   mu = a + b * growth;
•
•   // likelihood including constants
•   vote ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std	
1	"b"	3.04361	0.0168508	0.671343	1
2	"a"	46.2857	0.0409229	1.56482	4
3	"sigma"	3.57977	0.012968	0.593587	2

```

• let
•   data = (N=nrow(hibbs), vote=hibbs.vote,
•         growth=hibbs.growth)
•   global m7_1s = SampleModel("hibbs",
•         stan7_1)
•   global rc7_1s = stan_sample(m7_1s; data)
•         success(rc7_1s) && describe(m7_1s)
end

```

```

/
c
Informational Message: The current Metropolis
j
ected because of the following issue:
Exception: normal_lpdf: Scale parameter is 0,
r/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T/
2, column 1 to column 26)
If this warning occurs sporadically, such as f
types like covariance matrices, then the sampl
but if this warning occurs often then your mod
conditioned or misspecified.

```

2095.17

```

• let
•   ss7_1s = describe(m7_1s)
•   ss7_1s[:sigma, :ess]
• end

```

2095.17

```

• let
•   ss7_1s = describe(m7_1s; showall=true)
•   ss7_1s[:sigma, :ess]
• end

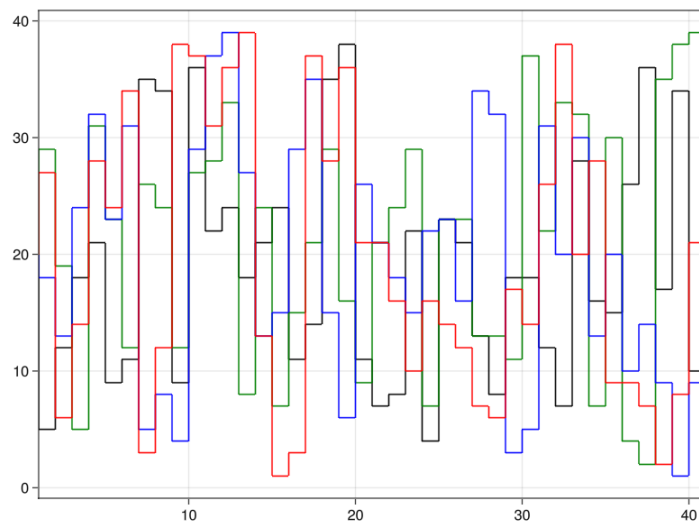
```

	parameters	mean	mcse	
1	"lp__"	-31.8687	0.0323058	1.2
2	"accept_stat__"	0.906547	0.00916869	0.1
3	"stepsize__"	0.440355	0.0325178	0.0
4	"treedepth__"	2.5685	0.0546857	0.6
5	"n_leapfrog__"	7.2355	0.550512	3.8
6	"divergent__"	0.0	NaN	0.0
7	"energy__"	33.3526	0.046026	1.7
8	"b"	3.04361	0.0168508	0.6
9	"a"	46.2857	0.0409229	1.5
10	"sigma"	3.57977	0.012968	0.5

```
• describe(m7_1s; showall=true)
```

post7_1s =	b	a	sigma
1	2.25991	46.2198	3.61128
2	2.6583	46.4954	3.61741
3	2.92437	46.0457	3.74011
4	3.06466	45.5862	3.46544
5	2.53448	48.3268	3.59055
6	2.62039	46.7069	2.99291
7	3.74541	44.948	2.76299
8	3.70214	45.8791	3.0808
9	2.49718	46.6594	3.00592
10	3.88758	44.4398	3.09775
⋮ more			
4000	3.43433	46.1684	3.08472

```
• post7_1s = success(rc7_1s) &&
  read_samples(m7_1s, :dataframe)
```



```
• trankplot(post7_1s, "b")
```

```
ms7_1s =
```

	parameters	median	mad_sd	mean	st
1	"a"	46.279	1.484	46.286	1.56
2	"b"	3.046	0.669	3.044	0.67
3	"sigma"	3.496	0.535	3.58	0.59

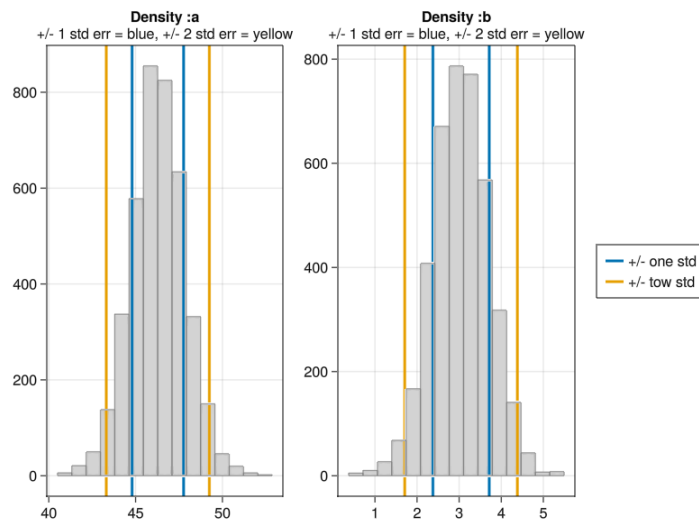
```
• ms7_1s = model_summary(post7_1s, [:a, :b, :sigma])
```

```
sims = 4000×3 Matrix{Float64}:
  2.25991  46.2198  3.61128
  2.6583   46.4954  3.61741
  2.92437  46.0457  3.74011
  3.06466  45.5862  3.46544
  2.53448  48.3268  3.59055
  2.62039  46.7069  2.99291
  3.74541  44.948   2.76299
  ⋮
  2.58214  46.6005  2.42432
  2.4851   47.9552  4.8092
  2.43727  47.225   4.6689
  2.81382  47.8815  4.89693
  3.07926  45.2446  3.94265
  3.43433  46.1684  3.08472
```

```
• sims = Array(post7_1s)
```

```
1x3 Matrix{Float64}:  
 3.04619 46.2789 3.49648
```

- `median(sims; dims=1)`

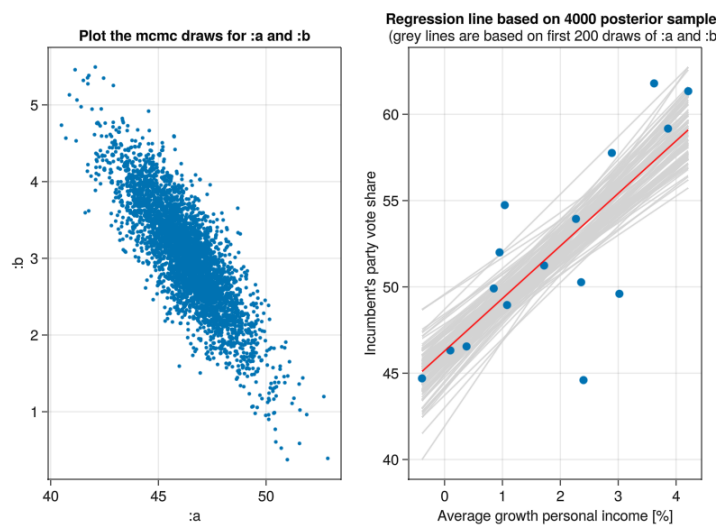


```

• let
•     f = Figure()
•     ax = Axis(f[1, 1]; title="Density :a",
•         subtitle="+/- 1 std err = blue, +/- 2
•             std err = yellow")
•     hist!(post7_1s.a; bins=15, color =
•         :lightgrey, strokewidth = 1, strokecolor
•             = :grey)
•     one = vlines!([ms7_1s[:a, :median] -
•         ms7_1s[:a, :mad_sd], ms7_1s[:a,
•             :median] + ms7_1s[:a, :mad_sd]];
•         linewidth=3)
•     two = vlines!([ms7_1s[:a, :median] -
•         2ms7_1s[:a, :mad_sd], ms7_1s[:a,
•             :median] + 2ms7_1s[:a, :mad_sd]];
•         linewidth=3)

•     ax = Axis(f[1, 2]; title="Density :b",
•         subtitle="+/- 1 std err = blue, +/- 2
•             std err = yellow")
•     hist!(post7_1s.b; bins=15, color =
•         :lightgrey, strokewidth = 1, strokecolor
•             = :grey)
•     vlines!([ms7_1s[:b, :median] -
•         ms7_1s[:b, :mad_sd], ms7_1s[:b,
•             :median] + ms7_1s[:b, :mad_sd]];
•         linewidth=3)
•     vlines!([ms7_1s[:b, :median] -
•         2ms7_1s[:b, :mad_sd], ms7_1s[:b,
•             :median] + 2ms7_1s[:b, :mad_sd]];
•         linewidth=3)
•     Legend(f[1, 3], [one, two], ["+/- one
•         std", "+/- tow std"])
•     f
• end

```



```

let
  growth_range =
  LinRange(minimum(hibbs.growth),
  maximum(hibbs.growth), 200)
  votes = mean.(link(post7_1s, (r,x) ->
  r.a + x * r.b, growth_range))

  xlabel = "Average growth personal
  income [%]"
  ylabel = "Incumbent's party vote share"

  fig = Figure()

  ax = Axis(fig[1, 1]; title="Plot the
  mcmc draws for :a and :b", xlabel=":a",
  ylabel=":b")
  scatter!(post7_1s.a, post7_1s.b;
  markersize=4)

  xlabel = "Average growth personal
  income [%]"
  ylabel="Incumbent's party vote share"
  ax = Axis(fig[1, 2]; title="Regression
  line based on 4000 posterior samples",
  subtitle = "(grey lines are based
  on first 200 draws of :a and :b)",
  xlabel, ylabel)
  for i in 1:100
    lines!(growth_range, post7_1s.a[i]
    .+ post7_1s.b[i] .* growth_range,
    color = :lightgrey)
  end
  scatter!(hibbs.growth, hibbs.vote)
  lines!(growth_range, votes, color =
  :red)
  fig

```

end

9.2 Prediction and uncertainty.

	x	y
1	-2.0	50
2	-1.0	44
3	0.0	50
4	1.0	47
5	2.0	56

```
• let
•   x = LinRange(-2, 2, 5)
•   y = [50, 44, 50, 47, 56]
•   global sexratio = DataFrame(x = x, y =
•   y)
• end
```

```

• stan9_1 = "
• data {
•   int<lower=1> N; // total number of
•   observations
•   vector[N] x;    // Independent
•   variable: growth
•   vector[N] y;    // Dependent variable:
•   votes
• }
• parameters {
•   real b;          // Coefficient
•   independent variable
•   real a;          // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(50, 5);
•   b ~ normal(0, 5);
•   sigma ~ uniform(0, 10);
•
•   mu = a + b * x;
•
•   // likelihood including constants
•   y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std	
1	"b"	1.34845	0.0468004	1.69112	-1
2	"a"	49.5631	0.0380842	2.20566	46
3	"sigma"	5.47735	0.103726	1.92207	2.

```

• let
•   data = (N = nrow(sexratio), x =
•   sexratio.x, y = sexratio.y)
•   global m9_1s = SampleModel("m9_1s",
•   stan9_1)
•   global rc9_1s = stan_sample(m9_1s; data)
•   success(rc9_1s) && describe(m9_1s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/7
d.

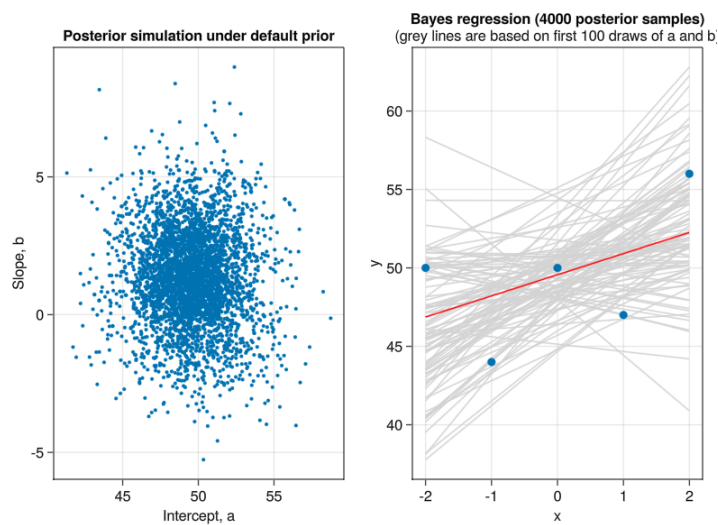
```

	parameters	median	mad_sd	mean	st
1	"a"	49.578	2.043	49.563	2.20
2	"b"	1.356	1.539	1.348	1.69
3	"sigma"	5.159	2.029	5.477	1.92

```

• if success(rc9_1s)
•   post9_1s = read_samples(m9_1s,
•   :dataframe)
•   sm9_1s = model_summary(post9_1s, [:a,
•   :b, :sigma])
end

```



```

let
  x_range = LinRange(minimum(sexratio.x),
    maximum(sexratio.x), 200)
  y = mean.(link(post9_1s, (r,x) -> r.a +
    x * r.b, x_range))

  xlabel = "x"
  ylabel = "y"

  fig = Figure()

  ax = Axis(fig[1, 1]; title="Posterior
simulation under default prior",
  xlabel="Intercept, a", ylabel="Slope,
b")
  scatter!(post9_1s.a, post9_1s.b;
  markersize=4)

  ax = Axis(fig[1, 2]; title="Bayes
regression (4000 posterior samples)",
  subtitle = "(grey lines are based
on first 100 draws of a and b)",
  xlabel, ylabel)
  for i in 1:100
    lines!(x_range, post9_1s.a[i] .+
      post9_1s.b[i] .* x_range, color =
      :lightgrey)
  end
  scatter!(sexratio.x, sexratio.y)
  lines!(x_range, y, color = :red)
  fig
end

```

```

• stan9_2 = "
• data {
•   int<lower=1> N; // total number of
•   observations
•   vector[N] x;    // Independent
•   variable: growth
•   vector[N] y;    // Dependent variable:
•   votes
• }
• parameters {
•   real b;          // Coefficient
•   independent variable
•   real a;          // Intercept
•   real<lower=0> sigma; // dispersion
•   parameter
• }
• model {
•   vector[N] mu;
•
•   // priors including constants
•   a ~ normal(48.8, 0.2);
•   b ~ normal(0, 0.2);
•   sigma ~ uniform(0, 10);
•
•   mu = a + b * x;
•
•   // likelihood including constants
•   y ~ normal(mu, sigma);
• }";

```

	parameters	mean	mcse	std
1	"b"	0.03449	0.00334336	0.190924
2	"a"	48.808	0.00383194	0.204215
3	"sigma"	5.07035	0.0355242	1.72614

```

• let
•   data = (N = nrow(sexratio), x =
•   sexratio.x, y = sexratio.y)
•   global m9_2s = SampleModel("m9_2s",
•   stan9_2)
•   global rc9_2s = stan_sample(m9_2s; data)
•   success(rc9_2s) && describe(m9_2s)
end

```

```

/var/folders/l7/pr04h0650q5dvqtnvs8s2c00000gn/T
d.

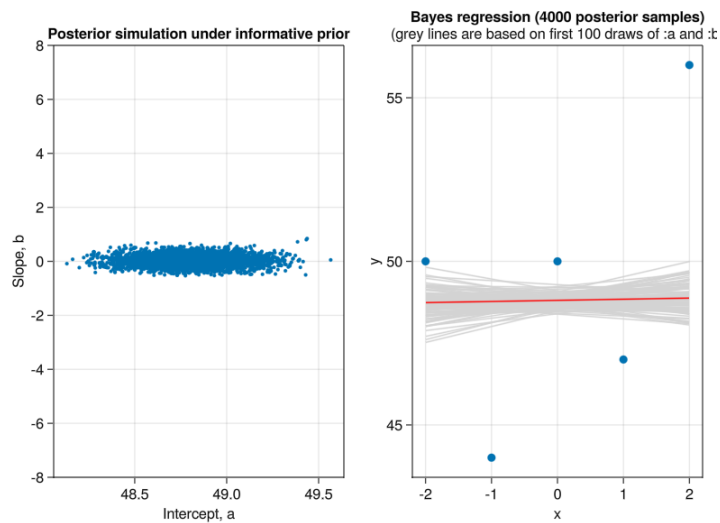
```

	parameters	median	mad_sd	mean	std
1	"a"	48.808	0.199	48.808	0.20
2	"b"	0.036	0.186	0.034	0.19
3	"sigma"	4.738	1.669	5.07	1.72

```

• if success(rc9_2s)
•   post9_2s = read_samples(m9_2s,
•   :dataframe)
•   sm9_2s = model_summary(post9_2s, [:a,
•   :b, :sigma])
end

```

```

let
  x_range = LinRange(minimum(sexratio.x),
    maximum(sexratio.x), 200)
  y = mean.(link(post9_2s, (r,x) -> r.a +
    x * r.b, x_range))

  xlabel = "x"
  ylabel = "y"

  fig = Figure()

  ax = Axis(fig[1, 1]; title="Posterior
simulation under informative prior",
  xlabel="Intercept, a", ylabel="Slope,
  b")
  ylims!(ax, -8, 8)
  scatter!(post9_2s.a, post9_2s.b;
  markersize=4)

  ax = Axis(fig[1, 2]; title="Bayes
regression (4000 posterior samples)",
  subtitle = "(grey lines are based
on first 100 draws of :a and :b)",
  xlabel, ylabel)
  for i in 1:100
    lines!(x_range, post9_2s.a[i] .+
      post9_2s.b[i] .* x_range, color =
      :lightgrey)
  end
  scatter!(sexratio.x, sexratio.y)
  lines!(x_range, y, color = :red)
  fig
end

```

9.3 Prior information and Bayesian synthesis.

Prior based on a previously-fitted model using economic and political condition.

```
• begin
•   theta_hat_prior = 0.524
•   se_prior = 0.041
• end;
```

Survey of 400 people, of whom 190 say they will vote for the Democratic candidate.

```
• begin
•   n = 400
•   y = 190
• end;
```

Data estimate.

```
theta_hat_data = 0.475
```

```
• theta_hat_data = y/n
```

```
se_data = 0.02496873044429772
```

```
• se_data =  $\sqrt{((\underline{y/n}) * (1 - \underline{y/n}) / \underline{n})}$ 
```

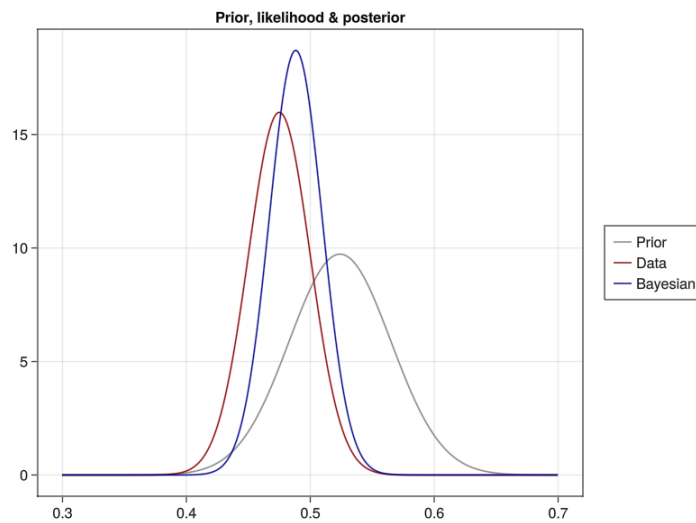
Bayes estimate.

```
theta_hat_bayes = 0.48825635323153693
```

```
• theta_hat_bayes =
•   (theta_hat_prior/se_prior2 +
•    theta_hat_data/se_data2) / (1/se_prior2
•    + 1/se_data2)
```

```
se_bayes = 0.02132543263776925
```

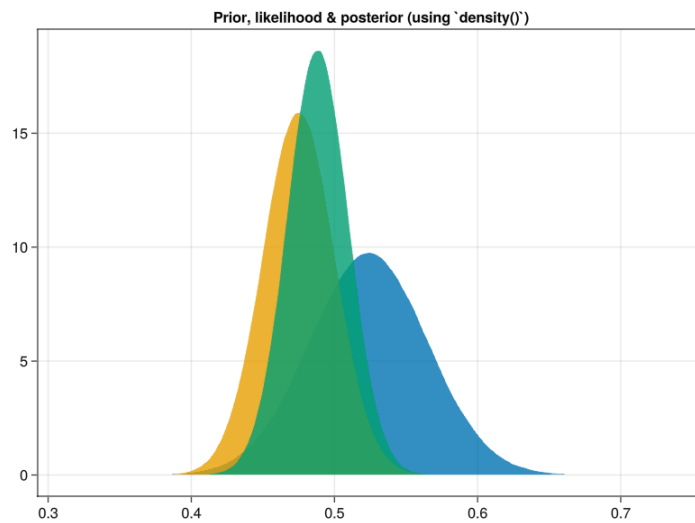
```
• se_bayes =  $\sqrt{1 / (1 / \underline{se\_prior}^2 + 1 / \underline{se\_data}^2)}$ 
```



```

• let
•   x = 0.3:0.001:0.7
•   f = Figure()
•   ax = Axis(f[1, 1], title="Prior,
•   likelihood & posterior")
•   prior = lines!(f[1, 1], x, pdf.
•   (Normal(theta_hat_prior, se_prior), x),
•   color=:gray)
•   data = lines!(x, pdf.
•   (Normal(theta_hat_data, se_data),
•   x),color=:darkred)
•   bayes = lines!(x, pdf.
•   (Normal(theta_hat_bayes, se_bayes), x),
•   color=:darkblue)
•   Legend(f[1, 2], [prior, data, bayes],
•   ["Prior", "Data", "Bayesian"])
•   current_figure()
end

```



```

• let
•   f = Figure()
•   ax = Axis(f[1, 1], title="Prior,
•   likelihood & posterior (using
•   `density()``)"
•   density!(rand(Normal(theta_hat_prior,
•   se_prior), Int(1e6)), lab="prior")
•   density!(rand(Normal(theta_hat_data,
•   se_data), Int(1e6)), lab="likelihood")
•   density!(rand(Normal(theta_hat_bayes,
•   se_bayes), Int(1e6)), lab="bayes")
•   current_figure()
• end

```

9.4 Example of Bayesian inference: beauty and sex ratio.